

國立臺灣大學電機資訊學院電子工程學研究所

博士論文

Graduate Institute of Electronics Engineering

College of Electrical Engineering and Computer Science

National Taiwan University

Doctor Thesis

合作, 競爭關係以及密集錯誤回復力的模型驗證

Model Checking Collaboration, Competition and Dense Fault

Resilience

黃重豪

Chung-Hao Huang

指導教授：王凡博士

Advisor: Farn Wang, Ph.D.

中華民國 105 年 1 月

January, 2016

國立臺灣大學  
電子工程學研究所

博士論文

合作、競爭關係以及密集錯誤回復力的模型驗證

黃重豪  
撰

# Abstract

In this thesis, I introduce BSIL(basic strategy-interaction logic) and TCL(temporal cooperation logic) which can help in formally define and verify the strategy interaction property of a game. The former, BSIL, is an extension to ATL (alternating-time logic) for the specification of strategies interaction of players in a system. BSIL is able to describe one system strategy that can cooperate with several strategies of the environment for different requirements. Such properties are important in practice and I show that such properties are not expressible in  $ATL^*$ , GL (game logic), and AMC (alternating  $\mu$ -calculus). Specifically, BSIL is more expressive than ATL but incomparable with  $ATL^*$ , GL, and AMC in expressiveness. I show that, for fulfilling a specification in BSIL, a memoryful strategy is necessary. I also show that the model-checking complexity of BSIL is PSPACE-complete and is of lower complexity than those of  $ATL^*$ , GL, AMC, and the general strategy logics. Which may imply that BSIL can be useful in closing the gap between large scale real-world projects and the time consuming game-theoretical results. I then show the feasibility of our techniques by implementation and experiment with our PSPACE model-checking algorithm for BSIL. On the other hand, TCL allows successive definition of strategies for agents and agencies. Like BSIL the expressiveness of TCL is still incomparable with  $ATL^*$ , GL and AMC. However, it can describe deterministic Nash equilibria while BSIL cannot. I prove that the model checking complexity of TCL is EXPTIME-complete. TCL enjoys this relatively cheap complexity by disallowing a too close en-

tanglement between cooperation and competition while allowing such entanglement leads to a non-elementary complexity. I have implemented a model-checker for TCL and shown the feasibility of model checking in the experiment on some benchmarks. Although BSIL and TCL have decent expressive power and benefit from relatively low complexity. PSPACE-complete and EXPTIME-complete is still not good enough for real problem. To adopt the game concept to real world problem, I introduce an algorithm to calculate the highest degree of fault tolerance a system can achieve with the control of a safety critical systems. Which can be reduced to solving a game between a malicious environment and a controller. During the game play, the environment tries to break the system through injecting failures while the controller tries to keep the system safe by making correct decisions. I found a new control objective which offers a better balance between complexity and precision for such systems: we seek systems that are  $k$ -resilient. A system is  $k$ -resilient means it is able to rapidly recover from a sequence of small number, up to  $k$ , of local faults infinitely many times if the blocks of up to  $k$  faults are separated by short recovery periods in which no fault occurs.  $k$ -resilience is a simple abstraction from the precise distribution of local faults, but I believe it is much more refined than the traditional objective to maximize the number of local faults. I will provide detail argument of why this is the right level of abstraction for safety critical systems when local faults are few and far between. I have proved, with respect to resilience, the computational complexity of constructing optimal control is low. And a demonstration of the feasibility through an implementation and experimental results will be in following chapters.

Key words:

# Contents

<b>Abstract</b>	<b>i</b>
<b>Contents</b>	<b>iii</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Related Work</b>	<b>12</b>
2.1 Existed Logics about Game and Strategy . . . . .	12
2.1.1 Prior to Strategy Logics . . . . .	12
2.1.2 With Strategy Logics . . . . .	13
2.2 Fault tolerance . . . . .	13
<b>3 Back Ground</b>	<b>15</b>
3.1 Concurrent Games . . . . .	15
3.2 Turn-Based Games . . . . .	17
3.3 Two-player concurrent game structures . . . . .	18
3.4 Fault Tolerance . . . . .	18
<b>4 Basic Strategy-interaction Logic</b>	<b>22</b>
4.1 Running Example . . . . .	22
4.1.1 Trying to Write Down a Correct Formal Specification . . . . .	22
4.1.2 Resorting to General Strategy Logics for a Correct Specification .	24
4.1.3 BSIL: New Strategy Modalities Expressive Enough for the Spec- ification . . . . .	25

4.1.4	Symbolic Strategy Names and Path Obligations . . . . .	25
4.1.5	Passing Down the Path Obligations While Observing the Restrictions Among S-Profiles . . . . .	27
4.1.6	Finding Finite Satisfying Evidence for a Formula in a Computation Tree . . . . .	29
4.2	BSIL . . . . .	30
4.2.1	Syntax . . . . .	30
4.2.2	semantics . . . . .	31
4.2.3	ATL+ . . . . .	32
4.2.4	Memory . . . . .	32
4.3	Expressive Power of BSIL . . . . .	33
4.3.1	Comparison with GL . . . . .	33
4.3.2	Comparison with AMC . . . . .	36
4.3.3	Comparison with ATL* . . . . .	39
4.4	Algorithm and Complexity . . . . .	39
4.4.1	Computing path obligations and passing them down the computation tree . . . . .	41
4.4.2	Procedures for checking BSIL properties . . . . .	46
4.4.3	Correctness proof of the algorithm . . . . .	48
4.4.4	Complexities of the algorithm . . . . .	50
4.4.5	Lower bound and completeness . . . . .	52
<b>5</b>	<b>Temporal Cooperation Logic</b>	<b>55</b>
5.1	TCL . . . . .	55
5.1.1	Syntax . . . . .	56
5.1.2	semantics . . . . .	57
5.2	Expressive Power of TCL . . . . .	60
5.3	Complexity of TCL . . . . .	63
<b>6</b>	<b>Software Resilience against Dense Errors</b>	<b>66</b>

6.1	Resilience in a Nutshell . . . . .	66
6.2	Safety resilience games . . . . .	70
6.3	Alternating-time $\mu$ -calculus with events . . . . .	75
6.3.1	Syntax . . . . .	77
6.3.2	semantics . . . . .	78
6.4	Resilience level checking algorithm . . . . .	80
6.4.1	High-level description of the algorithm . . . . .	80
6.4.2	Realization with AMCE model-checking . . . . .	80
6.4.3	Complexity . . . . .	85
<b>7</b>	<b>Implementation Experiment</b>	<b>90</b>
7.1	BSIL and TCL . . . . .	90
7.1.1	Implementation . . . . .	90
7.1.2	Benchmarks and their experiment report . . . . .	90
7.2	Resilience . . . . .	94
7.3	Tool implementation . . . . .	94
7.3.1	Benchmarks and their experiment report . . . . .	95
7.3.2	Modelling of the fault-tolerant systems . . . . .	96
<b>8</b>	<b>Conclusions</b>	<b>103</b>
8.1	Conclusion . . . . .	103
	<b>Bibliography</b>	<b>105</b>

# Chapter 1

## Introduction

There are many open systems such as bank system, web service and HMI in the factory which take input from users. Currently, the specification and verification of such systems focuses on the design of system interfaces that allow for the fulfillment of some objectives of the users while enforce certain safety policies of the systems. Intuitively, such system can be modeled as games played by users and the system and the verification of such systems can benefit from game-based techniques. From the game theoretical perspective, the design problem of such open systems can be modeled as multi-agent games. Some agents represent the system, while other agents represent the environment (or the users). The system wins the game in an execution if the user's objectives allowed by the system are indeed fulfilled and his disallowed objectives do not happen. The goal of system design is to design a computable strategy of the system that enforces all the system objectives and also allows the users to achieve those objectives consistent with the system objectives. Such a strategy is called a *winning strategy* for the system.

To achieve such goal, we create 2 temporal logic which include strategy quantifier in the syntax, *BSIL* (*basic strategy-interaction logic*) and *TCL* (*temporal cooperation logic*). At the moment, there are various logic that express such properties of strategic power of agents, including *ATL* (*alternating-time logic*), *ATL\**, *AMC* (*alternating  $\mu$ -calculus*), *GL* (*game logic*) [3], and *SL* (*strategy logics*) [4–6], for the specification of open systems. Each language also comes with a verification algorithm that helps to decide whether or not a winning strategy for the system exists. There is, however, a gap between the indus-



trial need for efficient algorithms (and solvers) and the available technology offered from previous research. Frankly speaking, none of those languages represents a proper combination of expressiveness for close interaction among agent strategies and efficiency for specification verification. ATL, ATL\*, AMC, and GL [3] allow us to specify that some players together have a strategy to satisfy some fully temporalised objective: strategy quantifiers mark the start of a state formula. As exemplified below, this is far from what the industry needs in specification.

Consider the example of a bank that need specify their information system embodied as a system security strategy and allowing a client to use a strategy to withdraw money, to use a strategy to deposit money, and to use a strategy to query for balance. Moreover, the same system strategy should forbid any illegal operation on the banking system. Specifically, the same system strategy must accommodate all strategies of the client for 'good behavior' (i.e, behavior in line with the specification), while blocking all strategies of the client that refer to undesired behavior, and thus preventing the client from damaging the system. We will show that *ATL*, *ATL\**, and *GL* [3] do not support such specifications. For example, it is not possible to specify in those languages that the system strategies used both in a withdrawal transaction and in a deposit transaction must be the same. Consequently, verification techniques for specifications in those languages cannot capture such real-world objectives for open systems.

To solve the expressiveness problem in the above example, strategy logics were proposed in [4–6] that allow for the flexible quantification of high-order strategy variables in logic formulas. However, the verification complexities are prohibitively high and hinder them from practical application. In retrospect, the strategy profiles in such strategy logics can be combined in unrestricted ways. For example, in [6], we can write down the following artificial and hypothetical property.

$$\langle\langle X \rangle\rangle \llbracket Y \rrbracket \langle\langle Z \rangle\rangle \Box (((1, X) \diamond p) \wedge \bigcirc (2, Y) \diamond q) \rightarrow (((2, Z) \diamond q) \wedge \neg (3, Z) \Box q))$$

Here  $\langle\langle X \rangle\rangle$  and  $\langle\langle Z \rangle\rangle$  declare the existence of strategies named  $X$  and  $Z$  respectively.  $\llbracket Y \rrbracket$  is a universal quantification on a strategy named  $Y$ . Then operator  $(1, X)$ ,  $(2, Y)$ ,  $(2, Z)$ , and  $(3, Z)$  respectively bind strategy  $X$ ,  $Y$ ,  $Z$ , and  $Z$  to agents 1, 2, 2, and 3. As can be seen, the

language is very free in style. According to the experiences in temporal logic development, usually proper restrictions in the modal operations can lead to a spectrum of sub-logics with different expressiveness and model-checking efficiency. The following examples shows that subclasses of temporal logics with proper balance between expressiveness and verification efficiency are not only theoretically interesting and can also be practically useful.

- The validity problem of  $\mathcal{L}$  (the first-order language of unary predicate symbols and the binary predicate symbol  $\leq$ ) is non-elementary [7] while PTL, with the same expressiveness as  $\mathcal{L}$ , is only PSPACE-complete [8].
- Between CTL [9] and CTL\* [10, 11], there are many subclasses of CTL\* with various balances between expressiveness and verification efficiency. Fair CTL [12], as a natural class between CTL and CTL\*, is expressive enough for many practical specifications and still enjoys a polynomial time model-checking complexity. There are also other subclasses of CTL\* with various balance considerations [11, 13–15].

Indeed, most specifications in real-world projects come in simple structures, for example, safety, liveness, etc. Thus it would be interesting to see what "natural" subclasses of strategy logics can be identified with a proper balance between expressiveness and model-checking efficiency. Moreover, it would be practical and appealing if the subclass can be characterized with elegant syntax. The purpose of this manuscript is to propose new natural modal operators for strategy collaborations and extend ATL for a subclass of strategy logic.

In the following, we use the classical prisoner's dilemma to explain how we design new modal operators for structured strategy collaboration to achieve a balance between expressiveness and model-checking efficiency.

**Example 1. 1 Prisoner's dilemma** Suppose the police is interrogating three suspects (prisoners). The police has very little evidence. A prisoner may cooperate (with his/her peers) and deny all charges made by the police. If all deny, they are all acquitted of all charges. However, each prisoner may choose to betray his/her peers and provide the

police with evidence. If more than one prisoner choose to betray their peers, all will be sentenced and stay in jail. If only one chooses to betray, then the other prisoners will stay in jail, while (s)he will be a ‘dirty witness’ and all charges against him/her will be dropped.

One may want to specify that the three prisoners can cooperate with each other (by denying all charges), and will not be in jail. Let  $j_a$  be the proposition for prisoner  $a$  in jail. This can be expressed in Alur, Henzinger, and Kupferman’s ATL\*, GL, or AMC [3], respectively, as follows.<sup>1</sup>

$$\begin{aligned} \text{ATL}^*: \quad & \langle 1, 2, 3 \rangle \wedge_{a \in [1,3]} \Diamond \neg j_a \\ \text{GL}: \quad & \exists \{1, 2, 3\}. \wedge_{a \in [1,3]} \forall \Diamond \neg j_a \\ \text{AMC}: \quad & \mathbf{lfp}x. \langle 1, 2, 3 \rangle \bigcirc \wedge_{a \in [1,3]} (x \vee \neg j_a) \end{aligned}$$

Here “ $\langle 1, 2, 3 \rangle$ ” and “ $\exists \{1, 2, 3\}$ ” are both existential quantifiers on the collaborative strategy among prisoners 1, 2, and 3. Such a quantifier is called a *strategy quantifier (SQ)* for convenience. Operator ‘**lfp**’ is the least fixed point operator. Even though one can specify strategies employed by sets of prisoners and there is a natural relationship (containment) between sets with such logics, there is no way to relate strategies to each other. For example, if prisoners 1 and 2 are unconditional loyal to prisoner 3, they can both deny the charges, make sure that prisoner 3 will not be in jail, and let prisoner 3 to decide whether they will be in jail. □

The research of strategies for related properties has a long tradition in game theory. It is easy to see the similarity and link between the specification problems for the prisoner’s dilemma and the banking system. This observation suggests that finding a language with an appropriate and natural balance between the expressive power and the verification complexity of a specification language is a central challenge.

To meet this challenge, we propose 2 extension logic of ATL, called *BSIL (basic strategy-interaction logic)* and *TCL (temporal cooperation logic)*. In a first step, we extend ATL to  $\text{ATL}^+$ , where  $\text{ATL}^+$  is the natural extension obtained by allowing for Boolean connectives of path quantifiers. (Cf. [13, 14] for the similar extension of CTL to  $\text{CTL}^+$ .) We then introduce a new modal operator called *strategy interaction quantifier (SIQ)*. In

---

<sup>1</sup>Note that the three example formulas are not equivalent.

the following, we use several examples in the prisoner's dilemma to explain BSIL and TCL, starting with the following specification in BSIL for the property discussed at the end of Example 1.

$$\langle 1, 2 \rangle ((\langle + \rangle \Diamond \neg j_3) \wedge (\langle +3 \rangle \Diamond \neg (j_1 \vee j_2)) \wedge \langle +3 \rangle \Box (j_1 \wedge j_2)) \quad (\text{A})$$

Here " $\langle +3 \rangle$ " is an existential SIQ that reasons over strategies of Prisoner 3 for collaborating with the strategies of Prisoners 1 and 2 introduced by the parent SQ " $\langle 1, 2 \rangle$ ". Similarly, " $\langle + \rangle$ " means that no collaboration of any prisoner is needed. (For conciseness, we omit " $\langle + \rangle$ " in the following.) We also call an SIQ an SQ. In BSIL formulas, we specifically require that no SIQ can appear as a topmost SQ in a path subformula. Further more, BSIL does not allow any temporal logic operator appears in the middle of 2 SIQs while TCL does. The loosen of the restriction brings TCL stronger expressive power as long as a higher verification complexity.

As can be seen, SIQ imposes a hierarchical style of strategy collaboration which seems natural for practical specification. Consider another example. If Prisoner 1 hates the others, he can always betray other prisoners, making sure that Prisoners 2 and 3 will be in jail, and let them decide whether he will be in jail. This property can be expressed in BSIL as follows.

$$\langle 1 \rangle ((\Box (j_2 \wedge j_3)) \wedge (\langle +2, 3 \rangle \Diamond \neg j_1) \wedge \langle +2, 3 \rangle \Box j_1) \quad (\text{B})$$

One restriction of BSIL and TCL is that no negation between an SIQ and its parent SIQ or SQ is allowed. This restriction then also forbids universal SIQ. While at first glance, it seems less than elegant in mathematics, it is both necessary for verification complexity and compatible with practical specification styles. When we take a closer look, in fact, there is an implicit universal SIQ at the end of every maximal syntax path from an SQ to its descendant SIQ. Thus, BSIL and TCL allows for the specification of the ways in combining system strategy profiles that can be computed statically to enforce the system policy against any hostile strategy profile of those agents not participating in the the system strategy profiles. If explicitly allowing for universal SIQs in BSIL and TCL, then the strategy profiles associated with existential SIQs in the scope of a universal SIQs

may have to be computed dynamically. Thus the explicit universal SIQs will not only necessarily blow up the verification complexity, but also contradict the main-stream of game theory for statically calculable strategies. In contrast, some strategy logics [6] allow for the specification of strategy profiles that are not statically calculable.

In this work, we establish that BSIL and TCL are incomparable with  $\text{ATL}^*$ , GL, and AMC in expressiveness. Although the strategy logics [4–6] are superclasses to BSIL and TCL with their flexible quantification of strategies and binding to strategy variables, their model-checking complexity are all doubly exponential time hard. In contrast, BSIL enjoys a PSPACE-complete model-checking complexity and TCL has a EXPTIME-complete model checking complexity for turn-based and concurrent game graphs. This may imply that BSIL and TCL could be a better balance between expressiveness and verification efficiency than  $\text{ATL}^*$ , GL, AMC [3], and SL [5, 6]. Further related work is the stochastic game logic (SGL) by Baier, Brázdil, Größer, and Kucera [16], which allows for expressing strategy interaction. However, for memoryful strategies, the model-checking problem of SGL is undecidable.

We also establish some other properties of BSIL and TCL. We show that the strategies for BSIL and TCL properties against turn-based games need to be memoryful. We prove that the BSIL model-checking problem is PSPACE-complete and TCL model-checking problem is EXPTIME-complete. However, the PSPACE model-checking algorithm of BSIL need enumerate the labellings on computation trees and may suffer from high time complexity. We present an alternative model-checking algorithm with time complexity quadratic in the size of a game graph and exponential only in the size of a BSIL specification. We also present an alternative model-checking algorithm with time complexity quadratic in the size of a game graph and double exponential in the size of a TCL specification. We also establish that the BSIL realisability problem is complete for doubly exponential time.

However, there can be million lines of code in today’s software system. Various empirical studies show that the defect density of commercial software system is around 1 to 20 defects in every 1000 lines of source code [1]. On such a scale of complexity, de-

$k$	0	1	2	3	4	5	6	...
$k$ errors	0.865	0.594	0.333	0.143	0.053	0.017	0.005	...
$k$ dense errors	0.865	$2 \cdot 10^{-4}$	$2 \cdot 10^{-9}$	$2 \cdot 10^{-14}$	$2 \cdot 10^{-19}$	$2 \cdot 10^{-24}$	$2 \cdot 10^{-29}$	...

Table 1.1: Probabilities of  $k$  dense errors

fects in the source codes are unavoidable and, even with the PSPACE model checking algorithm of BSIL, the computing power of today's machine is far from enough to verify the systems. To lower the verification complexity, we try to focus on a specific type of property which is highly connected to the game theoretical perspective, the fault tolerance problem. To make the system more reliable, developers of systems have already created many engineering techniques to contain the damage that could be caused by such defects. For example, a software system may have several measures to its disposal to avoid system failure, including resending the request, resetting the server, clearing the communication buffers, and etc when observing that a critical service request is not acknowledged. However, in general, since all the recovery cost time and money it is important to estimate how to organize the measures for the maximal resilience of the system against realistic errors. At the moment, an automated support which can suggest defence techniques to development teams is missing. We created a game theoretic approach to study this problem and carried out experiments to show how this approach can be helpful in synthesizing the most resilient defence of software systems against multiple errors.

The naive way to measure the safety level of a system is to find the number of errors that it can endure before running into failure state. But in second thought, no non-trivial system can handle unlimited errors without degrading to inevitable system failure. Thus, it would be meaningless to analyse the resilience level of the systems to software errors can proceed without creating a realistic error model in which practical control mechanism can be devised to defend the systems against errors. In this work, we are interested in defending the system against a more restricted error model, but still let the error model has a quantifiable level of power in order to simulate different error scenarios. Further more, we think a reasonable foundation need to take into consider that the life-time of a software system is much longer than the duration needed for a reasonably designed software system

to recover from an error. Therefore, we propose to evaluate control mechanism of software systems on how many errors the control can endure before recovery to safe states. We then present an algorithm to find a control strategy that can handle the maximum number of such errors.

Let us standardize the basic terms before proceeding further. A design defect in software or hardware is called a *fault* in embedded systems. An *error* (sometimes called component failure in the literature) is the effect of a fault that causes a difference between the expected and the actual behavior of a software system, e.g., measurement errors, read/write errors, etc. An *error* does not always lead to a system failure, but may instead be repaired by, e.g., a defence mechanism in the software. That is, an *error* may be detected and fixed/neutralized before it creates any harm to the system or its users. A *failure* is the fact that users can observe the faulty behavior created by *errors*.

Our specific goal is to develop a technique for finding a control mechanism of a software system which can against the maximal number of dense errors without degrading to failure. Our inspiration is from methods for resilient avionic systems [2], where fault tolerance is designed to recover from a bounded number of errors. The number of errors a system needs to tolerate is calculated from the mean time between errors of individual components and the maximal duration of the system. We use the quality guarantees one obtains for an airplane(the system) as an example to demonstrate the difference between the objective to tolerate up to  $k$  *errors* and sequences of separated blocks of up to  $k$  *dense errors* in a short period. Assuming the operating time of the system is 20 hours, the mean time between exponentially distributed errors is 10 hours and the repair time is 3.6 seconds. The mean time between dense errors (consecutive errors before system recovery) is calculated in Table 1.1.

The figures for  $k$  errors (component failures) are simply the values for the Poisson distribution with coefficient 2. To explain the figures for  $k$  dense errors, consider the density of 2 dense errors occurring in close succession. If an error occurs, the chance that the next error occurs within the repair time (3.6 seconds) is approximately  $\frac{1}{10000}$ . The goal to tolerate an arbitrary number of up to  $k$ -dense errors is much harder than the goal of

tolerating up to  $k$  errors, but, as the example shows, the number  $k$  can be much smaller. Tolerating an arbitrary number of errors (with a distance of at least 3.6 seconds between them) creates the same likelihood to result in a system failure as tolerating up to 9 errors overall, and tolerating up to 15 errors still results in a 70% higher likelihood of a system failure than tolerating blocks of up to 2 errors in this example. Only errors for which this is the case could cause a system failure. The mean time between blocks of two dense errors is therefore not ten hours, but 100,000 hours. Likewise, it increases to 1,000,000,000 (one billion) hours for blocks of three dense errors, and so forth.

Maximizing the number of dense errors that are permitted before full recovery is therefore a natural design goal. The system is allowed again the same number of errors after full recovery. Now, if the *mean time between errors* (*MTBE*) is huge compared to the time the system needs to fully recover, then the mean time between system failures (*MTBF*) grows immensely.

We treat the problem of designing a resilient control mechanism against dense errors as a two-player game, called *safety resilience game*, between the system (protagonist<sup>2</sup>) and a hostile agent (antagonist<sup>3</sup>) that injects errors into the system under execution. Intuitively, the protagonist wants to keep the system from failure in the presence of errors, while the antagonist wants to derail the system to failure.

Specifically, system designers may model their system, defence mechanism, and error model as a finite game graph. The nodes in the graph represent system states. These system states are partitioned into three classes: the safe states, the failure states, and the recovery states. Some transitions are labelled with errors while others are considered normal transitions. The game is played with respect to a resilience level  $k$ . If a play ever enters a failure state, then the antagonist wins in the play. Otherwise, the protagonist wins.

On a certain state in the game, the protagonist plays by selecting a move, intuitively the corresponding 'normal' event that should happen next (unless an error is injected). The antagonist can then decide to trigger an error transition (injecting an error) with the intention to eventually deflect the system into a failure state. The error model, however,

---

<sup>2</sup>In game theory, a protagonist sometimes is also called *player 1*.

<sup>3</sup>In game theory, an antagonist sometimes is also called *player 2*.



restricts the antagonist to inject at most  $k$  errors before she allows for a long period of time that the system may use to recover to the safe states. (If the antagonist decides to use less than  $k$  errors, the protagonist does not know about this. It proves that this information is not required, as we will show that the protagonist can play memoryless.) After full recovery by the protagonist to the safe states, the antagonist is allowed again to inject the same number of errors, and so forth.

If the system can win this game, then the system is called *k-resilient*. For *k-resilient* systems, there exists a control strategy—even one that does not use memory—to make the system resilient in the presence of blocks of up to  $k$  dense errors. The idea is, if the component MTBF is huge compared to the time the system needs to fully recover, then the expected time for system breakdown grows immensely.

Besides formally defining safety resilience games, we also present algorithms for answering the following questions.

- Given an integer  $k$ , a set  $F$  of failure states, and a set  $S$  of safe states (disjoint from  $F$ ), is there a recovery mechanism that can endure up to  $k$  dense errors, effectively avoid entering  $F$ , and quickly direct the system back to  $S$ . Sometimes, the system designers may have designated parts of the state space for the recovery mechanism. The answer to this question thus also implicitly tells whether the recovery mechanism is fully functional in the recovery process.
- Given an integer  $k$  and the set of failure states, what is the maximal set of safe states, for which the system has a strategy to maintain  $k$ -resilience? In game theory, this means that safety resilience games can be used for synthesizing safety regions for a given bound on consecutive errors before the system is fully recovered.

The question can be extended to not only partition the states into safety, recovery, and failure states, but also for providing memoryless control on the safety and recovery states.

- Given a set of failure states, what is the maximal resilience level of the system that can be achieved with proper control? We argue that this maximal resilience level is a

well-defined and plausible indicator of the defence strength of a control mechanism against a realistic error model.

With the algorithm, software engineers and system designers can focus on maximizing the number of dense errors that the system can tolerate infinitely often, providing that they are grouped into blocks that are separated by a short period of time, which is sufficient for recovery.

We investigate how to analyse the game with existing techniques. We present an extension to alternating-time  $\mu$ -calculus (AMC) and propose to use the AMC model-checking algorithm on concurrent games to check resilience levels of embedded systems. We present reduction from safety resilience games to AMC formulas and concurrent game structures. Then we present a PTIME algorithm for answering whether the system can be controlled to tolerate up to a given number of errors in a row. The algorithm can then be used to find the maximal resilience level that can be achieved of the system. The evaluation is constructive: it provides a control strategy for the protagonist, which can be used to control a system to meet this predefined resilience level.

# Chapter 2

## Related Work

### 2.1 Existed Logics about Game and Strategy

#### 2.1.1 Prior to Strategy Logics

Kupferman, Vardi, and Wolper proposed Module checking, a famous framework for checking whether open systems can satisfy temporal logic properties [58]. In the framework, the open systems are modeled as a 2-agent turn-based games, one agent is the system and the other is the environment. The specification properties can be in CTL, CTL\*, and LTL. Since there are only two agents, non-trivial strategy collaboration does not exist and the techniques in our model-checking algorithms cannot contribute to the improvement of their algorithms.

Alur, Henzinger, and Kupferman presented ATL (alternating-time temporal logic), ATL\*, AMC (alternating-time  $\mu$ -calculus), and GL (game logic) with strategy quantifier  $\langle\langle M \rangle\rangle$  [3].

Brihaye, Lopes, Laroussinie, and Markey introduced a very expressive extension to ATL\* with other players' strategies and memory constraints [59]. They also showed that the model-checking problem of this extension is decidable.

Pinchinat introduced a way to specify expressive constraints on strategies in concurrent games by extending  $\mu$ -calculus with decision modalities [60].

Laroussinie and Markey reported decidability of satisfiability problems in this direc-

tion with new context constraints, e.g., the number of moves by agents are bound [61].

### 2.1.2 With Strategy Logics

Chatterjee, Henzinger, and Piterman introduced a strategy logic allowing for first-order quantification over strategies [5]. The decision procedure is however non-elementary.

Mogavero, Murano, Perelli, Sauro, Vardi et al also identified fragments of strategy logics which enjoys a doubly exponential time complexity model-checking algorithm [6, 62, 63]. For example, in [63], conjunction and disjunction cannot happen in the same scope of strategy quantification. BSIL is also a fragment of strategy logic with restriction on the hierarchy of SIQs and sometimes can be handy in expressing boolean relation among strategies. Moreover, the restriction on BSIL results in a much lower model-checking complexity of PSPACE. We releases the hierarchy restriction in TCL. Thus TCL can describe the specifications which requires strategy quantifiers be placed after temporal operators. And the TCL model-checking complexity is higher than BSIL, EXPTIME.

## 2.2 Fault tolerance

Traditionally, fault tolerance refers to various basic fault models [43], such as a limited number of errors [50]. These traditional fault models are subsumed by more general synthesis or control objectives [43, 51, 52]; as simple objectives with practical relevance, they have triggered the development of specialized tools [44, 45].

Dijkstra's self-stabilization criterion [53, 54] suggests to build systems that eventually recover to a 'good state', from where the program commences normally. Instead of *constructing* a system to satisfy such a goal, one might want to apply control theory to *restrict* the execution of an existing system to achieve an additional goal. Our control objective is a recovery mechanism for up to  $k$  errors. After recovery, the system has to tolerate up to  $k$  errors again, and so forth. In this work, we suggest a mechanism to synthesize a recovery mechanism for a given fault model and recovery primitives.

In [55], an interesting notion of robustness based on Hamming and Lewenstein dis-

tance related to the number of past states is defined. It establishes a connection between these distances with a notion of synchronization that characterizes the ability of the system to reset for combinatorial systems. In [56], 'ratio games' are discussed, where the objective is to minimize the ratio between failures induced by the environment and system errors caused by them.

Besides using our simple game model that neither refer directly to time, nor to probabilities, one can also consider models that make these aspects explicit. Their analysis is far more complex (with [57] offering the best complexity bounds), and so are the resulting strategies. If we, for example, return to the example of airplanes with an operation time of 20 hours referred to in Table 1.1, then an optimal timed model would take the remaining operation time into account. When the remaining time is two minutes, the balance between being resilient against waves of two errors and being resilient against 5 errors looks very different, and the optimal control would change over time rather than being static. Another implication of more complex models would be that the error model would have to be more detailed. Even if one assumes that a simple concept like safe states persists, it depends on the fineties of such a model if a two step path back to it where an error after step one leads to system failure is preferable over a much longer path, say through 10,000 intermediate states, where one error can be tolerated during recovery.

We believe that the independence from such details is an advantage of our technique, partly because it is simpler and cheaper, and partly because the further advantages one can obtain from more detailed error models rely heavily on very knowledge of (or, realistically, on very detailed assumptions on) how errors are distributed.

In [18, 19] the resilience model we have introduced [20] has been applied for synthesising robust control in an assume-guarantee setting to produce robustness against occasional noncompliance of the environment with the assumptions of its behavior.

# Chapter 3

## Back Ground

### 3.1 Concurrent Games

A *concurrent game* is played by many agents that make their moves concurrently. Such a game can be formalized with the following definition.

**Definition 1.** A concurrent game graph is a tuple  $A = \langle m, Q, r, P, \lambda, R, \Delta, \delta \rangle$  with the following restrictions.

- $m$  is the number of agents in the game.
- $Q$  is a finite set of states.
- $r \in Q$  is the *initial state* of  $A$ .
- $P$  is a finite set of atomic propositions.
- Function  $\lambda : Q \mapsto 2^P$  labels each state in  $Q$  with a set of atomic propositions.
- $R \subseteq Q \times Q$  is the set of transitions.
- $\Delta$  is a set of tokens that can be issued by the agents during transitions.
- $\delta : (R \times [1, m]) \mapsto \Delta$  is a function that specifies the token (move symbol) issued by each agent in a transition.

During a transition, each agent selects a token. If there is no transition matching all the tokens selected by the agents, then there is no transition. Otherwise, the matching transition takes place. □

In Figure 3.1, we have the graphical representation of a concurrent game graph. The




Figure 3.1: A concurrent game graph

ovals represent states while the arcs represent state transitions. We also put down the  $\lambda$  values inside the corresponding states. On each edge, we label the tokens issued by the agents. Specifically, the label on arrow  $(q, q')$  is  $[\delta((q, q'), 1), \dots, \delta((q, q'), m)]$ . For example, in Figure 3.1, on edge  $(v, u)$ , we have label  $[a, b]$  meaning that to make the transition, agent 1 has to choose token  $a$  while agent 2 has to choose  $b$ .

For convenience, in the remaining part of the manuscript, we assume that we are always in the context of a given game graph  $\mathcal{G} = \langle m, Q, r, P, \lambda, R, \Delta, \delta \rangle$ . Thus, when we write  $Q, r, P, \lambda, R, \Delta$ , and  $\delta$  we respectively refer to the corresponding components of this  $\mathcal{G}$ .

A *state predicate* of  $P$  is a Boolean combination of elements in  $P$ . The satisfaction of a state predicate  $\eta$  at a state  $q$ , in symbols  $q \models \eta$ , is defined in the standard way.

A *play* is an infinite path in a game graph. A play is *initial* if it begins with the initial state. Given a play  $\rho = \bar{q}_0 \bar{q}_1 \dots$ , for every  $k \geq 0$ , we let  $\rho(k) = \bar{q}_k$ . Also, given  $h \leq k$ , we let  $\rho[h, k]$  denote  $\rho(h) \dots \rho(k)$  and  $\rho[h, \infty)$  denote the infinite tail of  $\rho$  from  $\rho(h)$ . A *play prefix* is a finite segment of a play from the beginning of the play. Given a play prefix  $\rho = \bar{q}_0 \bar{q}_1 \dots \bar{q}_n$ , we use  $|\rho| = n + 1$  for the *length* of  $\rho$ . For convenience, we use  $\text{last}(\rho)$  to denote the last state in  $\rho$ , i.e.,  $\rho(|\rho| - 1)$ .

Let  $Q^*$  be the set of finite sequences of states in  $Q$ . For an agent  $a \in [1, m]$ , a *strategy*  $\sigma$  for  $a$  is a function from  $Q^*$  to  $\Delta$ . An *agency*  $A$  of  $[1, m]$  is an integer subset of  $[1, m]$ . For example, “ $\{1, 3, 4\}$ ” represents the agency that consists of agents 1, 3, and 4. A *strategy profile* (or *S-profile*)  $\Sigma$  of an agency  $A \subseteq [1, m]$  is a partial function from  $[1, m]$  to the set of strategies such that, for every  $a \in [1, m]$ ,  $a \in A$  iff  $\Sigma(a)$  is defined. The composition of two S-profiles  $\Sigma, \Pi$ , in symbols  $\Sigma \circ \Pi$ , is defined with the following restrictions for every

$a \in [1, m]$ .


- If  $\Pi(a)$  is defined, then  $\Sigma \circ \Pi(a) = \Pi(a)$ .
- If  $\Sigma(a)$  is defined and  $\Pi(a)$  is undefined, then  $\Sigma \circ \Pi(a) = \Sigma(a)$ .
- If  $\Sigma(a)$  and  $\Pi(a)$  are both undefined, then  $\Sigma \circ \Pi(a)$  is also undefined.

We will use composition of S-profiles to model inheritance of strategy bindings from ancestor formulas.

A play  $\rho$  is compatible with a strategy  $\sigma$  of an agent  $a \in [1, m]$  iff for every  $k \in [0, \infty)$ ,  $\delta((\rho(k), \rho(k+1)), a) = \sigma(\rho[0, k])$ . The play is compatible with an S-profile  $\Sigma$  of agency  $A$  iff for every  $a \in A$ , the play is compatible with  $\Sigma(a)$  of agent  $a$ .

## 3.2 Turn-Based Games

Another popular game structure is *turn-based game* in which at each state, at most one agent gets to decide the next state. For example, in figure 3.2, we have the graphical representation of a turn-based game graph with initial state  $v$ . The ovals and squares



○ belongs to Agent 1 and □ belongs to Agent 2.

Figure 3.2: A turn-based game graph

represent states respectively of agent 1 and agent 2. The arcs represent state transitions.

In fact, every turn-based game can be represented as a special case of concurrent games. Specifically, a turn-based game  $\mathcal{G} = \langle m, Q, r, P, \lambda, E, \Delta, \delta \rangle$  can be viewed as a concurrent game with the following restrictions.

- $\Delta = Q \cup \{\perp\}$  where  $\perp$  denotes a dummy move not in  $Q$ .
- For every  $(q, q') \in R$ , if  $q$  belongs to agent  $a$ , then  $\delta((q, q'), a) = q'$  and for every  $a' \neq a$ ,  $\delta((q, q'), a') = \perp$ .






Figure 3.3: Framework of resilience design

- For every  $(q_1, q_2), (q_3, q_4) \in R$  and agent  $a$ ,  $\delta((q_1, q_2), a) = \perp$  if and only if  $\delta((q_3, q_4), a) = \perp$ .

This restriction says that every state can be owned by only one agent.

For convenience, for a turn-based game, the owner of a state  $q$ ,  $\omega(q)$  in symbols, is defined as agent  $a$  with  $\forall(q, q') \in E(\delta((q, q'), a) = q')$ . For ease of notation, we denote with  $Q_a = \{q \in Q \mid \omega(q) = a\}$  the states owned by an agent  $a$ .

In the investigation of many research issues, turn-based game graphs are easier to handle than concurrent game graphs. So in latter sections, we sometimes use turn-based game graphs in examples and explanation of the theory as we see fit.

### 3.3 Two-player concurrent game structures

To facilitate the explanation of resilience analysis in a game's perspective, we start by defining a special case of the concurrent game,  $\mathcal{K} = \langle 2, Q, r, P, \lambda, R, E_1, E_2, \delta \rangle$ . The first difference is that there are only 2 players in the game. For convenience, we divide the token set that player can choose,  $\Delta$ , into 2 subsets  $E_1$  and  $E_2$ . Intuitively,  $E_1$  are the tokens set that can be picked by protagonist and  $E_2$  are the tokens set that can be picked by antagonist. Moreover, each member of  $E_2$  are labelled by either *error* or *non-error* moves.

### 3.4 Fault Tolerance

Resilience to errors in computer systems is usually achieved through error recovery design as illustrated in Figure 3.3. The system states can be partitioned into three regions safe, recovery, and failure. The left part of the figure represents the safety region. The

states in this region represent those for 'normal' operation. When an error occurs, the system goes through a recovery process, in which it follows some recovery mechanism. This is shown as the "recovering" area in Figure 3.3. In this region, the system tries to repair the effects caused by an error and thus to go back to the safety region.

In the recovery region, however, errors may still happen. In general, fault-tolerant systems are built under the assumption that error detection and recovery is speedy and that there can only be a few errors during the process of recovery. If the recovery mechanism is not resilient enough, a few errors may drive the system into failure.

We use following 3 examples from computer architecture, OS interrupt handler and security system to illustrate our point.

**Example 2. (Fault-tolerant computer architectures):** In computer architectures, fault-tolerance is usually achieved via hardware duplication. Consider an example of a multi-processor system that includes  $n$  processor copies and  $m$  memory copies. The  $n$  processors each can follow the instructions of the original system, or be engaged in memory recovery. When a copy of the memory fails, a processor can be assigned to recover it. Majority check can be used to detect that a processor is faulty or that memory copy is faulty (often, both would happen at the same time). For recovery, we can set a free processor to recover some memory copy, or make a processor follow the code of the majority of processors.

The key to error resilience is to decide whether to make a processor follow the execution of the majority, or to assign it to recover faulty memory. If too many errors occur in a short while before the errors can be recovered from, then there may be no more processors left to carry out any more recovery. When such a critical situation arises, the system enters failure state when another error is induced.

The recovery mechanism described above is typical in the design of fault-tolerant systems [17]. As explained, a practical recovery mechanism usually does not rely on the detailed structure of the system. Instead, error-detection techniques such as parity checks, voting (for majority checks), etc., are usually employed. In fact, the number of duplicates is usually critical to the resilience of the system to errors. As long as the majority of the duplicate modules can be recovered in time (i.e., before the next wave of errors), resilience

of the system can be achieved. □

**Example 3. (Exception handling):** At the operating system level, errors are usually signalled via interrupt lines and handled with routines called handlers. The first thing that needs to be done by a handler is to save the CPU state of the interrupted process. In some operating systems, a static memory space is used for this purpose for each handler. In such a scheme, if the same error happens again while executing the error handler, then the system can run into the risk that the CPU states of the interrupted handler can be overwritten and destroyed.

Another scheme is to use a stack to save the CPU states of the interrupted processes. Such a scheme seems resilient to errors that happen during the execution of error handlers. Still, too many errors that happen during the execution of error handlers can deny critical functions of the system and incur failures, including missed timer updates and priority inversions. Thus, a proper assumption on the timely error recovery by the error handling routines is critical to the design of error resilience in such cases. □

**Example 4. (Security attacks):** Security in the Internet also relies on resilience to attacks of hackers, viruses, malware, etc. For example, one common technique of attacks to communication modules is to overflow the communication buffers. In such attacks, the sizes of the buffers and the ability of the security procedures to detect and recover from such overflowing attacks is crucial to the resilience design. □

These examples show that recovery is a crucial concept for designing systems that are resilient to errors. When system errors are detected in such a system, the system activates a recovery mechanism to remove the effect of the errors. When designing such systems, the system designers usually know what errors and failures the systems can expect, according to the specification. To avoid failures in the occurrence of dense errors, the system designers usually incorporate many error recovery mechanism in the system, e.g., exception handlers and hardware/software redundancy. But, in general, it would be difficult for the designers to evaluate how effective their recovery mechanism is to dense errors. To overcome this difficulty, we believe that it is important to support them with automated analytical tools with a solid foundation.

Resilience has also been used in [18, 19] with a similar goal. When synthesising code, one relies on assumptions of the behavior of the environment, and the formal specification would only ask for the provision of guarantees under the condition that the assumptions are satisfied. When evaluating the quality of an implementation, the behavior in cases where the environment does not comply with the assumption matters. In [18, 19], the resilience model we have introduced in the conference version [20] of this paper has been followed up upon, and proven to be well suited for reactive synthesis.

In this work, we use these observations to design a theoretical framework for synthesizing a control mechanism that provides the maximal resilience against software errors in a realistic error model.

# Chapter 4

## Basic Strategy-interaction Logic

### 4.1 Running Example

We use a banking example to explain the idea of this work. Suppose that a bank want to provide the following services.

- No client is allowed to check other's password
- Depositing to an account by a client from the root screen.
- Transferring money from an account in the bank to an account in another bank, also from the root screen.
- There is only one bank system, which means the strategy of the bank system which fulfilled above specification should be identical.

As we can see, the deposit service involves the interaction between a client and the bank while the transfer service involves at least three parties: a client, the bank, and a partner bank. Also in the meantime, the bank wants to forbid any client and partner from checking sensitive information of other clients, .e.g., checking the password of another client.

#### 4.1.1 Trying to Write Down a Correct Formal Specification

To develop the services, the bank manager in charge of the project need specify the services and make sure that the specification is not erroneous. If they turn to the literature, the bank manager may choose ATL or its extensions, e.g., fair ATL, ATL\*, AMC, GL, etc.,

to specify the services. The choice is plausible at first glance since ATL and fair ATL both have a polynomial time model-checking algorithm and could support the verification of the services. So the manager could easily specify the services with the following formula.

$$\langle 1 \rangle \left( \begin{array}{l} \Box \neg \text{checkOthersPassword} \\ \wedge \langle 2 \rangle \Diamond \text{depositDone} \\ \wedge \langle 2, 3 \rangle \Diamond \text{transferDone} \end{array} \right) \quad (C)$$

Here we use agent index 1 for the bank, 2 for the client, and 3 for the partner bank. But on a second look, we can see that this specification is too restrictive. Subformula  $\langle 2 \rangle \Diamond \text{depositDone}$  says that the client can force a deposit transaction no matter how the banking system responds. In practice, there are many factors beyond the control of the client for the success of a transaction. For example, when the banking system is in maintenance or out of order, then the transaction may fail.

After realizing the problem with formula (C), the manager decides to rewrite the specification as follows.

$$\begin{array}{l} \langle 1 \rangle \Box \neg \text{checkOthersPassword} \\ \wedge \langle 1, 2 \rangle \Diamond \text{depositDone} \\ \wedge \langle 1, 2, 3 \rangle \Diamond \text{transferDone} \end{array} \quad (D)$$

This specification fixes the issue observed in formula (C) since now  $\langle 1, 2 \rangle \Diamond \text{depositDone}$  says with the collaboration of the banking system, the client can finish a deposit transaction. However, there is a subtle issue which nullifies the formula. In the collaboration between the banking system and the client for property  $\langle 1, 2 \rangle \Diamond \text{depositDone}$ , the banking system is no longer required to maintain property  $\Box \neg \text{checkOthersPassword}$  since according to the semantics of ATL, the specification does not require the banking system to use the same strategy to enforce both  $\langle 1 \rangle \Box \neg \text{checkOthersPassword}$  and  $\langle 1, 2 \rangle \Diamond \text{depositDone}$ . The same issue also appears in  $\langle 1, 2, 3 \rangle \Diamond \text{transferDone}$ . That is, the banking system's strategy in enforcing  $\langle 1, 2, 3 \rangle \Diamond \text{transferDone}$  is not required to be the same one in enforcing  $\langle 1 \rangle \Box \neg \text{checkOthersPassword}$ . Thus, the bank cannot sue their contractors for developing the service system if the systems leaks clients' passwords in transferring funds.

### 4.1.2 Resorting to General Strategy Logics for a Correct Specification

After several trials, the manager eventually finds out that the specification cannot be expressed in ATL, ATL\*, GL, and AMC [3]. Then the manager turns to strategy logics [5, 6] in the literature and finds out that the specification can be expressed as follows.

$$\langle x \rangle \langle y \rangle \langle z \rangle \langle w \rangle \left( \begin{array}{l} (1, x) \Box \neg \text{checkOthersPassword} \\ \wedge (1, x)(2, y) \Diamond \text{depositDone} \\ \wedge (1, x)(2, z)(3, w) \Diamond \text{transferDone} \end{array} \right) \quad (\text{E})$$

This formula declares four existentially quantified strategies:  $x$ ,  $y$ ,  $z$ , and  $w$ , and says the following.

- On using strategy  $x$ , the banking system can enforce  $\Box \neg \text{checkOthersPassword}$ .
- On using strategy  $x$  and  $y$  respectively, the banking system and a client can finish the deposit transaction.
- On using strategy  $x$ ,  $z$ , and  $w$  respectively, the banking system, the client, and a partner bank can also finish a fund transfer transaction.

The manager is happy with the elegant specification and wants to verify the specification on the model of the development delivery. Then he finds out that there is no tractable algorithm and working tools for checking this specification or synthesize the strategies. In fact, the complexity reported in the literature is at best doubly exponential [5, 6]. So the manager finds himself in the dilemma between expressiveness and verification efficiency.

But when the manager bangs his head to the theoretical challenges, he soon realizes that strategy logics offer more expressiveness than he expects. For example, he could write down the following formula.

$$\langle x \rangle \langle y \rangle [z] \langle w \rangle \left( \begin{array}{l} (1, x) \Box \neg \text{checkOthersPassword} \\ \wedge (1, x) \Box (2, y) \Diamond \text{depositDone} \\ \wedge (1, x) \Box (2, z)(3, w) \Diamond \text{transferDone} \end{array} \right) \quad (\text{F})$$

Here  $[z]$  universally quantifies a strategy named  $z$ . Thus subformula  $(1, x)(2, z)(3, w) \Diamond \text{transferDone}$  in fact says that the banking system has a strategy  $x$  such that for all strategies  $z$  of the

client, the partner bank have a strategy  $w$  to help enforcing  $\Diamond transferDone$ . Considering that a strategy is an action decision function on the history, this in fact means that strategy  $w$  is clairvoyant and contradicts the implementation assumption of all practical banking systems. Indeed the huge verification complexity of strategy logics can be attributed to such free-style binding operations of strategy names to agents. Thus, the manager asks himself whether there is a subclass of strategy logics that allows for elegant expressions of natural and practical specifications while supporting verification with less complexity. Similar examples in fact can appear in different projects with services that rely on collaboration of multiple agents.

#### 4.1.3 BSIL: New Strategy Modalities Expressive Enough for the Specification

In fact, in conceivable applications, the developers cannot implement and are not interested at clairvoyant strategies. In other words, for practical specifications, it is likely that existential SIQs are expressive enough. In BSIL, formula (E) can be rewritten as follows.

$$\langle 1 \rangle \left( \begin{array}{l} \Box \neg checkOthersPassword \\ \wedge \langle +2 \rangle \Diamond depositDone \\ \wedge \langle +2, 3 \rangle \Diamond transferDone \end{array} \right) \quad (G)$$

This formula says that the banking system has a strategy that at any instant, the system can ensure that no password is leaked, the system allows a client to finish a deposit transaction, and the system allows a client to transfer funds from or to a partner bank.

#### 4.1.4 Symbolic Strategy Names and Path Obligations

But to verify such a BSIL property, we need new techniques to check that along plays, the same decision is made for the same strategy. Consider the model of the banking system in figure 4.1. We need to find strategies of the agents to fulfill  $\Box \neg checkOthersPassword$ ,  $\Diamond depositDone$ , and  $\Diamond transferDone$ . To assist in our explanation of our verification algorithms, we use symbolic strategy names  $x, y, z, w$  and the binding notations in formula (F).



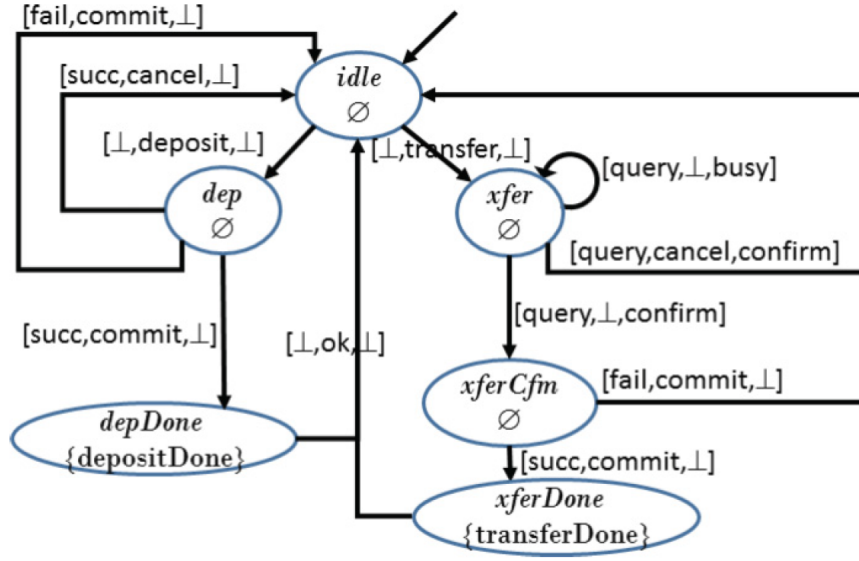


Figure 4.1: Game graph of a banking system

In fact, formula (G) is exactly the same as formula (E). Thus by interpreting all symbolic strategy names as existential quantified, formula (G) can be rewritten as follows.

$$\begin{aligned}
 & ((1, x) \Box \neg \text{checkOthersPassword}) \\
 & \wedge ((1, x)(2, y) \Diamond \text{depositDone}) \\
 & \wedge ((1, x)(2, z)(3, w) \Diamond \text{transferDone})
 \end{aligned} \tag{H}$$

Note that since we only allow explicit existential strategy quantification, the number of strategy bindings are exactly determined by the number and sizes of SQs and SIQs in the formula.

If we label  $(1, x) \Box \neg \text{checkOthersPassword}$ ,  $(1, x)(2, y) \Diamond \text{depositDone}$ , and  $(1, x)(2, z)(3, w) \Diamond \text{transferDone}$  on nodes of a computation tree in a labelling procedure, then the symbolic names precisely reflect the constraints on the decisions along plays. For example, suppose that both  $(1, x) \Box \neg \text{checkOthersPassword}$  and  $(1, x)(2, y) \Diamond \text{depositDone}$  are labelled on a state of location *dep*. This state has three successor states: for convenience C1 of location *idle* via  $[\text{fail}, \text{commit}, \perp]$ , C2 of location *idle* via  $[\text{succ}, \text{cancel}, \perp]$ , and C3 of location *depDone* via  $[\text{succ}, \text{commit}, \perp]$ . If strategy  $x$  chooses action *succ*, then obligation  $(1, x) \Box \neg \text{checkOthersPassword}$  will be passed down to C2 and C3. In this situation, whether obligation  $(1, x)(2, y) \Diamond \text{depositDone}$  is labeled on C2 or C3 relies on the action decision of strategy  $y$ . But the constraint on strategy decision is that

if  $(1, x)(2, y) \diamond \text{depositDone}$  is labeled on C2, then  $(1, x) \Box \neg \text{checkOthersPassword}$  must also be labeled on C2; and if  $(1, x)(2, y) \diamond \text{depositDone}$  is labeled on C3, then  $(1, x) \Box \neg \text{checkOthersPassword}$  must also be labeled on C3. The reason is that the path obligations are enforced by the same strategy of the bank.

#### 4.1.5 Passing Down the Path Obligations While Observing the Restrictions Among S-Profiles

The above observation points out how to judge whether a passing-down scheme of path obligations from a parent state to its child states are consistent with the strategy quantifications. For example, suppose that we have the following characteristics of the strategy profile  $(1, x)(2, y)(2, z)(3, w)$ .

- Strategy  $x$  of the banking system never issues action fail.
- Strategy  $y$  and  $z$  never cancel transactions for its client.
- Strategy  $w$  never shows busy message for the partner bank.

Then the verification problem of strategy interaction can be visualized as finding strategies  $x, y, z$ , and  $w$  that selecting paths in a computation tree to satisfy three path obligations:  $\Box \neg \text{checkOthersPassword}$ ,  $\diamond \text{depositDone}$ , and  $\diamond \text{transferDone}$ . Figure 4.2 shows part of a computation tree when strategy  $x$  for the banking system is fixed with the following restriction.

- At location *idle*, *depDone*, and *xferDone*, always issue action  $\perp$ .
- At location *dep* and *xferCfm*, only issue action succ.

The leaves of the tree all end at location *idle* at which the same tree structure can be replicated. We can start synthesizing strategy  $y, z$ , and  $w$  in the context of this strategy  $x$  which, for convenience of explanation, is memoryless (or positional).

As we have suggested, we can label the path obligations on the nodes of the computation tree and examine whether a strategy profile can fulfill all the path obligations. In the literature, we can name the nodes by their branching paths. For example, the root is named  $\varepsilon$  which is the empty string. Then the node at location *xferDone* is named '100' and the rightmost *idle* node is named '12.' We begin by label node ' $\varepsilon$ ' with the three path obli-




Figure 4.2: Part of the computation tree of the banking system

gations:  $(1, x) \Box \neg \text{checkOthersPassword}$ ,  $(1, x)(2, y) \Diamond \text{depositDone}$ , and  $(1, x)(2, z)(3, w) \Diamond \text{transferDone}$ . Then we can pass the path obligations to the children in the following way.

- We must pass  $(1, x) \Box \neg \text{checkOthersPassword}$  to both child '0' and '1' of the root since the two children are both selected by action  $\perp$  of the banking system.
- We can pass  $(1, x)(2, y) \Diamond \text{depositDone}$  and  $(1, x)(2, z)(3, w) \Diamond \text{transferDone}$  respectively to child '0' and '1' of the root since the two children are selected with the same strategy  $x$  of the banking system and different strategies ( $y$  and  $z$ ) of the client.

Note that we can check whether the passing down of the path obligation is consistent with the strategy quantifications in the input formula (G) just by checking the labelling of path obligations with strategy name bindings since all consistency information are maintained there.

Similarly, from node '0', strategy  $x$  and  $y$  can together choose to pass  $(1, x)(2, y) \Diamond \text{depositDone}$  to node '00' while strategy  $x$  must unilaterally pass  $(1, x) \Box \neg \text{checkOthersPassword}$  to both '00' and '01.' Then at node '00,'  $(1, x)(2, y) \Diamond \text{depositDone}$  is fulfilled eventually and no longer need be passed down. The passing down of path obligations from node '1' and fulfillment of  $(1, x)(2, z)(3, w) \Diamond \text{transferDone}$  can then be reasoned similarly.

#### 4.1.6 Finding Finite Satisfying Evidence for a Formula in a Computation Tree

As we can see, the existence of the tree top in figure 4.2 is sufficient for synthesizing strategy profile  $(1, x)(2, y)(2, z)(3, w)$  to satisfy formula  $(G)$ . From this example, it is easy to see that the requirement for such a tree top. First, all eventual-formulas (or until-formulas) from the root SQ to the next level of SQs have to be fulfilled by the strategy profile. Second, the leaves of the tree top do not contain any eventual-formula (or until-formula) labels. Thus our PSPACE algorithm for model-checking BSIL formulas is actually a non-deterministic one that guess and check the consistency of the tree top and strategic actions at the nodes in the tree top.

Then the final question regarding our algorithms is the terminating condition for searching the tree tops and strategy profiles. Let us consider a tree node labelled with path obligations and how to (non-deterministically) explore for a tree top and a strategy profile. Note that the path obligations labelled on a child will be no more than those of its parent. In fact, the path obligations can stop being passed down in a path when it is fulfilled. Path obligations are simply passed down and not generated along the paths. As a result, along a path longer than the size of the game graph, either the number of path obligations decreases or two nodes, say  $v$  and  $v'$  (for convenience, we assume  $v$  is an ancestor of  $v'$ ), with identical location and labels must appear. If the tree top is an evidence of existence of a strategy profile and the latter happens, then we can replace the subtree rooted at  $v$  with the one rooted at  $v'$  and the result tree top is still an evidence. This observation implies that we can focus on tree top such that along any path, no duplication of nodes with the same location and the same path obligation labels. This implies that we only have to explore tree tops of depth no greater than the size of the game graph times the number of temporal modalities in the given BSIL formula.

## 4.2 BSIL

### 4.2.1 Syntax

For concurrent game graph  $\mathcal{G}$  of  $m$  agents, we have three types of formulas: *state formulas*, *tree formulas*, and *path formulas*. State formulas describe properties of states. Tree formulas describe interaction of strategies. Path formulas describe properties of plays. BSIL formulas are constructed with the following three syntax rules, respectively, for state formulas  $\phi$ , tree formulas  $\tau$ , and path formulas  $\theta$ .

$$\begin{aligned}\phi &::= p \mid \neg\phi_1 \mid \phi_1 \vee \phi_2 \mid \langle A \rangle \tau \mid \langle A \rangle \theta \\ \tau &::= \tau_1 \vee \tau_2 \mid \tau_1 \wedge \tau_2 \mid \langle +A \rangle \tau_1 \mid \langle +A \rangle \theta \\ \theta &::= \neg\theta_1 \mid \theta_1 \vee \theta_2 \mid \bigcirc \phi_1 \mid \phi_1 \text{U} \phi_2\end{aligned}$$

Here,  $p$  is an atomic proposition in  $P$  and  $A$  is an agency of  $[1, m]$ .  $\langle A \rangle$  is a *strategy quantifier* (SQ) and  $\langle +A \rangle$  is a *strategy interaction quantifier* (SIQ).  $\langle A \rangle \psi$  means that there exists an S-profile for the agency  $A$  that makes all plays satisfy  $\psi$ . Formulas of the form  $\langle +B \rangle \psi_1$  must happen within an SQ. Intuitively, they mean that there exists an S-profile of  $B$  that collaborates with the strategies declared with ancestor formulas to make  $\psi_1$  true. For convenience, we view SQs as special cases of SIQs. Also, for conciseness, we omit null SIQs  $\langle + \rangle$ .

State formulas  $\phi$  are called *BSIL formulas*. From now on, we assume that we are always in the context of a given BSIL formula  $\chi$ . Note that we strictly require that strategy interaction cannot cross path modal operators. This restriction is important and allows us to analyze the interaction of strategies locally in a state and then enforce the interaction along all paths from this state.

For convenience, we also use the common shorthands.

$$\begin{aligned}\text{true} &\equiv p \vee (\neg p) & \text{false} &\equiv \neg \text{true} & \psi_1 \Rightarrow \psi_2 &\equiv (\neg \psi_1) \vee \psi_2 \\ \Diamond \phi_1 &\equiv \text{true} \text{U} \phi_1 & \Box \phi_1 &\equiv \neg \Diamond \neg \phi_1\end{aligned}$$

The SQs and SIQs introduced above are all existential in that they are satisfied with one S-profile. Note that there is no universal SQs and SIQs in BSIL. This is purely for the

complexity of the model-checking algorithm (and problem) that we are going to present later. Thus BSIL can be used to specify the different ways of combining S-profiles to enforce system policy.

For ease of notations, we may abbreviate  $\langle \{a_1, \dots, a_n\} \rangle$  and  $\langle +\{a_1, \dots, a_n\} \rangle$  as  $\langle a_1, \dots, a_n \rangle$  and  $\langle +a_1, \dots, a_n \rangle$ , respectively.

### 4.2.2 semantics

BSIL subformulas are interpreted with respect to S-profiles. A state or a tree formula  $\phi$  is satisfied at a state  $q$  with S-profile  $\Sigma$ , denoted  $\mathcal{G}, q \models_{\Sigma} \phi$ , if, and only if, the following inductive constraints are satisfied.

- $\mathcal{G}, q \models_{\Sigma} p$  iff  $p \in \lambda(q)$ .
- For state formula  $\phi_1$ ,  $\mathcal{G}, q \models_{\Sigma} \neg\phi_1$  iff  $\mathcal{G}, q \models_{\Sigma} \phi_1$  is false.
- For state or tree formulas  $\psi_1$  and  $\psi_2$ ,  $\mathcal{G}, q \models_{\Sigma} \psi_1 \wedge \psi_2$  iff  $\mathcal{G}, q \models_{\Sigma} \psi_1$  and  $\mathcal{G}, q \models_{\Sigma} \psi_2$ .
- For state or tree formulas  $\psi_1$  and  $\psi_2$ ,  $\mathcal{G}, q \models_{\Sigma} \psi_1 \vee \psi_2$  iff either  $\mathcal{G}, q \models_{\Sigma} \psi_1$  or  $\mathcal{G}, q \models_{\Sigma} \psi_2$ .
- $\mathcal{G}, q \models_{\Sigma} \langle A \rangle \tau$  iff there exists an S-profile  $\Pi$  of  $A$  with  $\mathcal{G}, q \models_{\Pi} \tau$ .
- $\mathcal{G}, q \models_{\Sigma} \langle +A \rangle \tau$  iff there exists an S-profile  $\Pi$  of  $A$  with  $\mathcal{G}, q \models_{\Sigma \circ \Pi} \tau$ . Here, the composition  $\Sigma \circ \Pi$  of the S-profiles  $\Sigma$  and  $\Pi$  models the inheritance of strategy bindings,  $\Sigma$ , from ancestor formulas.
- $\mathcal{G}, q \models_{\Sigma} \langle A \rangle \theta$  iff there exists an S-profile  $\Pi$  of  $A$  such that, for all plays  $\rho$  from  $q$  compatible with  $\Pi$ ,  $\rho \models_{\Pi} \theta$  holds. Intuitively, this means that  $\rho$  satisfies path formula  $\theta$  with S-profile  $\Pi$ .
- $\mathcal{G}, q \models_{\Sigma} \langle +A \rangle \theta$  iff there exists an S-profile  $\Pi$  of  $A$  such that, for all plays  $\rho$  from  $q$  compatible with  $\Sigma \circ \Pi$ ,  $\rho \models_{\Sigma \circ \Pi} \theta$  holds.

A play  $\rho$  satisfies a path formula  $\theta$  with S-profile  $\Sigma$ , in symbols  $\rho \models_{\Sigma} \theta$ , iff the following restrictions hold.

- For a path formula  $\theta_1$ ,  $\rho \models_{\Sigma} \neg\theta_1$  iff it is not the case that  $\rho \models_{\Sigma} \theta_1$ .
- For path formulas  $\theta_1$  and  $\theta_2$ ,  $\rho \models_{\Sigma} \theta_1 \vee \theta_2$  iff either  $\rho \models_{\Sigma} \theta_1$  or  $\rho \models_{\Sigma} \theta_2$ .
- $\rho \models_{\Sigma} \bigcirc \psi_1$  iff  $\mathcal{G}, \rho(1) \models_{\Sigma} \psi_1$ .




Figure 4.3: A simple turn-based game that demands memoryful strategies

- $\rho \models_{\Sigma} \psi_1 U \psi_2$  iff there exists an  $h \geq 0$  with  $\mathcal{G}, \rho(h) \models_{\Sigma} \psi_2$  and for all  $j \in [0, h)$ ,  $\mathcal{G}, \rho(j) \models_{\Sigma} \psi_1$ .

For convenience, we let  $\perp$  be a null S-profile, i.e., a function that is undefined on everything. If  $\phi_1$  is a BSIL (state) formula and  $\mathcal{G}, q \models_{\perp} \phi_1$ , then we may simply write  $\mathcal{G}, q \models \phi_1$ . If  $\mathcal{G}, r \models \phi_1$ , then we also write  $\mathcal{G} \models \phi_1$ .

### 4.2.3 ATL+

ATL<sup>+</sup> is the syntactic fragment of BSIL given by the following grammar

$$\begin{aligned} \phi &::= p \mid \neg\phi_1 \mid \phi_1 \vee \phi_2 \mid \langle A \rangle \theta \\ \theta &::= \neg\theta_1 \mid \theta_1 \vee \theta_2 \mid \bigcirc \phi_1 \mid \phi_1 U \phi_2 \end{aligned}$$

ATL<sup>+</sup> can also be viewed as an extension of ATL [3] that, similar to ...’s extension of CTL to CTL<sup>+</sup> [11, 13], allows for the Boolean combination of path formulas. As such, all complexities of ATL<sup>+</sup> must reside between those of BSIL and ATL as well as between those of BSIL and CTL<sup>+</sup>, which we will use to establish the lower bounds for ATL<sup>+</sup> and BSIL.

### 4.2.4 Memory

In this subsection, we show a simple example in which the agents need memory to achieve their objective for ATL<sup>+</sup> specifications. This is exemplified by the simplest possible case: the turn based game in Figure 4.3 with one agent, two states, one atomic proposition, and two memoryless strategies that do not count the unreachable states in the histories. For the ATL<sup>+</sup> sentence  $\langle 1 \rangle ((\neg \bigcirc p) \wedge \Diamond p)$ , apparently agent 1 needs memory to enforce it.

**Lemma 5.** The strategies of the agents in ATL<sup>+</sup> (and thus in BSIL) specifications need memory. This even holds for the single agent case. □

### 4.3 Expressive Power of BSIL

In this section, we establish that BSIL is incomparable with  $ATL^*$ , AMC, and GL [3] in expressiveness. In fact, we shall first establish the incomparability between BSIL with GL and AMC. Then the incomparability with BSIL follows since GL and AMC are super-classes of  $ATL^*$  [3].

#### 4.3.1 Comparison with GL

GL separates strategy quantifications from path quantifications. In comparison,  $ATL^*$  and BSIL combines these two quantifications into SQs and SIQs. Thus with GL, we can specify that, for all S-profiles of  $A$ , there exists a play satisfying  $\psi_1$ . The (existential) strategy quantification for agency  $A$  of GL is of the form  $\exists A.\psi_1$ . The path quantifications are just ordinary CTL modalities:  $\forall\Box$ ,  $\forall U$ ,  $\exists\Box$ , and  $\exists U$ .

The following two lemmas show the relation between GL and BSIL. Lemma 6 uses two inductive families of game graphs to show that GL is not as expressive as BSIL. The base cases,  $G_1$  and  $H_1$ , are in figure 4.4 for three agents. The inductive cases,  $G_{k+1}$  and  $H_{k+1}$ , are respectively constructed out of  $G_k$  and  $H_k$  as in figure 4.5.


**Lemma 6.** Every GL formula  $\phi$  with  $k$  ( $k > 0$ ) SQs cannot distinguish  $G_k$  and  $H_k$  while  $\langle 1 \rangle((\langle +2 \rangle \Box p) \wedge \langle +2 \rangle \Box q)$  can.

**Proof :** It is clear that  $\langle 1 \rangle((\langle +2 \rangle \Box p) \wedge (\langle +2 \rangle \Box q))$  can  $G_k$  and  $H_k$  no matter what value  $k$  is. The proof continues by an induction  $k$ , the number of SQs in  $\phi$ .


**Base case:** Assume that  $\phi$  has only one modal operator. Then there are the following case analysis of GL formulas.

- **Case 1:**  $\phi$  is  $\exists A.\phi_1$  where  $\phi_1$  is a Boolean combination of formulas of the form  $\exists\psi$  or  $\forall\psi$  where  $\psi$  is a path formula. Note that  $\phi_1$  characterizes a set of states that either start a play satisfying  $\psi$  or start only plays satisfying  $\psi$ . The following case analysis shows that, for every trace subset  $S$  and every agency  $A$ , there exists a strategy of  $A$  to characterize  $S$  in figure 4.4(a) iff there exists a strategy for  $A$  to characterize  $S$  in figure 4.4(b).





(a)  $G_1$ , a game graph for base case.




(b)  $H_1$ , another game graph for base case.

○ belongs to Agent 1; □ belongs to Agent 2; and △ belongs to Agent 3.

Figure 4.4: Base cases for the expressiveness of BSIL over  $ATL^*$

- **Case 1a:**  $\phi$  is  $\exists\emptyset.\phi_1$ . In this case, there is no strategy and the sets of traces imposed by no strategy in the two state graphs are both  $\{\square p, \square q\}$ . Thus,  $\exists\emptyset.\phi_1$  cannot distinguish the trace sets of the two state graphs.
- **Case 1b:**  $\phi$  is  $\exists\{1\}.\phi_1$ . From figure 4.4, no matter what strategies agency  $\{1\}$  may choose, the trace sets for the two game graphs are both  $\{\square p, \square q\}$ . Thus,  $\exists\{1\}.\phi_1$  cannot distinguish the trace sets of the two state graphs.
- **Case 1c:**  $\phi$  is  $\exists\{2\}.\phi_1$ . This case is similar to case 1b.
- **Case 1d:**  $\phi$  is  $\exists\{3\}.\phi_1$ . This case is similar to case 1b.
- **Case 1e:**  $\phi$  is  $\exists\{1, 2\}.\phi_1$ . Agency  $\{1, 2\}$  can cooperate to force three trace sets:  $\{\square p\}$ ,  $\{\square q\}$ ,  $\{\square p, \square q\}$  in both of the game graphs in figure 4.4. Thus,  $\exists\{1, 2\}.\phi_1$  cannot distinguish the trace sets of the two state graphs.
- **Case 1f:**  $\phi$  is  $\exists\{1, 3\}.\phi_1$ . This case is similar to case 1e.



○ belongs to Agent 1; □ belongs to Agent 2; and △ belongs to Agent 3.

Figure 4.5: Inductive cases for the expressiveness of BSIL over  $ATL^*$

- **Case 1g:**  $\phi$  is  $\exists\{2, 3\}.\phi_1$ . This case is similar to case 1e.
- **Case 1h:**  $\phi$  is  $\exists\{1, 2, 3\}.\phi_1$ . Agency  $\{1, 2, 3\}$  can cooperate to force three trace sets:  $\{\Box p\}$ ,  $\{\Box q\}$ ,  $\{\Box p, \Box q\}$  in both of the game graphs in figure 4.4. Thus,  $\exists\{1, 2, 3\}.\phi_1$  cannot distinguish the trace sets of the two state graphs.
- **Case 2:**  $\phi$  is a Boolean combination of formulas in case 1. Since case 1 does not distinguish the two game graphs, this case can neither do it.

Thus the base case is proven.

**Induction step:** If  $\phi$  has  $k$  modal operators, to tell the difference between  $G_k$  and  $H_k$ , we need a modal subformula of  $\phi$  that can tell the difference of  $G_{k-1}$  and  $H_{k-1}$ . But according to the inductive hypothesis, this is impossible. Thus the lemma is proven.  $\square$

**Lemma 7.** GL formula  $\exists\{1\}.((\exists\Box p) \wedge \exists\Box q)$  is not equivalent to any BSIL formula.

**Proof :** The proof basically follows the same argument in [3] that  $\exists\{1\}.((\exists\Box p) \wedge \exists\Box q)$  is not equivalent to any  $ATL^*$  formula.  $\square$

Lemmas 6 and 7 together show that GL and BSIL are not comparable in expressiveness.

### 4.3.2 Comparison with AMC

AMC is an extension from  $\mu$ -calculus and allows for multiple fixpoints interleaved together [3]. An AMC formula contains fixpoint operators on state set variables. The only modality of AMC is of the form  $\langle A \rangle \bigcirc \psi$  and least fixpoint **lfp** $x.\psi_1(x)$  where  $\psi_1(x)$  is a Boolean function of atomic propositions and state set variables (including  $x$ ). It is required that every occurrence of  $x$  in  $\psi_1$  is under even number of negations. The duality of the least fixpoint operator is the greatest fixpoint operator **gfp**. Formula **gfp** $x.\psi_1$  is defined as  $\neg \mathbf{lfp}x.\neg\psi_1(x)$ .

To establish that AMC is not as expressive as BSIL, we basically follow the proof style for Lemma 13 and use the same two families of game graphs. The statement of the lemma requires notations for state set variables and other details in AMC. We need to define the domain of values for the free state set variables in AMC formulas. Let  $X$  be the set of state set variables. Without loss of generality, we assume that no two subformulas of the form **lfp** $x.\phi$  in a given AMC formula share the same quantified name of  $x$ . Given a subformula **lfp** $x_i.\phi$  with free variables  $x_1, \dots, x_n$  in  $\phi$  and no modal operator  $\langle \dots \rangle \bigcirc$  in  $\phi$ , we define  $\phi$  as a *base template* for  $x_i$ . Then we can define the *base formula domain* of  $x_i$ , denoted  $F_0(x_i)$ , as the smallest set with the following restrictions. We let  $\phi[x_1 \mapsto \eta_1, \dots, x_n \mapsto \eta_n]$  be the AMC formula identical to  $\phi$  except that every occurrence of  $x_i$  in  $\phi$  are respectively replaced with  $\eta_i$ .

- For each  $x_i \in X$  with base template  $\phi$ ,  $\phi[x_1 \mapsto \text{false}, \dots, x_n \mapsto \text{false}] \in F_0(x_i)$ .
- For each  $x_i \in X$  with base template  $\phi$  and  $\phi_1 \in F_0(x_1), \dots, \phi_n \in F_0(x_n)$ ,  $\phi[x_1 \mapsto \phi_1, \dots, x_n \mapsto \phi_n] \in F_0(x_i)$ .

Note that there is neither variables nor ‘**lfp**’ operators in  $F_0(x)$  for every  $x$ . Thus we can define the characterization  $\kappa$  of a formula  $\phi_1$  in  $F_0(x)$  for  $\mathcal{G}$ , in symbols  $\kappa(\mathcal{G}, \phi_1)$ , as follows.

- For each atomic proposition  $p \in P$ ,  $\kappa(\mathcal{G}, p) = \{q \mid p \in \lambda(q)\}$ .
- For each  $\phi \in F_0(x)$ ,  $\kappa(\mathcal{G}, \neg\phi) = Q - \kappa(\mathcal{G}, \phi)$ .
- For each  $\phi_1, \phi_2 \in F_0(x)$ ,  $\kappa(\mathcal{G}, \phi_1 \vee \phi_2) = \kappa(\mathcal{G}, \phi_1) \cup \kappa(\mathcal{G}, \phi_2)$ .

A *valuation*  $\nu$  of variables in  $X$  for  $\mathcal{G}$  is a mapping from  $X$  such that, for each  $x \in X$ ,

there exists a base domain formula  $\phi$  of  $x$  such that  $\nu(x) = \kappa(\mathcal{G}, \phi)$ .

The expressiveness comparison between AMC and BSIL relies on the following lemma.

**Lemma 8.** For every AMC formula  $\phi$  without modal operator of the form  $\langle \dots \rangle \bigcirc$ , state set variable  $x$ , and a formula  $\phi_1$  in  $F_0(x)$ ,  $r_0 \in \kappa(G_0, \phi_1)$  iff  $s_0 \in \kappa(H_0, \phi_1)$  in figure 4.4.

**Proof :** We can prove this with an structural induction on  $\phi_1$ . The base case is straightforward. The inductive step follows since Boolean combinations of subformulas that cannot distinguish  $G_0$  and  $H_0$  can neither distinguish the two game graphs.  $\square$

Now we want to classify AMC formulas according to the nesting depths of operators of the form  $\langle \dots \rangle \bigcirc$  in a formula. Specifically, we let  $\text{AMC}^{(k)}$  be the set of AMC formulas with exactly nesting depth  $k$  of operator  $\langle \dots \rangle \bigcirc$ . For example,  $\text{Ifp}x.\langle 1 \rangle \bigcirc (p \rightarrow \text{Ifp}y.\langle 2 \rangle (x \vee y \wedge \bigcirc q))$  is in  $\text{AMC}^{(2)}$ . Then,  $\text{AMC}^{(0)}$  is the smallest set with the following restrictions.

- **Case 1a:** For each atomic proposition  $p \in P$ ,  $p \in \text{AMC}^{(0)}$ .
- **Case 1b:** For each proposition variable  $x \in X$ ,  $x \in \text{AMC}^{(0)}$ .
- **Case 1c:** For each  $\phi \in \text{AMC}^{(0)}$ ,  $\neg\phi \in \text{AMC}^{(0)}$ .
- **Case 1d:** For each  $\phi_1, \phi_2 \in \text{AMC}^{(0)}$ ,  $\phi_1 \vee \phi_2 \in \text{AMC}^{(0)}$ .
- **Case 1e:** For each  $\phi \in \text{AMC}^{(0)}$  and proposition variable  $x \in X$ ,  $\text{Ifp}x.\phi \in \text{AMC}^{(0)}$ .

Then  $\text{AMC}^{(k)}$ ,  $k > 0$ , is the smallest set with the following restrictions.

- **Case 2a:** For each  $A \subseteq [1, m]$  and  $\phi \in \text{AMC}^{(k-1)}$ ,  $\langle A \rangle \bigcirc \phi \in \text{AMC}^{(k)}$ .
- **Case 2b:** For each  $\phi \in \text{AMC}^{(k)}$ ,  $\neg\phi \in \text{AMC}^{(k)}$ .
- **Case 2c:** For each  $\phi_1 \in \text{AMC}^{(k)}$  and  $\phi_2 \in \bigcup_{h \leq k} \text{AMC}^{(h)}$ ,  $\phi_1 \vee \phi_2 \in \text{AMC}^{(k)}$ .
- **Case 2d:** For each  $\phi_1 \in \bigcup_{h \leq k} \text{AMC}^{(h)}$  and  $\phi_2 \in \text{AMC}^{(k)}$ ,  $\phi_1 \vee \phi_2 \in \text{AMC}^{(k)}$ .
- **Case 2e:** For each  $\phi \in \text{AMC}^{(k)}$  and proposition variable  $x \in X$ ,  $\text{Ifp}x.\phi \in \text{AMC}^{(k)}$ .

Note that there could be free variables in the formulas classified in the above. The evaluation of such formulas for a game graph depends on the valuation of the free variables.

Given two game graphs  $G, H$  and an AMC formulas  $\phi$ , we say that two valuations of  $\nu$  and  $\nu'$  respectively of  $G$  and  $H$  are *consistent* if for every  $x \in X$ , there exists a  $\phi_1 \in F_0(x)$  such that  $\nu(x) = \kappa(G, \phi_1)$  and  $\nu'(x) = \kappa(H, \phi_1)$ . In the following, we adopt the AMC semantic notations in [3]. Given a game graph  $G$ , an AMC formula  $\phi$ , and a valuation  $\nu$  of

state set variables in  $X$ ,  $(\phi)^G(\nu)$  denotes the set of states of  $G$  that satisfy  $\phi$  with valuation  $\nu$ .

**Lemma 9.** Assume that  $G_k$  and  $H_k$  are defined in figures 4.4 and 4.5. For every  $k$ , AMC formula  $\phi \in \text{AMC}^{(k)}$ , and two consistent valuations  $\nu$  and  $\nu'$  respectively of  $G_k$  and  $H_k$ ,  $r_k \in (\phi)^{G_k}(\nu)$  iff  $s_k \in (\phi)^{H_k}(\nu')$ .

**Proof :** We use an induction on  $k$  to prove the lemma.

**Base case:** When  $\phi$  is in case 1a through 1d, the lemma follows straightforwardly. In case 1e,  $(\text{Ifp}x.\psi)^{G_1}(\nu)$  can be expanded as follows. We let

- $\psi^{G_1, \nu, (0)}$  be  $(\psi[x \mapsto \text{false}])^{G_1}(\nu)$  and
- for each  $h > 0$ ,  $\psi^{G_1, \nu, (h)}$  be  $(\phi[x \mapsto \psi_1^{G_1, (h-1)}])(\nu)$ .

Then  $(\text{Ifp}x.\phi)^{G_1}(\nu) = \bigcup_{h \geq 0} \psi^{G_1, \nu, (h)}$  according to the semantics of AMC. Similarly,  $(\text{Ifp}x.\phi)^{H_1}(\nu') = \bigcup_{h \geq 0} \psi^{H_1, \nu', (h)}$ . According to the same argument for cases 1a through 1d, for each  $h \geq 0$ ,  $r_0 \in \psi^{G_1, \nu, (h)}$  iff  $s_0 \in \psi^{H_1, \nu', (h)}$ . Thus it is clear that  $r_0 \in (\text{Ifp}x.\phi)^{G_1}(\nu)$  iff  $s_0 \in (\text{Ifp}x.\phi)^{H_1}(\nu')$ . Thus the lemma is proven in this case.

**Induction step:** To tell the difference between  $G_k$  and  $H_k$ , we need a formula with the following structure.

- At least one nesting of operators like  $\langle \dots \rangle \bigcirc$  in a least fixpoint operation to infer the reachability of  $G_{k-1}$  and  $H_{k-1}$ .
- A modal subformula nested inside a  $\langle \dots \rangle \bigcirc$  modal operator of  $\phi$  that can tell the difference of  $G_{k-1}$  and  $H_{k-1}$ . But according to the inductive hypothesis, this is impossible.

Thus the lemma is proven. □

Then with Lemma 9, we conclude the proof for Lemma 10 in the following.

**Lemma 10.** For every AMC formula  $\phi$ , there are two game graphs that  $\phi$  cannot distinguish while  $\langle 1 \rangle (\langle +2 \rangle \Box p) \wedge \langle +2 \rangle \Box q$  can.

**Proof :** In the proof for Lemma 9, it is apparent that  $\phi$  with  $\phi \in \text{AMC}^{(k)}$  cannot tell  $G_k$  and  $H_k$ . □

By the same argument in [3], for one-agent game, BSIL coincides with CTL and is not as expressive as AMC.

**Lemma 11.** For game graphs of one agent, AMC is strictly more expressive than BSIL.

**Proof :** For one-agent games, AMC is equivalent to  $\mu$ -calculus and BSIL is equivalent to CTL which is strictly less expressive than  $\mu$ -calculus.  $\square$

A comment on Lemmas 6 and 10 is that the path modal formulas in the lemmas can be changed independently to  $\Diamond\neg p$  and  $\Diamond\neg q$  without affecting the validity of the lemma. This can be used to show that the example properties in the introduction are indeed inexpressible in  $\text{ATL}^*$ , GL, and AMC.

### 4.3.3 Comparison with $\text{ATL}^*$

It is easy to see that BSIL is a super-class of ATL. Thus we have the following lemma.

**Lemma 12.** BSIL is at least as expressive as ATL.  $\square$

Then Lemmas 6 and 10 lead to the fact that there are some BSIL properties that  $\text{ATL}^*$  cannot express since GL and AMC are both super-class of  $\text{ATL}^*$  [3].

**Lemma 13.** For every  $\text{ATL}^*$  formula  $\phi$ , there are two game graphs that  $\phi$  cannot distinguish while  $\langle 1 \rangle (\langle +2 \rangle \Box p) \wedge \langle +2 \rangle \Box q$  can.  $\square$

Lemmas 12 and 13 together establish that ATL is strictly less expressive than BSIL. Then the following lemma shows the reverse direction.

**Lemma 14.**  $\text{ATL}^*$  formula  $\langle 1 \rangle \Box \Diamond p$  is not equivalent to any BSIL formula.

**Proof :** The proof is similar to the proof for the inexpressibility of  $\langle 1 \rangle \Box \Diamond p$  with ATL [3].

$\square$

Lemmas 13 and 14 together establish that  $\text{ATL}^*$  and BSIL are not comparable in expressiveness.

## 4.4 Algorithm and Complexity

The model-checking problem of BSIL is contained in PSPACE mainly due to the restriction that disallows negation in tree formulas. As in the model-checking algorithms of ATL [3], we can evaluate the proper state subformulas independently and then treat them

as auxiliary propositions. Moreover, as in the evaluation of  $\Diamond$ -formulas in ATL model-checking, if a  $\Diamond$ -formula can be enforced with an S-profile, it can be enforced in a finite number of steps along every play compatible with the strategy in a computation tree. Once a bound  $b$  for this finite number of steps is determined, we can enumerate all strategies embedded in the computation tree up to depth  $b$  and try to find one that enforces a BSIL formula.

As explained in section ??, our algorithm also labels subformulas and their symbolic S-profiles on the nodes in a computation tree. The formula is satisfied if and only if we can find a finite tree top with labels consistent with the SQs and SIQs in the input formula and can be extended to an infinite computation tree. In Subsection 4.4.1, we derive the labels, i.e., subformulas and their symbolic S-profiles, that are sufficient for our model-checking algorithm and present two procedures:

- `localEval()` that checks whether the satisfaction of formulas at a state can be decided locally.
- `sucSet()` that nondeterministically chooses a scheme to pass down the subformulas and their symbolic S-profiles to the child nodes without violating the restrictions on the S-profiles declared with the SQs and SIQs.

There are, however, severe differences between exploring a computation tree for BSIL and exploring one for ATL [3]. For BSIL, we have to take the interaction of strategies into account. For example, we may have to enforce a subformula  $\langle 1 \rangle ((\langle +2 \rangle \Diamond p) \wedge \langle +2 \rangle (\Box q \vee \Diamond r))$ . Then, when exploring the computation tree, we may follow two strategies of Agent 2, one to enforce  $\Diamond p$  and the other to enforce  $\Box q$  or  $\Diamond r$ . There are the following situation for the interaction between these two strategies. The two strategies may make the same decision all the way until we reach a tree node  $v$ . (For turn-based games,  $v$  has to be owned by Agent 2.) This can be conceptualized as passing the obligations of  $\Diamond p$  and  $\Box q \vee \Diamond r$  along the path from the root to  $v$ . Then, at node  $v$ , the two strategies may differ in their decisions and pass down the two obligations to different branches.

Then, in Subsection 4.4.2, we present our algorithm in two parts, one for model-checking BSIL state formulas and the other for model-checking BSIL tree formulas. In

Subsection 4.4.3, we prove the correctness of the algorithm. In Subsection 4.4.4, we show that our algorithm is in PSPACE. Together with Lemma 22 in Subsection 4.4.5, we then establish the PSPACE-completeness of the BSIL and  $ATL^+$  model-checking problems.

#### 4.4.1 Computing path obligations and passing them down the computation tree

We use  $\{a_1 \mapsto s_1, \dots, a_n \mapsto s_n\}$  to denote a partial function that maps  $a_i$  to  $s_i$  for each  $i \in [1, n]$ . Given a partial function  $f$ , we denote the domain of  $f$  by  $def(f)$ . Inheriting the notations in [6], we may also represent the mapping as  $(a_1, s_1)(a_2, s_2) \dots (a_n, s_n)$ .

We need some special techniques in checking tree formulas. We adopt the concept of strategy variables from [5,6]. A *strategy variable binding* (*SV-binding* for short) is a partial function from  $[1, m]$  to strategy variables. Given an SV-binding  $\Lambda$ ,  $\Lambda \circ (a_1, s_1) \dots (a_n, s_n)$  is the SV-binding that is identical to  $\Lambda$  except that Agent  $a_i$  is bound to  $s_i$  for every  $i \in [1, n]$ .

Suppose that we are given SV-bindings  $\Lambda_1, \dots, \Lambda_n$  and S-profiles  $\Sigma_1, \dots, \Sigma_n$ . We say that  $\Lambda_1, \dots, \Lambda_n$  *matches*  $\Sigma_1, \dots, \Sigma_n$  if, and only if, for every  $a \in [1, m]$  and  $i, j \in [1, n]$  with  $a \in def(\Sigma_i) \cap def(\Sigma_j)$ ,  $\Sigma_i(a) = \Sigma_j(a)$  if, and only if,  $\Lambda_i(a) = \Lambda_j(a)$ .

Given an SV-binding  $\Lambda$  and a state, tree, or path formula  $\psi$ ,  $\Lambda\psi$  is called a *bound formula*.  $\Lambda\psi$  is a *bound path obligation* (*BP-constraint*) if  $\psi$  is a Boolean combination of path formulas. A Boolean combination of BP-obligations is called a *Boolean bound formula* (*BB-formula*). The strategy variables in BB-formula are only used to tell whether or not two path properties are to be enforced with the same strategy. For example, the property  $\langle 1 \rangle((\langle +2 \rangle \Diamond p) \wedge \langle +2 \rangle(\Box q \vee \Diamond r))$  can be rewritten as BB-formula  $((1, s_1)(2, s_2) \Diamond p) \wedge (1, s_1)(2, s_3) \Box q \vee \Diamond r$ , which says that Agent 1 must use the same strategy to fulfill both  $\Diamond p$  and  $\Box q \vee \Diamond r$ , while Agent 2 may use different strategies to fulfill these two path properties.

Suppose we are given a function  $\pi$  that maps symbolic strategy names to strategies. Similar to the semantics of strategy logics [6] with strategy variables, we can also define the satisfaction of BB-formulas  $\Lambda\psi$  at a state  $q$  with  $\pi$ , in symbols  $\mathcal{G}, q \models^\pi \Lambda\psi$ , as follows.



Table 4.1: Rewriting rules for BB-formulas

$bf(\Lambda \neg \neg \phi)$	$\equiv$	$bf(\Lambda \phi)$
$bf(\Lambda(\tau_1 \vee \tau_2))$	$\equiv$	$bf(\Lambda \tau_1) \vee bf(\Lambda \tau_2)$
$bf(\Lambda(\tau_1 \wedge \tau_2))$	$\equiv$	$bf(\Lambda \tau_1) \wedge bf(\Lambda \tau_2)$
$bf(\Lambda \langle a_1, \dots, a_n \rangle \psi)$	$\equiv$	$bf(\{a_1 \mapsto newVar(), \dots, a_n \mapsto newVar()\} \psi)$
$bf(\Lambda \langle +a_1, \dots, a_n \rangle \psi)$	$\equiv$	$bf(\Lambda \circ \{a_1 \mapsto newVar(), \dots, a_n \mapsto newVar()\} \psi)$
$bf(\Lambda \bigcirc \phi_1)$	$\equiv$	$\Lambda \bigcirc bf(\emptyset \phi_1)$
$bf(\Lambda \neg \bigcirc \phi_1)$	$\equiv$	$\Lambda \bigcirc bf(\emptyset \neg \phi_1)$
$bf(\Lambda \phi_1 \mathbf{U} \phi_2)$	$\equiv$	$\Lambda bf(\emptyset \phi_1) \mathbf{U} bf(\emptyset \phi_2)$
$bf(\Lambda \neg \phi_1 \mathbf{U} \phi_2)$	$\equiv$	$\Lambda((bf(\emptyset \phi_1) \mathbf{U} bf(\neg \emptyset(\phi_1 \vee \phi_2))) \vee \Box bf(\emptyset \neg \phi_2))$
$bf(\Lambda p) \equiv p$	$;$	$bf(\Lambda \neg p) \equiv \neg p$
$bf(\Lambda true) \equiv true$	$;$	$bf(\Lambda \neg true) \equiv false$
$bf(\Lambda false) \equiv false$	$;$	$bf(\Lambda \neg false) \equiv true$

$\phi_1, \phi_2$ : state or path formulas.  $\tau_1, \tau_2$ : tree formulas.  $\psi_1, \psi_2$ : tree or path formulas.

- $\mathcal{G}, q \models^\pi \Lambda_1 \psi_1 \vee \Lambda_2 \psi_2$  iff  $\mathcal{G}, q \models^\pi \Lambda_1 \phi_1$  or  $\mathcal{G}, q \models^\pi \Lambda_2 \phi_2$  holds.
- $\mathcal{G}, q \models^\pi \Lambda_1 \phi_1 \wedge \Lambda_2 \phi_2$  iff both  $\mathcal{G}, q \models^\pi \Lambda_1 \phi_1$  and  $\mathcal{G}, q \models^\pi \Lambda_2 \phi_2$  hold.
- Given an SV-binding  $\Lambda$  and a path formula  $\psi_1$  with an S-profile  $\Sigma = \{a \mapsto \pi(\Lambda(a)) \mid a \in def(\Lambda)\}$ ,  $\mathcal{G}, q \models^\pi \Lambda \psi_1$  iff, for all plays  $\rho$  compatible with  $\Sigma$  from  $q$ ,  $\rho \models_\Sigma \psi_1$  holds.

In Table 4.1, we present equivalence rules to rewrite state, tree, and path formulas to BB-formulas using the procedure  $bf()$ . For convenience, we use a procedure  $newVar()$  that returns a strategy variable that has not been used before. In general, the semantics of BSIL deals with the satisfaction of a set of subformulas bound to different S-profiles. The following two lemmas relate the rules in Table 4.1 with the semantics of BSIL formulas.

**Lemma 15.** Suppose we are given a state  $q$ , a BSIL subformula  $\psi$ , and a S-profile  $\Sigma$  such that  $\mathcal{G}, q \models_\Sigma \psi$ . There exist an SV-binding  $\Lambda$  and a function  $\pi$  such that  $\mathcal{G}, q \models^\pi bf(\Lambda \psi)$ .

**Proof :** We construct  $\Lambda$  and  $\pi$  as follows. Without loss of generality, we assume  $\psi$  is unique in the input formula. For every  $a \in def(\Sigma)$ , we let  $\Lambda(a) = s_a^\psi$  and  $\pi(s_a^\psi) = \Sigma(a)$ . It is clear that the functional composition of  $\Lambda$  and  $\pi$  is actually  $\Sigma$ . Then, according to the semantics of  $\mathcal{G}, q \models^\pi bf(\Lambda \psi)$  presented in the above,  $\mathcal{G}, q \models^\pi bf(\Lambda \psi)$  since for all play  $\rho$  compatible with  $\Sigma$  from  $q$ ,  $\rho \models_\Sigma \psi$ . Thus the lemma is proven.  $\square$

**Lemma 16.** Suppose we are given a state  $q$ , a BSIL subformula  $\psi$ , a SV-bindings  $\Lambda$ , and a

function  $\pi$  such that  $\mathcal{G}, q \models^\pi bf(\Lambda\psi)$ . Then there exist an S-profile  $\Sigma$  such that  $\mathcal{G}, q \models_\Sigma \psi$ .

**Proof :** We can construct  $\Sigma$  by defining, for all  $a \in def(\Lambda)$ ,  $\Sigma(a) = \pi(\Lambda(a))$ . Thus,  $\Sigma$  is the functional composition of  $\Lambda$  and  $\pi$ . Then according to the semantics of  $\mathcal{G}, q \models^\pi bf(\Lambda\psi)$  presented in the above,  $\mathcal{G}, q \models^\pi bf(\Lambda\psi)$  implies  $\mathcal{G}, q \models_\Sigma \psi$ . Thus the lemma is proven.  $\square$

To ease the presentation of our algorithms, we also assume that there is a procedure that rewrites a BB-formula to an equivalent BB-formula in disjunctive normal form. Specifically, a *disjunctive normal BB-formula (DNBB-formula)* is the disjunction of conjunctions of BP-obligations. The rewriting of a BB-formula  $\phi$  to a DNBB-formula can be done by repeatedly applying the distribution law of conjunctions of disjunctions until a DNBB-formula is obtained.

**Example 17.** DNBB-formula rewriting: We have the following rewriting process for a BSIL formula for five agents.

$$\begin{aligned}
& bf(\emptyset\langle 1, 2 \rangle (\langle +3 \rangle (\Box p \vee \Diamond q) \wedge \langle +3 \rangle (\langle +2 \rangle \Diamond r \vee \langle +4 \rangle \Box q))) \\
& \equiv bf((1, s_1)(2, s_2) (\langle +3 \rangle (\Box p \vee \Diamond q) \wedge \langle +3 \rangle (\langle +2 \rangle \Diamond r \vee \langle +4 \rangle \Box q))) \\
& \equiv bf((1, s_1)(2, s_2) \langle +3 \rangle (\Box p \vee \Diamond q)) \wedge bf((1, s_1)(2, s_2) \langle +3 \rangle (\langle +2 \rangle \Diamond r \vee \langle +4 \rangle \Box q)) \\
& \equiv bf((1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q)) \wedge bf((1, s_1)(2, s_2)(3, s_4)(\langle +2 \rangle \Diamond r \vee \langle +4 \rangle \Box q)) \\
& \equiv (1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q) \\
& \quad \wedge ((1, s_1)(2, s_5)(3, s_4)\Diamond r \vee (1, s_1)(2, s_2)(3, s_4)(4, s_6)\Box q) \\
& \equiv ((1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q) \wedge (1, s_1)(2, s_5)(3, s_4)\Diamond r) \\
& \quad \vee ((1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q) \wedge (1, s_1)(2, s_2)(3, s_4)(4, s_6)\Box q)
\end{aligned}$$

This DNBB-formula sheds some light on the analysis of BSIL formulas. As can be seen, the formula is satisfied iff one of the two outermost disjuncts is satisfied. Without loss of generality, we examine the first disjunct:

$$\eta_1 \equiv (1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q) \wedge (1, s_1)(2, s_5)(3, s_4)\Diamond r$$

There are the following two S-profiles involved in the satisfaction of the formula.

- $\Sigma_1$  for  $(1, s_1)(2, s_2)(3, s_3)$  of  $\{1, 2, 3\}$  used to satisfy  $\Box p \vee \Diamond q$ .
- $\Sigma_2$  for  $(1, s_1)(2, s_5)(3, s_4)$  of  $\{1, 2, 3\}$  used to satisfy  $\Diamond r$ .

This disjunct imposes the restrictions that  $\Sigma_1$  and  $\Sigma_2$  must agree in their moves by agent

1. (Or for turn-based games, they must agree in their choices at nodes owned by Agent

1.) Similarly, we can examine

$$\eta_2 \equiv (1, s_1)(2, s_2)(3, s_3)(\Box p \vee \Diamond q) \wedge (1, s_1)(2, s_2)(3, s_4)(4, s_6)\Box q$$

There is a new S-profile introduced.

- $\Sigma_3$  for  $(1, s_1)(2, s_2)(3, s_4)(4, s_6)$  of  $\{1, 2, 3, 4\}$  used to satisfy  $\Box q$ .

This disjunct imposes the restrictions that  $\Sigma_1$  and  $\Sigma_3$  must agree in their moves by Agents 1 and 2. In the following, we use the observation in this example to construct structures from DNBB-formulas for the model-checking of conjunctive DNBB-formulas.  $\square$

For the ease of notation, we represent a conjunctive DNBB-formula  $\eta$  as a set of BP-obligations in our algorithms. Our goal is to design a computation tree exploration procedure that given a set  $C$  of BP-obligations, labels each node in the tree with a subset of  $C$  for the set of path formulas that some S-profiles have to enforce without violating the restrictions of strategy interaction imposed in  $C$  through the strategy variables. In the design of the procedure, one central component is how to label the children of a node with appropriate sets of BP-obligations as inherited path obligations from  $C$ . We need two basic procedures for this purpose. The first is to evaluate the truth values of path literals in BP-obligations with a proposition interpretation when possible. Specifically, when we can deduce the truth values of U-formulas and  $\Box$ -formulas from the truth values of propositions (or state subformulas) at a state, the procedure changes the respective U-formula and  $\Box$ -formula to their respective truth values. The procedure is as follows.

---

`localEval( $W, \theta$ ) //  $\lambda()$`  has been extended with satisfied state subformulas at each state.

- 1: **switch** ( $\theta$ )
- 2: **case** *true* or *false*: **return**  $\theta$
- 3: **case**  $p$ : **if**  $p \in W$  **then** **return** *true* **else** **return** *false* **end if**
- 4: **case**  $\neg p$ : **if**  $p \in W$  **then** **return** *false* **else** **return** *true* **end if**
- 5: **case**  $\theta_1 \vee \theta_2$ :
- 6:   Let  $\theta_1$  be `localEval( $W, \theta_1$ )` and  $\theta_2$  be `localEval( $W, \theta_2$ )`.

```

7:  if  $\theta_1$  is true or  $\theta_2$  is true then return true.
8:  else if  $\theta_1$  is false then return  $\theta_2$ . else if  $\theta_2$  is false then return  $\theta_1$ . else return
     $\theta_1 \vee \theta_2$ .
9:  end if
10: case  $\theta_1 \wedge \theta_2$ :
11:   Let  $\theta_1$  be localEval( $W, \theta_1$ ) and  $\theta_2$  be localEval( $W, \theta_2$ ).
12:   if  $\theta_1$  is false or  $\theta_2$  is false then return false.
13:   else if  $\theta_1$  is true then return  $\theta_2$ . else if  $\theta_2$  is true then return  $\theta_1$ . else return
     $\theta_1 \wedge \theta_2$ .
14:  end if
15: case  $\bigcirc \phi_1$ : return  $\theta$ 
16: case  $\Box \phi_1$ : if  $\phi_1 \notin W$  then return false else return  $\theta$  end if
17: case  $\phi_1 \cup \phi_2$ : if  $\phi_2 \in W$  then return true else if  $\phi_1 \notin P$  then return false else
    return  $\theta$  end if
18: end switch

```

---

Statements 17 checks if  $\Lambda\theta$  is fulfilled. When it is fulfilled,  $\theta$  is changed to *true*. Statements 17 and 16 also check if  $\Lambda\theta$  is violated. When a violation happens,  $\theta$  is changed to *false*.

Then we need a procedure, `next()`, that calculates the BP-obligations passed down from a previous state. This is simply done by replacing every  $\bigcirc\phi$  by  $\phi$  in the BP-obligations.

With the two basic procedures defined above, we now present a procedure that nondeterministically calculates sets of BP-obligations passed down to the successor states. This is accomplished with the procedure `sucSet`( $q, C$ ) in the following. Given a node  $q$  in the computation tree and a set  $C$ , the procedure nondeterministically returns an assignment of BP-obligations to children of  $q$  to enforce the BP-obligations in  $C$  without violating the strategy interaction of BP-obligations.

---

`sucSet`( $q, C$ ) //  $\lambda()$  has been extended with satisfied state subformulas at each state.

```

1: Convert  $C$  to  $\{\Lambda_{\text{localEval}}(\lambda(q), \theta) \mid \Lambda\theta \in C\}$ .
2: if  $\Lambda_{\text{false}} \in C$  then return  $\emptyset$  end if
3: Let  $S$  be the set of all symbolic strategy variables in  $C$ . That is,  $S = \{s \mid a \mapsto s \in \Lambda, \Lambda\theta \in C\}$ .
4: Nondeterministically pick an  $\alpha_s \in \Delta$  for each  $s \in S$ .
5: Let  $K$  be  $\{(q', \emptyset) \mid (q, q') \in R\}$ .
6: for each  $\Lambda\theta \in C$  do
7:   if for all  $(q, q')$ , there is an  $a \in \text{def}(\Lambda)$  with  $\delta((q, q'), a) \neq s_{\Lambda(a)}$  then return  $\emptyset$  end if
8:   for  $(q', C') \in \Delta$  with  $\forall a \in \text{def}(\Lambda) (\delta((q, q'), a) = s_{\Lambda(a)})$  do
9:     Replace  $(q', C')$  with  $(q', C' \cup \{\Lambda_{\text{next}}(\theta)\})$  in  $K$ .
10:   end for
11: end for
12: return  $K$ .

```

---

The nondeterministic choices at statement 4 make sure that one symbolic strategy variable is mapped to exactly one move. The loop at statement 6 iterates through all the path obligations at the current node and passes them down to the children if necessary. The if-statement at line 7 checks whether all obligations can be passed down to some children. If some obligations are not passed due to mismatch between moves of the strategies and the labels on the transitions, then we return with failure. Otherwise, statement 8 passes the obligations to all children with matching transition labels. The obligations to children are recorded in  $K$  which is returned with success at statement 12.

#### 4.4.2 Procedures for checking BSIL properties

The procedure in the following checks a BSIL state property  $\phi$  at a state  $q$  of  $A$ .

---

$\text{checkBSIL}(q, \phi)$

```

1: if  $\phi$  is  $p$  then if  $\phi \in \lambda(q)$  then return true. else return false. end if
2: else if  $\phi$  is  $\phi_1 \vee \phi_2$  then return  $\text{checkBSIL}(q, \phi_1) \vee \text{checkBSIL}(q, \phi_2)$ 
3: else if  $\phi$  is  $\neg\phi_1$  then return  $\neg\text{checkBSIL}(q, \phi_1)$ 

```

4: **else if**  $\phi$  is  $\langle A \rangle \tau$  for a tree or path formula  $\tau$  **then return**  $\text{checkTree}(q, \langle A \rangle \tau)$   
 5: **end if**

---

The procedure is straightforward and works inductively on the structure of the input formula. For convenience, we need procedure  $\text{checkSetOfBSIL}(Q', \phi_1)$  in the following that checks a BSIL property  $\phi_1$  at each state in  $Q'$ .

---

$\text{checkSetOfBSIL}(Q', \phi_1)$

---

1: **if**  $\phi_1 \notin P \cup \{true, false\}$  **then**  
 2:   **for each**  $q' \in Q'$  **do**  
 3:     **if**  $\text{checkBSIL}(q', \phi_1)$  **then** Let  $\lambda(q')$  be  $(\lambda(q') \cup \{\phi_1\}) - \{\neg\phi_1\}$ .  
 4:     **else** Let  $\lambda(q')$  be  $(\lambda(q') - \{\phi_1\}) \cup \{\neg\phi_1\}$ . **end if**  
 5:   **end for**  
 6: **end if**

---

Then we use procedure  $\text{checkTree}(q, \langle A \rangle \tau)$  in the following to check if a state  $q$  satisfies  $\langle A \rangle \tau$ .

---

$\text{checkTree}(q, \langle A \rangle \tau)$

---

1: Rewrite  $bf(\emptyset \langle A \rangle \tau)$  to DNBB-formula  $\eta_1 \vee \dots \vee \eta_n$ .  
 2: **for**  $i \in [1, n]$  **do**  
 3:   Represent  $\eta_i$  as a set  $C$  of BP-obligations.  
 4:   **for each**  $\Lambda\theta$  in  $C$ . **do**  
 5:     **if**  $\theta$  is  $\bigcirc\phi_1$  **then**  $\text{checkSetOfBSIL}(\{q' \mid (q, q') \in \mathcal{R}\}, \phi_1)$ .  
 6:     **else if**  $\theta$  is  $\phi_1 U \phi_2$  **then**  $\text{checkSetOfBSIL}(Q, \phi_1); \text{checkSetOfBSIL}(Q, \phi_2);$   
       **end if**  
 7:   **end for**  
 8:   **if**  $\text{recTree}(q, C)$  **then return true. end if**  
 9: **end for**  
 10: **return false.**

---

We first rewrite  $\langle A \rangle \tau$  to its DNBB-formula at statement 1 by calling  $bf(\emptyset \langle A \rangle \tau)$  and using the distribution law of conjunctions over disjunctions. (In practice, to contain the complexity in PSPACE, we only need to enumerate the disjuncts of the DNBB-formula in

PSPACE.) We then iteratively check with the loop starting from statement 2 if  $\langle A \rangle \tau$  is satisfied due to one of its conjunctive DNBB-formula components of  $\langle A \rangle \tau$ . At statement 3, we construct the set  $C$  of BP-obligations of the component. We evaluate the subformulas with the inner loop starting at statement 4. Finally at statement 8, we explore the computation tree, with procedure  $\text{recTree}(q, C)$  in the following, and pass down the path obligations to the children according to the restrictions of the SV-binding in  $C$ .

---

$\text{recTree}(q, C)$

```

1: if  $(q, C)$  coincides with an ancestor in the exploration then
2:   if there is no  $\Lambda\phi_1 \cup \phi_2$  in  $C$  then return true; else return false. end if
3: end if
4: if  $\text{sucSet}(q, C)$  is empty then return false end if
5: for each  $(q', C') \in \text{sucSet}(q, C)$  with  $C' \neq \emptyset$  do
6:   if  $\text{recTree}(q', C')$  is false then return false. end if
7: end for
8: return true.

```

---

Note that procedure  $\text{recTree}(q, C)$  is nondeterministic since it employs  $\text{sucSet}(q, C)$  to nondeterministically calculate an assignment  $\Delta$  of path obligations to the children of  $q$ .

#### 4.4.3 Correctness proof of the algorithm

In order to prove the correctness of this algorithm, we define *obligation distribution trees* (OD-trees) in the following. An OD-tree for a set  $C$  of BP-obligations and game graph  $\mathcal{G}$  from a state  $q_0 \in Q$  is a labeled computation tree  $\langle V, \bar{r}, \alpha, E, \beta \rangle$  with the following restrictions.

- $V$  is the set of nodes in the tree.
- $\bar{r} \in V$  is the root of the tree.
- $\alpha : V \mapsto Q$  labels each tree node with a state. Also  $\alpha(\bar{r}) = q_0$ .
- $E \subseteq V \times V$  is the set of arcs of the tree such that, for each  $(q, q') \in R$ , there exists an  $(v, v') \in E$  with  $\alpha(v) = q$  and  $\alpha(v') = q'$ .

- $\beta : V \mapsto 2^C$  labels each node with a subset of  $C$  for path formulas in  $\chi$  that need to be fulfilled at a node. Moreover, we have the following restrictions on  $\beta$ .
  - $C = \beta(r)$ .
  - For every  $v \in V$ , there exists a  $\Delta = \text{sucSet}(\alpha(v), \beta(v))$  such that, for every  $(q', C') \in \Delta$ , there exists a  $(v, v') \in E$  with  $\alpha(v') = q'$  and  $\beta(v') = C'$ .

The OD-tree is *fulfilled* iff, for every path  $v_0v_1 \dots v_k \dots$  along the tree from the root, there exists an  $h \geq 0$  such that, for every  $j \geq h$ , there is no  $\Lambda\phi_1 U \phi_2 \in \beta(v_j)$ . We have the following connection between an OD-tree and an execution of procedure  $\text{recTree}(q, C)$  from the root of an OD-tree.

**Lemma 18.** For a set  $C$  of BP-obligations,  $\text{recTree}(q, C)$  returns true iff there exists a fulfilled OD-tree for  $C$  and  $\mathcal{G}$  from  $q$ .

**Proof :** In order to prove the lemma, we show both directions.

$(\Rightarrow)$  : It is straightforward to see that  $\text{recTree}(q, C)$  returns true only if a finite tree has been constructed with leafs duplicating their ancestors. According to statement 2 of  $\text{recTree}(q, C)$ , it is clear that along the path from that ancestor to a leaf, no node is labeled with a BP-obligation of the form  $\Lambda\phi_1 U \phi_2$  by  $\beta$ . Thus we can extend the leaves by duplicating the subtree rooted at their duplicating ancestors. In this way, we can extend the finite tree to a fulfilled OD-tree.

$(\Leftarrow)$  : Suppose there exists a fulfilled OD-tree for  $C$  and  $\mathcal{G}$  from  $q$ . Since all infinite paths from the root stabilize to suffices without index labels by  $\beta$  for until-formulas (as the tree is finitely branching, it would otherwise contain an infinite path with standing utility by Königs lemma), we can repeatedly replace every subtree  $T$  with a subtree  $T'$  of  $T$  such that the root of  $T$  and  $T'$  have the same  $\alpha$  and  $\beta$  labels. We can repeat this replacement until no node labeled with an until-formula has the same  $\alpha$  and  $\beta$  labels as one of its descendants. The existence of such an OD-tree after the replacements implies that  $\text{recTree}(q, C)$  eventually explores such a tree, finds the termination condition at all leafs, and returns *true*. □

**Lemma 19.** Given a conjunctive DNBB-formula  $\eta$  represented as a set  $C$  of BP-obligations, there exists a function  $\pi$  on strategy variables in  $\eta$  with  $\mathcal{G}, q \models^\pi \eta$  iff there exists a fulfilled



OD-tree for  $\mathcal{G}$  and  $C$  from  $q$ .

**Proof :** The lemma can also be proven in two directions. In the forward direction, we can use  $\pi$  to construct S-profiles to enforce  $\eta$ . The S-profiles can then be used to construct a fulfilled OD-tree for  $\mathcal{G}$  and  $C$  from  $q$ .

In the backward direction, we can follow the paths and obligations that are passed-down in the OD-tree and construct S-profiles that enforce  $\eta$ . Then, from these S-profiles, due to the one-to-one correspondence between the strategy variables and the strategies in the S-profiles, we can define a  $\pi$  with  $\mathcal{G}, q \models^\pi \eta$ .  $\square$

The correctness of procedure `recTree`( $q, C$ ) then directly follows from Lemmas 18 and 19. Then the correctness of procedure `checkBSIL`( $q, \phi$ ) follows by a structural induction on a given BSIL formula and the correctness of procedure `recTree`( $q, C$ ).

**Lemma 20.** Given a state  $q$  in  $\mathcal{G}$ , `checkBSIL`( $q, \chi$ ) iff  $\mathcal{G}, q \models_\perp \chi$ .  $\square$

#### 4.4.4 Complexities of the algorithm

The algorithm that we presented in Subsections 4.4.1 and 4.4.2 can run in PSPACE mainly because we can enumerate the conjuncts in a DNF in PSPACE and can implement procedure `recTree`( $q, C$ ) with a stack of polynomial height. To see this, please recall that we use the procedure `sucSet`( $q, C$ ) to calculate the assignment of BP-obligations to the children to  $q$  in the computation tree. Specifically, procedure `sucSet`( $q, C$ ) non-deterministically returns a set  $\Delta$  with elements of the form  $(q', C')$  such that  $(q, q') \in R$  and  $C' \subseteq C$  since in procedure `sucSet`( $q, C$ ), a path obligation  $\Lambda\theta$  is passed down to a child and recorded in the corresponding  $C'$  only when it matches the for loop condition at statement 8. Thus, along any path in the OD-tree, the sets of literal bounds never increase. Moreover, when there is a node in the exploration of OD-tree that coincides with an ancestor, we backtrack in the exploration. This implies that, along any path segment longer than  $|Q|$ , one of the following two conditions hold.

- A backtracking happens at the end of the segment.
- The sets of BP-obligations along the segment must decrease in size at least once.

These conditions lead to the observation that, with procedure `sucSet`( $q, C$ ), the recursive

exploration of a path can grow no longer than  $1 + |C| \cdot |Q|$ . This leads to the following lemma.

**Lemma 21.** The BSIL model-checking algorithm in Subsections 4.4.1 and 4.4.2 is in PSPACE.

**Proof :** For convenience, we let  $\#(\chi)$  be the number of modal formulas in  $\chi$ . Following the argument from above, it is straightforward to check that to explore an OD-tree, we only need a stack of at most  $1 + \#(\chi) \cdot |Q|$  frames. In each frame, we only need to record a state in  $Q$ , a subset of  $[1, \#(\chi)]$  for the path obligations, and a  $\Delta$  returned from procedure  $\text{sucSet}(q, C)$ . Procedure  $\text{sucSet}(q, C)$  can be nondeterministically computed by randomly assigning the obligations in  $C$  to the successors of  $q$  and check if the assignment satisfies strategy interaction restriction of the BP-obligations. Procedure  $\text{checkBSIL}(q, \phi)$  can then be executed in space cubic in the size of  $\phi$  for the  $\Delta$ 's at nodes along the path. Thus, we conclude that the algorithm is a PSPACE algorithm.  $\square$

A rough analysis of the time complexity of our algorithm follows. Let  $|\chi|$  be the length of a BSIL formula  $\chi$ . At each call to  $\text{sucSet}()$ , the size of  $C$  is at most  $|\chi|$ . The number of root-to-leaf paths in an OD-tree is at most  $|\chi|$  since we only have to pass down  $|\chi|$  BP-obligations. We can use the positions of the common ancestors of the leaves of such paths to analyze the number of the configurations of such OD-trees. The common ancestors can happen anywhere along the root-to-leaf paths. Thus, there are  $(1 + |\chi| \cdot |Q|)^{|\chi|}$  ways to arrange the positions of the common ancestors since the length of paths are at most  $1 + |\chi| \cdot |Q|$ . The number of ways that the BP-obligations can be assigned to the leaves is at most  $|\chi|^{|\chi|}$ . The number of state labeling of the nodes on the paths is at most  $|Q|^{|\chi| \cdot (1 + |\chi| \cdot |Q|)}$ . Thus, given a  $C$ , the total number of different OD-trees is in  $O(|Q|^{|\chi| \cdot (1 + |\chi| \cdot |Q|)} |\chi|^{|\chi|} (1 + |\chi| \cdot |Q|)^{|\chi|}) = O(|Q|^{|\chi| \cdot (2 + |\chi| \cdot |Q|)} |\chi|^{2|\chi|})$ . There are  $O(2^{|\chi|})$  different possible values of  $C$ . There are at most  $|\chi|$  OD-trees to construct for the model-checking task. Thus, the total time complexity of our algorithm is in  $O(|\chi|^{2|\chi|} |Q|^{|\chi| \cdot (2 + |\chi| \cdot |Q|)} |\chi|^{2|\chi|})$ .

#### 4.4.5 Lower bound and completeness

We close by establishing the PSPACE lower bound for  $\text{ATL}^+$  model-checking, and hence the PSPACE-completeness of  $\text{ATL}^+$  and BSIL model-checking. This is done by reduction from the prenex QBF (quantified Boolean formula) satisfiability problem [64] to an  $\text{ATL}^+$  model-checking problem for a 2-agent game graph, where both the game graph and the  $\text{ATL}^+$  specification are linear in the prenex QBF problem we reduce from. We assume a QBF property  $\eta \equiv \nabla_1 p_1 \dots \nabla_l p_l (C_1 \wedge C_2 \wedge \dots \wedge C_n)$  with a set  $P = \{p_1, \dots, p_l\}$  of atomic propositions and the following restrictions.

- For each  $k \in [1, l]$ ,  $\nabla_k$  is either  $\exists$  or  $\forall$ .
- For each  $k \in [1, n]$ ,  $C_k$  is a clause  $l_{k,1} \vee \dots \vee l_{k,h_k}$ , where, for each  $j \in [1, h_k]$ ,  $l_{k,j}$  is a *literal*, i.e., either an atomic proposition or a negated atomic proposition.

Intuitively, the reduction is to translate the QBF formula to a game graph and an  $\text{ATL}^+$  formula for a traversal requirement on the game graph. The atomic propositions are then encoded as path constraints on the game graph. The interesting thing about the reduction is that the  $\text{ATL}^+$  formula (naturally) contains no SIQs at all. This reduction may be interpreted as that we get the SIQ in the expressiveness without paying extra computation complexity.

Suppose that  $\Gamma_p$  represents the subgraphs for the truth of an atomic proposition  $p$ . The rest of the game graph is partitioned into subgraphs  $\Omega_p$  responsible for the interpretation of atomic proposition  $p$  for all  $p \in P$ . Then the prenex QBF formula actually can be interpreted as a requirement for covering those  $\Gamma_p$ 's with the decisions in those  $\Omega_p$ 's. For example, the following formula  $\eta \equiv \exists p \forall q \exists r ((p \vee q \vee r) \wedge (\neg p \vee \neg r))$  can be read as *there exists a decision in  $\Omega_p$  such that for every decision in  $\Omega_q$ , there exists a decision in  $\Omega_r$  such that*

- *one of  $\Gamma_p$ ,  $\Gamma_q$ , and  $\Gamma_r$  is covered; and*
- *either  $\Gamma_p$  or  $\Gamma_r$  is not covered.*

The details of constructing those  $\Gamma_p$ 's and  $\Omega_p$ 's can be found in the proof of the following lemma that establishes the PSPACE complexity lower-bound.

**Lemma 22.** The  $\text{ATL}^+$  model-checking problem for turn-based game graphs is PSPACE-

hard.

**Proof :** Suppose we are given a prenex QBF  $\eta$  with propositions in  $P$ . We assume without loss of generality that all propositions are bound variables. (Note that, for the satisfiability problem, we can simply bind all free propositions by leading existential quantifiers.) We also use  $P$  for the atomic proposition set of  $\mathcal{G}$ . The idea is to use, for each  $p \in P$ , “ $\Diamond p$ ” to encode that  $p$  is true, while “ $\Box \neg p$ ” is used to encode that  $p$  is false. (Note that  $\Box \neg p \equiv \neg \Diamond p$ .) Then we construct a two-agent turn-based game graph  $G_\eta$  as shown in Figure 4.6, which reflects the structure of  $\eta$ : there is a sequence of (true) decisions (the states  $u_i$  and the state  $v_{n+1}$  have only one outgoing transition, and no true decision is taken there), which refer to the truth of the individual  $\Box p_i$ . These decisions are taken in the order given by the prenex quantifiers of  $\eta$ , and the existential decisions are taken by Agent 1 while the universal decisions are taken by Agent 2.

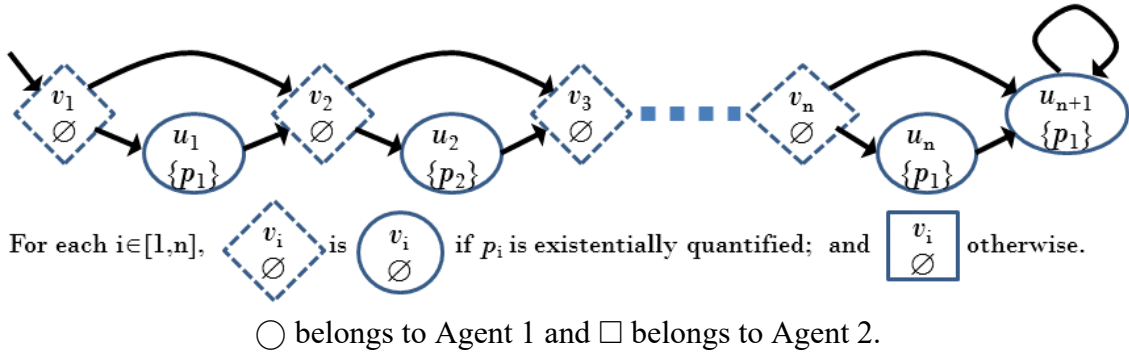


Figure 4.6: Game graph for the PSPACE-hardness proof of a Boolean formula with  $n$  propositions

In the graph, we use oval nodes for states owned by Agent 1 and square nodes for states owned by Agent 2. In each state, we put down its name and the set of atomic propositions that are true at the node. For each atomic proposition  $p_i \in P$ , we have a corresponding subgraph consisting of nodes  $v_i$  and  $u_i$ . The subgraph of  $u_i$  corresponds to  $\Gamma_{p_i}$  and that of  $v_i, u_i$  together corresponds to  $\Omega_{p_i}$ .

The design of the graph allows only at most one visit to states  $v_1, \dots, v_n$ . For all  $i \in [1, n]$ , state  $v_i$  is owned by Agent 2 if  $p_i$  is universally quantified; and owned by Agent 1 otherwise. If  $p_i$  is owned by Agent 1, then Agent 1 can choose either  $(v_i, u_i)$  or  $(v_i, v_{i+1})$ . If  $p_i$  is owned by Agent 2, then both choices of Agent 2 at node  $v_i$  must yield satisfaction

of  $\eta$ .

We construct an  $\text{ATL}^+$  formula  $\phi_\eta$  as  $\langle 1 \rangle \bigwedge_{1 \leq i \leq k} \bigvee_{1 \leq j \leq h_k} \tau(l_{i,j})$  where  $\tau(p) \stackrel{\text{def}}{=} \Diamond p$  and  $\tau(\neg p) \stackrel{\text{def}}{=} \Box \neg p$ . I.e.,  $\phi_\eta$  is obtained from  $\eta$  by replacing the leading quantifiers in  $\eta$  by  $\langle 1 \rangle$ .

We now show that  $\mathcal{G}_\eta \models \phi_\eta$  if, and only if,  $\eta$  is satisfied (i.e., if  $\eta$  is a tautology). The latter is the case if there is a winning strategy for a ‘satisfier’ in the following game between a ‘satisfier’ and a ‘refuter’: following the order of the bound variables, the ‘satisfier’ and ‘refuter’ choose the truth value of the existentially and universally bound variables, respectively. When all variables are assigned truth values, the ‘satisfier’ wins if the CNF formula is satisfied with these values. Otherwise, the refuter wins.

Taking a winning strategy of the satisfier in this game obviously provides a winning strategy for Agent 1 and, vice versa, a winning strategy of Agent 1 in the model-checking game can be used as a winning strategy for the ‘satisfier’ in the satisfaction game.  $\square$

We have an example for the reduction in the proof of Lemma 22.

**Example 23.** : Given  $\eta \equiv \exists p \forall q \exists r ((p \vee q \vee r) \wedge (\neg p \vee \neg r))$ , according to the construction in the proof of Lemma 22, we have the following  $\text{ATL}^+$  formula:  $\phi_\eta \stackrel{\text{def}}{=} \langle 1 \rangle (((\Diamond p) \vee (\Diamond q) \vee (\Diamond r)) \wedge ((\Box \neg p) \vee (\Box \neg r)))$ .  $\square$

Following Lemmas 22 and 21, we obtain the complexity of our model-checking problems.

**Lemma 24.** The BSIL and  $\text{ATL}^+$  model-checking problems are PSPACE-complete.  $\square$

# Chapter 5

## Temporal Cooperation Logic

### 5.1 TCL

TCL is an extension of BSIL which allow SQs appear after temporal operators. Thus, TCL is able to describe following complicate properties while BSIL is not.

Let's reuse the example of prisoner dilemma1. We can specify that prisoner 1 has a *nice* strategy which means he will not betray other players until been betrayed in TCL by following expression.

$$\langle 1 \rangle \Box ((\langle + \rangle \bigcirc \neg \text{betray}_1) \vee \bigvee_{a \neq 1} \text{betray}_a) \quad (\text{I})$$

Similarly, we can describe more involved strategies, such as 'Prisoner 2 will always betray when she does not have the power to force Player 1 to always play nice.'

$$\langle 2 \rangle (((\langle + \rangle \Box \text{betray}_2) \vee \langle + \rangle \Box ((\langle + \rangle \bigcirc \neg \text{betray}_1) \vee \bigvee_{a \neq 1} \text{betray}_a)) \quad (\text{J})$$

Moreover, we can describe deterministic Nash equilibria like 'the three agents can cooperate such that every agent either eventually leaves prison, or stays forever in prison regardless of her own strategy under the current strategies of the remaining prisoners' in TCL.

$$\langle 1, 2, 3 \rangle \wedge_{a \in [1,3]} ((\langle + \emptyset \rangle \Diamond \neg \text{jail}_a) \vee \langle -a \rangle \Box \text{jail}_a) \quad (\text{K})$$

### 5.1.1 Syntax

A TCL formula  $\phi$  is constructed with the following three syntax rules.

$$\begin{aligned}
\phi &::= p \mid \neg\phi_1 \mid \phi_1 \vee \phi_2 \mid \langle A \rangle \psi \\
\psi &::= \phi \mid \eta \mid \psi_1 \vee \psi_2 \mid \psi_1 \wedge \psi_2 \mid \langle +A \rangle \psi_1 \mid \langle +A \rangle \bigcirc \psi_1 \mid \langle +A \rangle \eta_1 \mathbf{U} \psi_1 \mid \langle +A \rangle \psi_1 \mathbf{R} \eta_1 \\
&\quad \mid \langle -A \rangle \psi_1 \mid \langle -A \rangle \bigcirc \psi_1 \mid \langle -A \rangle \eta_1 \mathbf{U} \psi_1 \mid \langle -A \rangle \psi_1 \mathbf{R} \eta_1 \\
\eta &::= \phi \mid \eta_1 \vee \eta_2 \mid \eta_1 \wedge \eta_2 \mid \langle + \rangle \bigcirc \eta_1 \mid \langle + \rangle \eta_1 \mathbf{U} \eta_2 \mid \langle + \rangle \eta_1 \mathbf{R} \eta_2 \\
&\quad \mid \langle -A \rangle \bigcirc \eta_1 \mid \langle -A \rangle \eta_1 \mathbf{U} \eta_2 \mid \langle -A \rangle \eta_1 \mathbf{R} \eta_2
\end{aligned}$$

Here,  $p$  is an atomic proposition in  $\mathcal{P}$  and  $A \subseteq \{1, \dots, m\}$  is an agency. Property  $\langle A \rangle \psi_1$  is an (existential) strategy quantification (SQ) specifying that there exist strategies of the agents in  $A$  that make all plays consistent with these strategies satisfy  $\psi_1$ . Property  $\langle +A \rangle \psi_1$  is an (existential) strategy interaction quantification (SIQ) and can only occur bound by an SQ. Intuitively,  $\langle +A \rangle \psi_1$  means that there exist strategies of the agents in  $A$  that work with the strategies introduced by the ancestor formulas. Likewise,  $\langle -A \rangle$  indicates a revocation of the strategy binding for the agents in  $A$ .  $\langle + \rangle$  is an abbreviation for  $\langle +\emptyset \rangle$  or, equivalently  $\langle -\emptyset \rangle$ . Thus, it neither binds nor revokes the binding of the strategy of any agent. Yet, it provides a temporalisation in that it provides a tree formula that can be interpreted at a particular point.

‘U’ is the *until* operator. The property  $\psi_1 \mathbf{U} \psi_2$  specifies a play along which  $\psi_1$  is true until  $\psi_2$  becomes true. Moreover, along the play,  $\psi_2$  must eventually be fulfilled. ‘R’ is the *release* operator. Property  $\psi_1 \mathbf{R} \psi_2$  specifies a play along which either  $\psi_2$  is always true or  $\psi_2 \mathbf{U} (\psi_1 \wedge \psi_2)$  is satisfied. (Release is dual to until:  $\neg(\phi_1 \mathbf{U} \phi_2) \Leftrightarrow \neg\phi_2 \mathbf{R} \neg\phi_1$ .)

In the following we may use  $\langle ?A \rangle \psi$  to conveniently denote an SQ or SIQ formula with ‘?’ is empty, ‘+’, or ‘-’. An SIQ  $\langle \pm A \rangle \psi$  is called non-trivial if  $A$  is not empty, and trivial otherwise.

Formulas  $\phi$  are called *TCL formulas*, *sentences*, or *state formulas*. Formulas  $\psi$  and  $\eta$  are called *tree formulas*. Note that we strictly require that non-trivial strategy interaction cannot cross path modal operators. This restriction is important because it offers a sufficient level of locality to efficiently model-check a system against a TCL property. To illustrate this and to provide a simple extension that offers more expressive power to the

cost of a much higher complexity, we informally discuss a small extension, *extended TCL* (ETCL), where the production rule of  $\psi$  also contains  $\neg\psi$  and show that it can be used to encode ATL\*, and the realisability problem of prenex QPTL can be reduced to ETCL model-checking.

For convenience, we also have the following shorthand notations.

$$\begin{array}{ll}
\text{true} \equiv p \vee (\neg p) & \text{false} \equiv \neg \text{true} \\
\phi_1 \wedge \phi_2 \equiv \neg((\neg\phi_1) \vee (\neg\phi_2)) & \phi_1 \Rightarrow \phi_2 \equiv (\neg\phi_1) \vee \phi_2 \\
\Diamond\phi_1 \equiv \text{true} \mathbf{U}\phi_1 & \Box\phi_1 \equiv \text{false} \mathbf{R}\phi_1 \\
\neg \bigcirc \phi_1 \equiv \bigcirc \neg\phi_1 & \langle A \rangle \bigcirc \psi_1 \equiv \langle A \rangle \langle + \rangle \bigcirc \psi_1 \\
\langle A \rangle \psi_1 \mathbf{U} \psi_2 \equiv \langle A \rangle \langle + \rangle \psi_1 \mathbf{U} \psi_2 & \langle A \rangle \psi_1 \mathbf{R} \psi_2 \equiv \langle A \rangle \langle + \rangle \psi_1 \mathbf{R} \psi_2
\end{array}$$

In general, it would also be nice to have the universal SQs and SIQs as duals of existential SQs and SIQs, respectively. Couldn't we add, or encode by pushing negations to state formulas, a property of the form  $[+A]\psi_1$ , meaning that, for all strategies of agency  $A$ ,  $\psi_1$  will be fulfilled? In principle, this is indeed no problem, and extending the semantics would be simple. This logic would be equivalent to allowing for negations in the production rule of  $\psi$ . The problem with this logic is that it is too succinct. We will briefly discuss in the following section that model-checking becomes non-elementary if we allow for such negations.

From now on, we assume that we are always in the context of a given TCL sentence.

### 5.1.2 semantics

In order to prepare the definition of a semantics for TCL formulas, we start with the definition of a semantics for sentences of the form  $\langle A \rangle \psi$ , where  $\psi$  does not contain any SQs. We call these formulas *primitive* TCL formulas.

Due to the design of TCL, strategy bindings can only effectively happen at non-trivial SQs  $\langle A \rangle$  and when a non-trivial SIQ  $\langle +B \rangle$  is interpreted. To ease referring to these strategies, we first define the *bound agency* of a subformulas  $\phi$  of a TCL sentence  $\chi$ , denoted  $\text{bnd}(\phi)$ , as follows.

- For state formulas  $\phi$ ,  $\text{bnd}(\phi) = \emptyset$ .



- For state formulas  $\langle A \rangle \psi$ ,  $\text{bnd}(\psi) = A$  (unless  $\psi$  is a state formula).
- For tree formulas  $\psi_1 = \langle +A \rangle \psi_2$ ,  $\text{bnd}(\psi_2) = \text{bnd}(\psi_1) \cup A$ .
- For tree formulas  $\psi_1 = \langle -A \rangle \psi_2$ ,  $\text{bnd}(\psi_2) = \text{bnd}(\psi_1) \setminus A$ .
- For all other tree formulas  $\psi_1$  or  $\psi_2$  with  $\psi = \psi_1 \text{OP} \psi_2$ , with  $\text{OP} \in \{\wedge, \vee, \mathcal{U}, \mathcal{R}\}$ , we have  $\text{bnd}(\psi_1) = \text{bnd}(\psi)$  or  $\text{bnd}(\psi_2) = \text{bnd}(\psi)$ , respectively.

$\text{bnd}$  shows, which agents have strategies assigned to them by an SIQ or SQ. Note that this leaves the  $\text{bnd}$  undefined for all state formulas not in the scope of an SQ formulas. For completeness, we could define  $\text{bnd}$  as empty in these cases, but a definition will not be required in the definition of the semantics.

As the introduction of additional strategies through non-trivial SIQ  $\langle +B \rangle$  is governed by a *positive* Boolean combination, all strategy selections can be performed concurrently. Such a design leads us to the concept of strategy schemes.

A *strategy scheme*  $\sigma$  is the set of strategies introduced by any non-trivial SQ  $\langle A \rangle$  or SIQ  $\langle +A \rangle$ . By abuse of notation, we use  $\sigma[\phi, a]$  to identify such a strategy. Read in this way,  $\sigma$  can be viewed as a partial function from subformulas and their bound agencies to strategies. Thus,  $\sigma[\phi, a]$  is defined if  $a \in \text{bnd}(\phi)$  is in the bound agency of  $\phi$ .

For example, given a strategy scheme  $\sigma$  for a TCL sentence  $\langle 1 \rangle \Diamond((\langle +2 \rangle \bigcirc p) \wedge \langle 2 \rangle \Box q)$ , the strategy used in  $\sigma$  by Agent 1 to enforce the whole formula can be referred to by

$$\sigma[\langle 1 \rangle \Diamond((\langle +2 \rangle \bigcirc p) \wedge \langle 2 \rangle \Box q), 1],$$

but also by  $\sigma[\langle +2 \rangle \bigcirc p, 1]$ , while  $\sigma[\langle 2 \rangle \Box q, 1]$  is undefined.

We use a simple tree semantics for TCL formulas. A (computation) tree  $T_r$  is obtained by unravelling  $\mathcal{G}$  from  $r$  and expand the ownership and labelling functions from  $\mathcal{G}$  to  $T_r$  in the natural way. Technically, we have the following definition.

**Definition: Computation Tree.** A *computation tree* for a turn based game  $\mathcal{G}$  from a state  $q$ , denoted  $T_q$ , is the smallest set of play prefixes that contains  $q$  and, for all  $\pi \in T$  and  $(\text{last}(\pi), q') \in \mathcal{E}$ ,  $\pi q' \in T$ .  $\square$

The *strategy-pruned* tree for a tree node  $\pi$ , a strategy scheme  $\sigma$ , and a subformula  $\psi_1$

of  $\chi$  from a state  $q$ , in symbols  $T_q\langle\pi, \sigma, \psi_1\rangle$ , is the smallest subset of  $T_q$  such that:

- $\pi \in T_q\langle\pi, \sigma, \psi_1\rangle$ ;
- for all  $\pi' \in T_q\langle\pi, \sigma, \psi_1\rangle$  with  $\omega((last(\pi')) \notin bnd(\psi_1)$  and  $(last(\pi'), q') \in \mathcal{E}$ ,  $\pi'q' \in T_q\langle\pi, \sigma, \psi_1\rangle$ ;
- for all  $\pi' \in T_q\langle\pi, \sigma, \psi_1\rangle$ ,  $a = \omega((last(\pi'))$ , and  $q' = \sigma[\psi_1, a](\pi')$  with  $a \in bnd(\psi_1)$ ,  $\pi'q' \in T_q\langle\pi, \sigma, \psi_1\rangle$ .

Given a computation tree or a strategy-pruned tree  $T$  and a node  $\pi \in T$ , for every  $\pi q \in T$ , we say that  $\pi q$  is a successor of  $\pi$  in  $T$ . A play  $\rho$  is a *limit of  $T$*  (or an infinite path in  $T$ ), in symbols  $\rho \overset{\infty}{\in} T$ , if there are infinitely many prefixes of  $\rho$  in  $T$ .

We now define the semantics of subformulas of primitive TCL formulas inductively as follows. Given the computation tree  $T_q$  of  $\mathcal{G}$ , a tree node  $\pi \in T_q$ , and a strategy scheme  $\sigma$ , we write  $T_q, \pi, \sigma \models \psi_1$  to denote that  $T_q$  satisfies  $\psi_1$  at node  $\pi$  with strategy scheme  $\sigma$ .

While the notation might seem heavy on first glance, note that the truth for state formulas merely depends on the state  $last(\pi)$  in which they are interpreted, and the tree formulas are simply interpreted on a strategy pruned tree rooted in  $\pi$  and defined by the strategy scheme.

- For state formulas  $\phi$  other than SQ formulas, we use the state formula semantics:

$T_q, \pi, \sigma \models \phi$  iff  $\mathcal{G}, last(\pi) \models \phi$ , with the usual definition.

- $\mathcal{G}, q \models p$  if, and only if,  $p \in \lambda(q)$ ,
- $\mathcal{G}, q \models \neg\phi$  if, and only if,  $\mathcal{G}, q \not\models \phi$ ,
- $\mathcal{G}, q \models \phi_1 \vee \phi_2$  if, and only if,  $\mathcal{G}, q \models \phi_1$  or  $\mathcal{G}, q \models \phi_2$ , and
- $\mathcal{G}, q \models \phi_1 \wedge \phi_2$  if, and only if,  $\mathcal{G}, q \models \phi_1$  and  $\mathcal{G}, q \models \phi_2$ .

(Note that this allows for using negation for state formulas.)

- $T_q, \pi, \sigma \models \psi_1 \vee \psi_2$  iff  $T_q, \pi, \sigma \models \psi_1$  or  $T_q, \pi, \sigma \models \psi_2$ . (The  $\psi_i$  are no state formulas.)
- $T_q, \pi, \sigma \models \psi_1 \wedge \psi_2$  iff  $T_q, \pi, \sigma \models \psi_1$  and  $T_q, \pi, \sigma \models \psi_2$  hold.
- $T_q, \pi, \sigma \models \langle \pm A \rangle \bigcirc \psi$  iff, for all successors  $\pi q'$  of  $\pi$  in  $T_q\langle\pi, \sigma, \langle \pm A \rangle \bigcirc \psi_1\rangle$ ,  $T_q, \pi q', \sigma \models \psi$  holds.
- $T_q, \pi, \sigma \models \langle \pm A \rangle \psi_1 \mathbf{U} \psi_2$  iff, for all limits  $\rho \overset{\infty}{\in} T_q\langle\pi, \sigma, \langle \pm A \rangle \psi_1 \mathbf{U} \psi_2\rangle$ , there is a  $k \geq |\pi| - 1$  such that  $T_q, \rho[0, k], \sigma \models \psi_2$  and, for all  $h \in [|\pi| - 1, k - 1]$ ,  $T_q, \rho[0, h], \sigma \models \psi_1$

hold.

- $T_q, \pi, \sigma \models \langle \pm A \rangle \psi_1 R \psi_2$  iff, for all limits  $\rho \in T_q \langle \pi, \sigma, \langle \pm A \rangle \psi_1 R \psi_2 \rangle$ , one of the following two restrictions are satisfied.
  - For all  $k \geq |\pi| - 1$ ,  $T_q, \rho[0, k], \sigma \models \psi_2$ .
  - There is a  $k \geq |\pi| - 1$  such that  $T_q, \rho[0, k], \sigma \models \psi_1 \wedge \psi_2$ , and, for all  $h \in [|\pi| - 1, k]$ ,  $T_q, \rho[0, h], \sigma \models \psi_2$ .
- $T_q, \pi, \sigma \models \langle \pm A \rangle \psi_1$  iff  $T_q, \pi, \sigma \models \psi_1$ .
- $\mathcal{G}, q \models \langle A \rangle \psi_1$  iff there is a strategy scheme  $\sigma$  such that  $T_q, q, \sigma \models \psi_1$ .

If  $\phi_1$  is a TCL sentence then we write  $\mathcal{G} \models \phi_1$  for  $\mathcal{G}, r \models \phi_1$ .

Note that, while asking for the existence of a strategy scheme refers to all strategies introduced by some SQ or SIQ in the TCL sentence, only the strategies introduced by the respective SQ and the SIQs in its scope are relevant.

The simplicity of the semantics is owed to the fact that it suffices to introduce new strategies at the points where eventualities become true for the first time. Thus, they do not really depend on the position in which they are invoked and we can guess them up-front. (Or, similarly, together with the points on the unravelling where they are invoked.) This is possible, simply because the validity of state formulas (and hence of TCL sentences) cannot depend on the validity of the left hand side of an until (or the right hand side of a release) *after* the first time it has been satisfied.

## 5.2 Expressive Power of TCL

Note that TCL is not a superclass of BSIL since BSIL allows for negation in front of SIQs while TCL does not. However, by examining the proofs in [76] for the inexpressibility of BSIL properties by  $ATL^*$ , GL, and AMC, we find that the BSIL sentence used in the proofs is also a TCL sentence. This leads to the conclusion that there are properties expressible in TCL but cannot be expressed in  $ATL^*$ , GL, and AMC.

**Lemma 25.** There are TCL sentences that cannot be expressed in any of  $ATL^*$ , GL, or AMC. □

TCL is, in fact, not only a powerful logic, but also contains important logics either as syntactical fragments or can embed them in a straight forward way. ATL and CTL can be viewed as syntactic fragments of TCL.

But it is also simple to embed LTL and even CTL\*. We start with  $\exists\text{LTL}$ , the less used variant where one is content if one path satisfies the formula. We then translate an LTL formula, which we assume w.l.o.g. to be in negative normal form (negations only in front of atomic propositions). Then “there is a path that satisfies  $\phi$ ” is equivalent to  $\langle 1, \dots, m \rangle \widehat{\phi}$ , where  $\widehat{\phi}$  is derived from  $\phi$  by replacing every occurrence of  $\bigcirc$ ,  $\text{U}$ , and  $\text{R}$  by  $\langle + \rangle \bigcirc$ ,  $\langle + \rangle \text{U}$ , and  $\langle + \rangle \text{R}$ , respectively. The simple translation is possible because the formula  $\widehat{\psi}$  is de-facto interpreted over a path, the path formed by the joint strategy of the agency  $[1, m]$ . The  $\langle + \rangle$  operators we have added have no effect on the semantics in such a case, just as a CTL formula can be interpreted as the LTL formula obtained by deleting all path quantifiers when interpreted over a word.

Consequently, we have the expected semantics for  $\forall\text{LTL}$ : “all paths satisfy  $\phi$ ” is equivalent to  $\neg \langle A \rangle \widehat{\neg \phi}$ , where  $\neg \phi$  is assumed to be re-written in negative normal form. The encoding of  $\exists\text{LTL}$  and  $\forall\text{LTL}$  can easily be extended to the encoding of CTL\*.

**Lemma 26.** TCL is more expressive than CTL\* and LTL. □

This encoding does not extend to ATL\*.  $\langle 1 \rangle ((\Box p) \vee \Box q)$  is an ATL\* property that cannot be expressed with TCL.

This is different from the ATL property  $(\langle 1 \rangle \Box p) \vee \langle 1 \rangle \Box q$  or the TCL property  $\langle 1 \rangle ((\langle + \rangle \Box p) \vee \langle + \rangle \Box q)$ . In fact, the proofs and examples in [76] can also be applied in this work to show that there are properties of ATL\* (or GL, or AMC) that cannot be expressed with TCL. This leads to the following lemma.

**Lemma 27.** TCL is incomparable in expressiveness with ATL\*, GL, and AMC. □

Note, however, that allowing for a negation in the definition of  $\psi$  would change the situation. Then an ATL\* formula  $\langle A \rangle \psi$  (assuming for the sake of simplicity that  $\psi$  is an LTL formula), would become  $\langle A \rangle \neg \langle +[1, m] \setminus A \rangle \widehat{\neg \psi}$  in the extended version of TCL. The translation extends to full ATL\*, but this example also demonstrates why negation is banned: even without nesting, we can, by encoding ATL\*, encode a 2EXPTIME complete




Figure 5.1: The turn-based game graph from the non-elementary hardness proof of extended TCL.

model-checking problem, losing the appealing tractability of our logic.

In fact, it is easy to reduce the realisability problem of prenex QPTL, and hence a non-elementary problem, to the model-checking problem of extended TCL. Using the game structure from Figure 5.1, we can encode the realisability of a prenex QPTL formula with  $n - 1$  variables, for simplicity of the form  $\forall p_2 \exists p_3 \forall p_4 \dots \exists p_n \phi$ , where  $p_2, \dots, p_n$  are all propositions occurring in  $\phi$ . We reduce this to model-checking the formula

$$\phi' = \langle 1 \rangle \neg \langle +2 \rangle \neg \langle +3 \rangle \neg \langle +4 \rangle \neg \dots \neg \langle +n \rangle (\psi_\phi \wedge \langle + \rangle \Box p_1),$$

where  $\psi_\phi$  can be obtained from  $\hat{\phi}$  by replacing

- every literal  $p_i$  by  $\langle -1 \rangle \langle +1 \rangle \bigcirc (p_i \wedge \langle + \rangle \bigcirc p_i)$ , and
- every literal  $\neg p_i$  by  $\langle -1 \rangle \langle +1 \rangle \bigcirc (p_i \wedge \langle + \rangle \bigcirc \neg p_i)$ .

These formulas are technically not extended TCL formulas as  $\langle +i \rangle \psi_1$  is not part of the production rule of  $\psi$ , but  $\langle +i \rangle \psi_1$  can be used as an abbreviation for  $\langle +i \rangle \text{false} \mathbf{U} \psi_1$ .

Checking satisfiability of  $\phi$  is equivalent to model-checking  $\phi'$  on the game shown in Figure 5.1. The game has  $n + 1$  nodes, agents, and atomic propositions. The nodes in Figure 5.1 are labeled with the agent that owned the nodes, and the atomic proposition  $p_i$  is true exactly in node  $i$ . From his state, Agent 1 can move to any other state, while all other agents can either stay in their state or return to the state owned by Agent 1.

The game starts in the node owned by Agent 1, and in order to comply with the specification, the outermost strategy profile chosen by Agent 1 must be to stay in the initial state for ever.  $\psi_\phi$  is chosen to align the truth of  $p_i$  at position  $j \in \mathbb{N}$  with the decision that Agent  $i$  makes on the history  $1^j i$ : *true* corresponds to staying in  $i$  and *false* with returning to 1.

It is not hard to establish a matching upper bound for model-checking extended TCL.




Figure 5.2: The turn-based game graph from the EXPTIME hardness proof.

### 5.3 Complexity of TCL

In this section, we show that model-checking TCL formulas is EXPTIME-complete in the formula and P-complete in the model (and for fixed formulas), while the satisfiability problem is 2EXPTIME-complete. As the proof of inclusion of the satisfiability problem in 2EXPTIME builds on the proof of the inclusion of model-checking in EXPTIME, we start with an outline of the EXPTIME hardness argument for the TCL model-checking problem and then continue with describing EXPTIME and 2EXPTIME decision procedures for the TCL model and satisfiability checking problem, respectively. 2EXPTIME hardness for TCL satisfiability is implied by the inclusion of CTL\* as a de-facto sub-language [77].

We show EXPTIME hardness by a reduction from the PEEK- $G_6$  [78] game. An instance of PEEK- $G_6$  consists of two disjoint sets of boolean variables,  $P_1 = \{p_1, \dots, p_h\}$  (owned by a safety agent) and  $P_2 = \{p_{h+1}, \dots, p_{h+k}\}$  (owned by a reachability agent), a subset  $I \subseteq P_1 \cup P_2$  of them that are initially *true*, and a boolean formula  $\gamma$  in CNF over  $P_1 \cup P_2$  that the reachability agent wants to become *true* eventually. The game is played in turns between the safety and the reachability agent (say, with the safety agent moving first), and each player can change the truth value of one of his or her variables in his/her turn.

**Lemma 28.** TCL model-checking is EXPTIME hard for primitive TCL formulas.

*Proof.* To reduce determining the winner of an instance of a PEEK- $G_6$  game to TCL model-checking, we introduce a 2-agent game  $\mathcal{G} = \langle 2, \mathcal{Q}, r, \omega, \mathcal{P}, \lambda, \mathcal{E} \rangle$  as shown in Figure 5.2, where Agent 1 (he, for convenience) represents the safety agent while Agent 2 (she, for convenience) represents the reachability agent.  $t_{h+k}$  and  $f_{h+k}$  are the only states owned by Agent 2.

The game is played in rounds, and a round starts each time the game is at state  $r$ . If the game goes through  $t_i$  this is identified with the variable  $p_i$  to be true. Likewise, going through  $f_i$  is identified with the variable being false.

It is simple to write a TCL specification that forces the safety player to toggle the value of exactly one of his variables in each round, and to toggle the value of the variable  $p_{h+i}$  of the reachability player defined by the state  $i$  she has previously moved to, while maintaining all other variable values. Requiring additionally that the safety agent can guarantee that the boolean formula is never satisfied provides the reduction.  $\square$   $\square$

The details of the construction are available in the full version. It is interesting that a game with only two agents suffices for the proof. Two agents are also sufficient to show P hardness for fixed formulas, as solving a reachability problem for AND-OR graphs [21] naturally reduces to showing  $\langle 1 \rangle \Diamond p$ .

**Lemma 29.** TCL model-checking for fixed formulas is P hard for primitive TCL formulas.

$\square$

In order to establish inclusion in EXPTIME and P, respectively, we use an automata based argument.

**Theorem 30.** The model-checking problem of TCL formulas against turn-based game graphs is EXPTIME-complete, and P-complete for fixed formulas.

*Proof.* We first show the claim for primitive TCL formulas  $\phi = \langle A \rangle \psi$ .

To keep the proof simple, we first consider a tree automaton  $\mathcal{U}$  that checks the acceptance of  $\psi$  for a given strategy scheme  $\sigma$ . That is,  $\mathcal{U}$  checks if  $T_q^+, q, \sigma \models \psi$  under the assumption that both  $\sigma$  and the truth values for the subformulas starting with a  $\langle \pm B \rangle$  are encoded in the nodes of  $T_q^+$ .

Such an automaton would merely have to run simple consistency checks, and it is simple to construct a suitable universal weak tree automaton  $\mathcal{U}$ , which is polynomial in the size of  $\phi$ . From there it is simple to infer a deterministic Büchi tree automaton  $\mathcal{D}$ , which is exponential in the weak universal tree automaton [79].

It is then a trivial step (projection) to *guess*  $\sigma$  and the truth annotation of the subformulas on the fly, turning the deterministic Büchi tree automaton  $\mathcal{D}$  that requires a correct annotation into a nondeterministic Büchi automaton  $\mathcal{N}$  of the same size that checks  $\mathcal{G}, q \models \phi$ . Acceptance can be checked in time quadratic in the size of the product of  $\mathcal{N}$  and  $\mathcal{G}$  [80].

To take the step to full TCL, we can model-check the truth of primitive TCL formulas and then use the result of this model-checking instead of the respective subformula.

Hardness is inherited from Lemmata 28 and 29. □ □

This argument shows more: the complexity of TCL model-checking for fixed formulas does not depend on the formula. It suffices to solve a number of Büchi games, where both the size of the game and the number of games to be played is linear in  $\mathcal{G}$ .

**Corollary 31.** Viewing the size of a TCL sentence as a parameter, TCL model-checking is fixed parameter tractable.

The automata construction from the proof of Theorem 30 extends to a construction for satisfiability checking.



# Chapter 6

## Software Resilience against Dense Errors

### 6.1 Resilience in a Nutshell

From Example 2 to 4 in the previous section, it is easy to see the common paradigm of error recovery in software systems.

When errors are detected, a recovery mechanism will be activated to avoid failures and try to get back to normal execution.

Moreover, such a recovery mechanism usually needs to operate under the assumption that more errors may also happen during the recovery process. In practice, system designers have already implemented many defensive modules, e.g., exception handlers, which are certainly good candidates for the recovery segments. Thus, the recovery scheme we discuss is likely to have arisen in an ad-hoc fashion as a natural concept when software architects and programmers designed recovery mechanisms for critical software.

The vast state spaces of critical systems make an automated support for and a solid foundation of evaluating design alternatives particularly valuable.

In the following, we will use the examples from the previous subsection as a motivation for defining a new game, called *safety resilience game*, between the recovery mechanism (the protagonist) and the error-injecting agent (the antagonist). The game is specified




Figure 6.1: Illustration of the recovery operation

with a set  $F$  of failure states, a set  $S$  of safe states (the safety region), the moves by the antagonist to inject errors, and the resilience level  $k$  that the designers want to achieve. The objective of the protagonist is to identify a control strategy so that the whole system can achieve the prescribed level (or the highest level)  $k$  of resilience for safety region  $S$  (a set of states) and failure state set  $F$ .

The game is played round by round. When the antagonist issues an error move, the play may be deflected into a recovery segment. If there are no more than  $k - 1$  errors in the recovery segment, then a  $k$ -resilient control mechanism must direct the recovery segment to end at a safe state. The above observation suggests that a safety region can be abstracted as a fixed point to the recovery procedure that transforms a safe state to another safe state via the recovery segment with at most  $k - 1$  errors. Conceptually, a fixed point to a procedure  $f(x)$  is a set  $S$  of elements in the domain of  $x$  such that  $S = \{f(x) \mid x \in S\}$ . To calculate the fixed point of the recovery procedure, we can use the greatest fixed point algorithm. The idea is to start from a superset of the recovery procedure fixed point. For convenience, we call a superset of the fixed point a *pseudo fixed point (PFP)*. Then we iteratively check every state  $q$  in the PFP and eliminate  $q$  from the PFP if, after at most  $k$  errors from  $q$ , the recovery mechanism either cannot avoid failure or cannot direct the system back to the PFP. As the iterative checking and elimination goes on, the PFP will shrink and eventually stabilize. Note that its size is always finite, since the initial PFP must be no bigger than  $Q$ . The final PFP is then a greatest fixed point to the recovery mechanism for  $k$ -resilience and is the legitimate safety region.

This recovery procedure can be illustrated as in Figure 6.1 for resilience to 2 errors. In this figure, the states in set  $S'$  are computed as the precondition of states in  $S$  through

those transitions in the figure. Each path from  $S'$  to  $S$  is a recovery segment.  $S$  and  $S'$  may overlap. The blue circles represent states in the recovery segments. If we calculate  $S'$  out of  $S$ , then, for each state  $q' \in S'$ , we can find a path from  $q' \in S'$  via a path in the recovery segment to another state  $q \in S$ . The maximal number of errors in a recovery segment is 2. Thus the protagonist has a strategy to recover from errors in  $S'$  to  $S$  even when 2 errors happen in the corresponding recovery segment. When  $S' = S$ , then  $S$  is a fixed point to the precondition operator through the recovery segments in the figure.

Now we formally define the concept that we explained with Figure 6.1.

**Definition 2. ( $k$ -safety):** Given a  $k \in \mathbb{N}$ , a state  $q$  is called  $k$ -safe with respect to a safety region  $S \subseteq Q \setminus F$  of non-failure states, denoted  $q \in \text{sfrch}_k(S)$ , if there is a strategy for the protagonist to guarantee that we can reach back to  $S$  from  $q$ , provided that the overall count of errors is at most  $k$ .  $\square$

However, the definition can be subtle in its interpretation. Specifically, the ability to stand against one wave of  $k$  errors is not the same as that against repeated recovery from waves of  $k$  errors. If the recovery mechanism is not designed properly, the system may gradually lose a bit of control after each wave of  $k$  errors and eventually degrade to system-level failure.

**Example 32. (Fault-tolerant computer architectures):** Consider Example 2 with  $2k+1$  processor copies, with the objective to maintain majority checks and to identify the bad processors. Indeed, according to the first, naïve solution, any safe state with a recovery strategy to  $Q \setminus F$  is good. After  $k$  processor copies fail, the majority checks are still capable to maintain the correctness of the combined behavior to follow the design of the original system. There seems to be nothing to do after  $k$  errors. Thus, naïvely, we can choose those states as the safety region if, at those states, majority checks still work.

However, there is no expectation that the system will be able to recover at any point in the future into a situation where it can bear another wave of  $k$  errors. It will fail and lose the function of majority checks just after one more error. In contrast, in this work, we aim to propose a dense error resilience criterion that given no more errors for enough time to allow recovery, the system will eventually recover to resilience to  $k$  dense errors again.  $\square$




Figure 6.2: An example for calculating  $\text{sfrch}_k$

To look at this issue in more detail, please consider the transition system with four states, including a single failure state (state 4, marked by a double line) shown in Figure 6.2. The controlled transitions are depicted as black solid arrows, the error transitions are depicted as red dashed arrows. For  $S = Q \setminus F = \{1, 2, 3\}$ , all states in  $S$  are in  $\text{sfrch}_0(S)$ . For all  $k \geq 1$ , we have  $\text{sfrch}_k(S) = \{1, 2\}$ : the protagonist can simply stay in  $\{1, 2\}$  during the safety phase of the game, and once the antagonist plays an error transition, the game progresses into the recovery segment, where the protagonist's objective is satisfied immediately. This outlines the difference between  $k$ -sfrch-ty and the linear time property of being able to repeatedly tolerate waves of up to  $k$  errors, which would only be satisfied by states 1 and 2 for  $k = 1$ , and only for state 1 for  $k = 2$ .

This difference raises the question if the rules of our game are depriving the antagonist of some of the  $k$  errors that she should intuitively be allowed to insert in a wave. The answer is that this is not the case if we use any fixed point of  $\text{sfrch}_k$  as  $S$ . In this case, the protagonist would regain the capability to endure a wave of  $k$  errors when reaching a safe state after recovery. Instead of depriving the antagonist, one could say that we reset the number of errors in any recovery segment that the antagonist can inject to  $k$ . Thus such a fixed point of  $\text{sfrch}_k$  should consist of states, from which we can use a control mechanism to fend off repetitive waves of  $k$  dense errors in the recovery segments. For convenience, we call states in such a fixed point of  $\text{sfrch}_k$  the  $k$ -resilient states.

For a state to be in  $\text{sfrch}_k(S)$ , the system (protagonist) has a strategy to recover to  $S$ , given that a long enough execution commenced without another round of  $k$  errors happening. We say that two successive errors are in the same *group of dense errors* if the sequence of states separating them was not long enough for recovery to the safety region. Vice versa, if two successive errors are far enough apart such that the protagonist can guarantee recovery in this separation, then they do not belong to the same group.

To check whether recovering to  $S$  by the protagonist (the fault-tolerance mechanism) is always possible, provided that at most  $k$  errors occurred during a recovery segment, observe that nesting  $\text{sfrch}_k$  once, i.e.,  $\text{sfrch}_k(\text{sfrch}_k(\cdot))$ , corresponds to tolerating up to two rounds of up to  $k$  dense errors, and so forth. Thus, for  $S$  to be a target of recovery for  $k$ -resilience,  $S$  must be a fixed point of the operator  $\text{sfrch}_k$  from Definition 2, or, equivalently,  $S = \text{sfrch}_k(S)$  must hold. Moreover, if  $S$  is the greatest fixed point to  $k$ -resilience, then we can apply  $\text{sfrch}_k(\cdot)$  any number of times to  $S$  and still obtain  $S$ . Computationally, the greatest fixed point of  $\text{sfrch}_k$  can be constructed as by executing

$$\text{sfrch}_k(\text{sfrch}_k(\text{sfrch}_k(\dots \text{sfrch}_k(S) \dots))),$$

using a sufficiently deep nesting that a fixed point is reached.

Note that this fixed point  $x$  to  $x = \text{sfrch}_k(x)$  is what we are really interested in, while  $\text{sfrch}_k(S)$  for a given  $S$  is an intermediate result that does not guarantee survival of the systems after waves of dense errors. If this greatest fixed point

$$R = \bigcup \{X \subseteq S \mid X = \text{sfrch}_k(X)\}$$

is non-empty, the protagonist's strategy for the fixed point (guaranteeing eventual recovery to a state in the fixed point within no more than  $k$  errors, i.e.,  $k$ -resilience) can be used to control the recovery mechanism, constraining its transitions to follow its winning strategy.

As explained in the introduction, there can be several natural control problems in our safety resilience game. First, the system designers may want to know whether the chosen safety region  $S$  can be supported by the recovery mechanism for resilience level  $k$ . Second, they may want to get design support for choosing the safety region for achieving resilience level  $k$ . Finally, they may want to know the maximal resilience level that they can achieve.

we will explain the algorithm for constructing  $\text{sfrch}_k(\cdot)$  and evaluating  $k$ -resilient states.

## 6.2 Safety resilience games

A system is  $k$ -resilient if it can be controlled to tolerate infinitely many groups of up to  $k$  dense errors, provided that the system is given enough time to recover between

these groups. As we have explained, in systems developed with defensive mechanism against errors, when errors are detected, recovery procedures should be activated. The major challenge is to decide given a set of failure states and a safety region, whether the recovery mechanism can support a resilience level required by the users. Our goal is to develop techniques with a solid foundation to assist the system designers in evaluating the resilience of their systems, to synthesize the controller strategy for the required resilience level, and to achieve the maximal resilience level.

We now formally define the safety resilience game played between a system (the protagonist) and an error-injector (the antagonist). Initially, the two players are given a 2-player concurrent game structure  $\mathcal{K}$ , a pebble in  $r$ , a set  $F \subseteq Q$  of failure states, and a safety region  $S \subseteq Q \setminus F$ . Then the recovery region consists of states in  $Q \setminus (F \cup S)$ . The two players together make decisions and move the pebble from state to state. The antagonist tries to deflect a play into  $F$  by injecting sufficiently many errors, while the protagonist tries to avoid that the pebble reaches  $F$ . To achieve this, the protagonist can use the recovery region as the safety buffer and try to get back to  $S$  as soon as the play is deflected from  $S$  to the recovery region. If a system is resilient to  $k$  errors, then it means that the protagonist can handle up to  $k - 1$  errors while in the recovery region. Thus when checking whether a system is resilient to  $k$  errors, we only need to check those recovery segments with no more than  $k - 1$  errors.

In the following, we formalize the concept.

**Definition 3. (Safety resilience game structure):** Such a structure is a pair  $\langle \mathcal{K}, F \rangle$  with the following restrictions.

- $\mathcal{K}$  is a 2-player concurrent game structure  $\langle Q, r, P, \lambda, E_1, E_2, \delta \rangle$ . Conceptually, the first player represents the system / the protagonist, while the second player represents the error model / the antagonist.
- $E_2$  is partitioned into error and non-error moves  $E_{error}$  and  $E_{noerr}$ , respectively. We require that only the 2nd player can issue *error* moves. Moreover,  $E_{noerr}$  must be non-empty.
- $F$  is the set of failure states in  $Q$  with  $r \notin F$ .

The antagonist can choose if she wants to respond on a move of the protagonist with an error move. We allow for different non-error moves to reflect ‘normal’ nondeterministic behavior, e.g., caused by abstraction. We allow for different error moves to reflect different errors that can occur in the same step.

We sometimes refer to transitions with *error* moves by the antagonist as *error transitions* and to transitions with *noerr* moves by the antagonist as *controlled transitions*.

For a party  $A \subseteq \{1, 2\}$ , we refer with  $\bar{A} = \{1, 2\} \setminus A$  to the players not in the party, and by  $E_A$  to the moves made by the players in  $A$ , that is,  $E_{\{1,2\}} = E_1 \times E_2$ ,  $E_{\{1\}} = E_1$ , etc.

The antagonist can use both error and non-error moves to influence the game. In a simple setting, the antagonist may only have the choice to insert error-moves, while there is only a single controlled transition. In this simple case, the protagonist can choose the successor state alone unless the antagonist plays an error transition. Specifically, a safety resilience game structure is *simple* if  $E_2$  contains only one error move. Considering simple safety resilience game structures leads to lower complexities, as it changes reductions from reachability in games (PTIME-complete [21]) to reachability in graphs (NL-complete [22]).  $\square$

Note that, in the game structure, only one system player and one error model player are allowed. This is purely for the simplicity of algorithm presentation. With proper reduction techniques, we can easily convert a game structure with more than one system player and more than one error model player to the structure in Definition 3. The standard technique would be using the transition rules of the product automata of the system players for the protagonist while using the transition rules of the product automata of the error model players for the antagonist. In fact, we indeed use this reduction technique in our experiment for analyzing the resilience levels of multi-agent systems.

From now on, we assume that we are in the context of a given safety resilience game structure  $\mathcal{G} = \langle \mathcal{K}, F \rangle$ .

**Definition 4. (Recovery segments):** We need to rigorously define *recovery segments*. A play prefix  $\rho$  is a *recovery segment* to safety region  $S \subseteq Q \setminus F$  if it satisfies the following

constraints.

- $\rho(0) \in S$ .
- If  $|\rho| = \infty$ , then all states in  $\rho[1, \infty)$  are in  $Q \setminus (S \cup F)$ . In this case,  $\rho$  is called a failed recovery segment.
- If  $|\rho| \neq \infty$ , then all states in  $\rho[1, |\rho| - 2]$  are in  $Q \setminus (S \cup F)$  and  $last(\rho) = \rho(|\rho| - 1)$  is either in  $F$  or  $S$ . If  $last(\rho) \in F$ ,  $\rho$  is also a failed recovery segment; otherwise, it is a successful one.

We use  $level(\rho, S)$  to denote the number of error moves between states in  $\rho$  with respect to the safety region  $S$ :  $level(\rho, S) \stackrel{\text{def}}{=} |\{i \in [0, |\rho| - 1) \mid \rho_e(i) \models E_{error}\}|$ .  $\square$

As stated in the introduction, we propose a game-theoretic foundation for resilience analysis of software systems. With this perspective, the protagonist acts as a maximizer, who wants to maximize the resilience levels along all plays. For this, the protagonist fixes a strategy that describe what he is going to do on each play prefix. The antagonist acts as a minimizer, who wants to minimize the resilience level. She can resolve nondeterminism and inject errors in order to achieve this, and (although this plays no major role in this setting) she knows the strategy the protagonist has fixed and can use this knowledge in principle.

The goal of the protagonist is therefore the same as the goal of the system designer: to obtain a strategy that offers a maximal level of resilience in a safety game. However, in order to avoid degenerate behavior where the protagonist benefits from being in the recovery phase and from the antagonist therefore being allowed less errors in the current wave of errors she may inject, we have to strengthen his obligation to eventually recover to the safe states when the environment chooses not to inject further errors. This way, the protagonist has no incentive to cycle in the recovery region. Consequently, he can recover to the safe region within  $|Q|$  moves after the antagonist has inserted the last error of the current wave, irrespective of whether the antagonist would be allowed to insert further errors in this wave. This is the key reason why memoryless optimal control exists for this error model, why it is reasonable to assume swift recovery, and, consequently, why it is a posteriori justified to leave the separation time between two waves implicit: the time to



traverse  $|Q|$  states suffices.

Besides obtaining this from intuition, we can also consider the tree of successful recoveries for any protagonist strategy that can endure  $k$  error moves by the antagonist. The tree of recoveries from up to  $k$  errors is finite according to the definition of successful recovery segments. Then for any subtree  $t$  in this tree of recoveries with a node  $v$  in  $t$  such that  $v$  is labeled with the same state as the root of  $t$  with no error on the path, we can always replace  $t$  with the subtree rooted at  $v$ . After the replacement, we have a tree of recoveries with no greater depth than the original one. After repeating such replacements, this immediately provides a translation from such a strategy with unrestricted memory to one with memory of size  $k$  (the resilience level). The restriction to memoryless strategies follows from the construction we give in Section ??, which does not depend on the memory and still yields a strategy, which is memoryless. Thus, in this work, we should define the resilience level of software systems based on *memoryless* protagonist strategies.

Based on the argument above, the gain of the protagonist in a play can be defined as follows.

**Definition 5. (Gain):** Given safety region  $S \subseteq Q \setminus F$ , the gain of a play  $\rho$  to  $S$ , in symbols  $gain(\rho, S)$ , denotes the maximal integer  $k \in \mathbb{N}$  such that, for all recovery segments  $\rho_r$  to  $S$  in  $\rho$ , if  $level(\rho_r, S) \leq k$ , then  $\rho_r$  is a successful recovery segment to  $S$ .  $\square$

The resilience level of a safety resilience game is defined as the maximum gain that the protagonist can guarantee in all plays with a memoryless strategy.

**Definition 6. (Safety resilience game):** Such a game is zero-sum and defined on a safety resilience game structure  $\mathcal{G} = \langle \mathcal{K}, F \rangle$  and a safety region  $S \subseteq Q \setminus F$ . The gain of  $\mathcal{G}$  to  $S$ , in symbols  $gain(\mathcal{G}, S)$ , is defined as the maximum gain that the protagonist can manage with memoryless strategies. Rigorously,

$$gain(\mathcal{G}, S) \stackrel{\text{def}}{=} \max_{\sigma \in \Sigma^{(0)}} \min_{\sigma' \in \Sigma} gain(play(r, \sigma, \sigma'), S)$$

Please be recall that  $play(r, \sigma, \sigma')$  is the play from  $r$  according to strategies  $\sigma$  and  $\sigma'$  respectively of the two players. Moreover  $\Sigma^{(0)}$  is the set of memoryless strategies.

We say that the resilience level of  $\mathcal{G}$  to  $S$  is  $\text{gain}(\mathcal{G}, S)$ . A strategy  $\omega$  for the protagonist is optimal to  $S$  if  $\min_{\sigma' \in \Sigma} \text{gain}(\text{play}(r, \omega, \sigma'), S) = \max_{\sigma \in \Sigma^{(0)}} \min_{\sigma' \in \Sigma} \text{gain}(\text{play}(r, \sigma, \sigma'), S)$ . When  $S$  is not given, we say that  $\mathcal{G}$  is  $k$ -resilient if there exists a non-empty  $S \subseteq Q \setminus F$  with  $\text{gain}(\mathcal{G}, S) \geq k$ .  $\square$

**Remark.** While the option of using memoryless strategies plays a minor role in the technical argument, it plays a paramount role in the usefulness of the resulting control strategy: choosing memoryless strategies implies that all recovery segments are short. In particular, all sub-paths (recovery segments) between two waves of dense errors injected by the antagonist are shorter—and usually significantly shorter—than the size of  $\mathcal{G}$ . In consequence, any time span long enough for traversing the recovery segment will lead to a full recovery. It is therefore sufficient for a temporal distance we have to assume between two waves of dense errors.

### 6.3 Alternating-time $\mu$ -calculus with events

We propose to solve our resilience game problems with an existing technology, i.e., model-checking of alternating-time  $\mu$ -calculus (AMC) formulas. AMC is a propositional temporal logic with fixed point operators. For example, the following formula

$$\mu X.(\text{safe} \vee \langle 1 \rangle \bigcirc X) \tag{A}$$

uses least fixed point operator  $\mu$  to declare a fixed point variable  $X$  for a set of states. Subformula  $\langle 1 \rangle \bigcirc \phi$  existentially quantifies over the protagonist strategies that can direct the plays to a successor state satisfying  $\phi$ . Together, the formula specifies a set  $X$  of states that can inductively reach a safe state with the control of the protagonist. Specifically, the formula says that a state is in  $X$  if either it is *safe* or the protagonist can direct to a successor state known to be in  $X$ . For our game structures, we only need strategy quantification of up to two players.

However, we need extend AMC with some simple syntax sugar. There are two extensions. The first is for Boolean combinations of path modalities in the scope of strategy quantification. For example, the following AMCE formula

$$\langle 1 \rangle ((\text{smoke} \Rightarrow \bigcirc \text{alarmOn}) \vee \bigcirc \text{windowClosed}) \quad (\text{B})$$

says that the protagonist can enforce either of the following two path properties with the same strategy.

- If there is smoke, then the alarm will be turned on in the next state.
- The window will always be closed in the next state.

Such a formula is not in ATL and AMC [3].

The second extension is for restricting transitions that may participate in the evaluation of path formulas. The restriction is via constraints on moves on transitions and can, in our extension to AMC, be specified with a move symbol set to the next-state modal operators. For example, the following AMCE formula

$$\langle 1 \rangle ((\bigcirc^{2:\text{error}} \text{alarmOn}) \wedge (\bigcirc^{\neg 2:\text{error}} \neg \text{alarmOn})) \quad (\text{C})$$

says that the protagonist can

- turn on the alarm when an error occurs; and
- keep the alarm silent when no error occurs.

Before we formally present AMCE, we need define expressions for constraints on moves of players in transitions. We adapt an idea from [23]. Specifically, a *move expression*  $\eta$  is of the following syntax.

$$\eta ::= a : e \mid \eta_1 \vee \eta_2 \mid \neg \eta_1$$

Here,  $a$  is a player index in  $\{1, 2\}$  and  $e$  is a move symbol in  $E_1 \cup E_2$ .  $\vee$  and  $\neg$  are standard disjunction and negation. Typical shorthands of Boolean operations can also be defined out of  $\vee$  and  $\neg$ . A total move vector can be expressed as  $[e_1, e_2]$  where for all  $a \in \{1, 2\}$ ,  $e_a \in E_a$  is the move by player  $a$  specified in the vector. We say  $[e_1, e_2]$  satisfies  $\eta$ , in symbols  $[e_1, e_2] \models \eta$ , if and only if the following constraints are satisfied.

- $[e_1, e_2] \models a : e$  if, and only if,  $e_a$  is  $e$ .
- $[e_1, e_2] \models \eta_1 \vee \eta_2$  if, and only if,  $[e_1, e_2] \models \eta_1$  or  $[e_1, e_2] \models \eta_2$ .
- $[e_1, e_2] \models \neg \eta_1$  if, and only if,  $[e_1, e_2] \not\models \eta_1$ .

### 6.3.1 Syntax

A formula  $\phi$  in AMCE has the following syntax.

$$\begin{aligned}\phi &::= p \mid X \mid \phi_1 \vee \phi_2 \mid \neg \phi_1 \mid \mu X. \phi_1 \mid \langle A \rangle \psi \\ \psi &::= \mid \psi_1 \vee \psi_2 \mid \neg \psi_1 \mid \bigcirc^\eta \phi_1\end{aligned}$$

Here,  $\phi$  is a state formula,  $\psi$  is a path formula,  $p$  is an atomic proposition symbol in  $P$  (atomic proposition set, as in Definition ??), and  $X$  is a set variable for subsets of  $Q$ . The Boolean connectors are the common ones:  $\vee$  for disjunction and  $\neg$  for negation. Note that we allow for Boolean combinations of the next operators  $\bigcirc$  under strategy quantification  $\langle A \rangle$ . This is one major difference of AMCE from AMC.

Formula  $\mu X. \phi_1$  is the usual least fixed point operation to  $\phi_1$ . According to the tradition in [3], we require that all free occurrences of  $X$  in  $\phi_1$  must occur within an even number of scopes of negations. This is because sentences with a negative occurrence, like  $\mu X. \neg X$ , have no natural semantics. A set variable  $X$  is *bound* in a formula  $\phi$  if it is inside a declaration scope of  $X$ . If it is not bound, then it is *free*. An AMCE sentence is an AMCE state formula without free set variables. In most cases, we are interested in specifications given as AMCE sentences.

The  $A$  in  $\langle A \rangle$  is a finite set of player indices in  $[1, 2]$ . Conceptually,  $\langle A \rangle \psi$  means that players in  $A$  can collaborate to make  $\psi$  true. For example,  $\langle \{1, 2\} \rangle \bigcirc p$  means that players 1 and 2 can collaborate to make  $p$  true in the next state. We follow the notations in [3] and omit the parentheses in formulas like  $\langle A \rangle \psi$ . For example,  $\langle \{2\} \rangle \bigcirc p$  and  $\langle \{1, 2\} \rangle \bigcirc p$  will be abbreviated as  $\langle 2 \rangle \bigcirc p$  and  $\langle 1, 2 \rangle \bigcirc p$  respectively.

We allow event restrictions as superscripts in  $\bigcirc^\eta \phi_1$  with a move expression  $\eta$ . The operator is important in supporting the evaluation of safety resilience levels with traditional model-checking technology. Note that since AMC [3] only allows for the next-state temporal modality, only the choice of moves to the next states of a strategy matters. Formula  $\bigcirc^\eta \phi_1$  is thus evaluated at states with respect to move vectors satisfying constraint  $\eta$ . The formula is true of a move vector  $[e_1, e_2]$  if and only if  $[e_1, e_2] \models \eta$  implies the satisfaction of  $\phi$  at state  $\delta(q, e_1, e_2)$ . Also  $\bigcirc^{1:E_1} \phi_1$  can be written as  $\bigcirc \phi_1$  in AMC [3] and the superscript to  $\bigcirc$  can be omitted.

We also adopt shorthands in the below. The  $\beta$  refers to state or path formulas.

$$\begin{aligned}
true &\stackrel{\text{def}}{=} p \vee \neg p \\
false &\stackrel{\text{def}}{=} \neg p \wedge p \\
\beta_1 \wedge \beta_2 &\stackrel{\text{def}}{=} \neg((\neg\beta_1) \vee (\neg\beta_2)) \\
\beta_1 \Rightarrow \beta_2 &\stackrel{\text{def}}{=} (\neg\beta_1) \vee \beta_2 \\
\nu X.\phi &\stackrel{\text{def}}{=} \neg\mu X.\neg\phi \\
[A]\psi &\stackrel{\text{def}}{=} \neg\langle A\rangle\neg\psi
\end{aligned}$$

### 6.3.2 semantics

In the following, we adapt the presentation style of [3] to define the semantics of AMCE inductively over the structure of the subformulas. The value of a state formula at a state is determined by the interpretation of the set variables. Such an interpretation  $I$  maps set variables to subsets of  $Q$ . In comparison, the value of a path formula at a state is determined by both the interpretation of the set variables and the move vector chosen by the players. For convenience and conciseness of presentation, we extend the definition of interpretation of [3] also to record the chosen move vector by some players. Specifically, we use an auxiliary variable “move” for the present chosen move vector in the evaluation of path formulas. Given an interpretation  $I$ ,  $I(\text{move})$  records the chosen move vector of all players in  $I$ . For example,  $I(\text{move}) = [\text{setAlarm}, \perp]$  means the chosen move vector that player 1 sets on an alarm while player 2 does nothing under interpretation  $I$ .

We need the following concept for collaborative choices of moves to the next states by some players. An *enforced move vector set* by  $A \subseteq [1, 2]$  is a maximal set of move vectors that agree on the choices of moves by players with indices in  $A$ . Specifically, given an enforced move vector set  $C$  by  $A$ , we require that, for every  $[e_1, e_2] \in C$ ,  $[e'_1, e'_2] \in C$ , and  $a \in A$ ,  $e_a = e'_a$ . For convenience, we let  $\Gamma^A$  denote the set of all enforced move sets by  $A$ .

Following the semantics style of [3], we can extend  $I$  to be an interpretation of all state and path formulas. Intuitively, given a state or path formula  $\beta$ ,  $I(\beta)$  is the set of states that satisfy  $\beta$  according to the assumption on values of set variable values and auxiliary

variable “move.” More precisely,  $I(\beta)$  is a subset of  $Q$  that satisfies the following inductive rules.

- $I(p) = \{q \mid p \in \lambda(q)\}$ .
- $I(\beta_1 \vee \beta_2) = I(\beta_1) \cup I(\beta_2)$ .
- $I(\neg\beta_1) = Q - I(\beta_1)$ .
- $I(\mu X.\phi_1)$  is the smallest set  $Y \subseteq Q$  with  $Y = I[X \mapsto Y](\phi_1)$ , where  $I[X \mapsto Y]$  is a new interpretation identical to  $I$  except that  $X$  is interpreted as  $Y$ .
- $I(\langle A \rangle \psi)$  is the set of states such that there is an enforced move vector set  $C$  by  $A$  such that, for all move vectors  $\epsilon \in C$ ,  $I[\text{move} \mapsto \epsilon](\psi)$  holds:

$$I(\langle A \rangle \psi) = \bigcup_{C \in \Gamma^A} \bigcap_{\epsilon \in C} I[\text{move} \mapsto \epsilon](\psi)$$

- Given  $I(\text{move}) = [e_1, e_2]$ , if  $[e_1, e_2] \models \eta$ , then  $I(\bigcirc^n \phi_1) = \{q \in Q \mid \delta(q, e_1, e_2) \in I(\phi_1)\}$ ; otherwise  $I(\bigcirc^n \phi_1) = Q$ .

A concurrent game structure is a model of an AMCE sentence  $\phi$ , if its initial state  $r$  is in the interpretation of  $\phi$  ( $r \in I(\phi)$ ) for any interpretation  $I$ .

Note that, strictly speaking, AMCE does not add much to the expressiveness of AMC. In the literature, propositions have often been used to record events. Intuitively, we would need one atomic proposition for each event to mark that it has just occurred. This event marker would be true exactly at states right after the event happened. (One would possibly have to create multiple copies of states to reflect this.)

As discussed in [24], such a modeling technique leads to an unnecessary blow up of the state space, which could be exponential in the number of players in general concurrent games. By properly selecting the transitions with respect to operators like  $\bigcirc^n$ , such auxiliary propositions are not necessary when encoding the state space. Thus, AMCE can also be of interest to practitioners for the efficient analysis and verification of general concurrent games.

## 6.4 Resilience level checking algorithm

In Subsection ??, we have proposed the idea of the  $\text{sfrch}_k(\cdot)$  operator and proposed to use its greatest fixed point for the evaluation of  $k$ -resilience. In the following, we first establish some properties of  $k$ -safety and then use AMC model-checking technology to solve the safety resilience games.

### 6.4.1 High-level description of the algorithm

The following lemma shows the sufficiency of  $k$ -safety as a building block for solving safety resilience games.

**Lemma 33.** For a safety resilience game  $\mathcal{G}$ ,  $\text{sfrch}_k(\cdot)$  has a greatest fixed point.

**Proof :** The lemma follows from the facts that the function  $\text{sfrch}_k$  is monotonic ( $S \subseteq S'$  implies  $\text{sfrch}_k(S) \subseteq \text{sfrch}_k(S')$  because a winning strategy for the protagonist for  $S$  is also a winning strategy for  $S'$  for all states in  $\text{sfrch}_k(S)$ ) and operates on a finite domain.  $\square$

For the example in Figure 6.2, considering  $S = \{1\}$  ( $\{1\} = \text{sfrch}_2(\{1, 2, 3\})$ ), the only state in  $S$ , state 1, is 2-resilient: it can recover with the recovery strategy to always go to the left.

The set of  $k$ -resilient states of  $\mathcal{G}$ , can be calculated as the greatest solution to  $S = \text{sfrch}_k(S)$  with  $S \subseteq Q \setminus F$ . Technically we can start the inductive calculation of the greatest fixed point from base case  $S_0 = Q \setminus F$ , and successively calculate  $S_{i+1} = \text{sfrch}_k(S_i)$ , for each  $i \geq 0$ . The set of  $k$ -resilient states is then the limit  $S_\infty$ . As soon as we have  $S_{i+1} = S_i$ , a fixed point is reached. We then have  $S_i = S_\infty$  and can stop the inductive construction. Since  $S_0$  is finite and  $S_{i+1} \subseteq S_i$  holds for all  $i \geq 0$ , we will eventually reach a  $j$  with  $S_{j+1} = S_j = S_\infty$ .

### 6.4.2 Realization with AMCE model-checking

We need formally define the interaction among strategies of players. We borrow the notation of function composition. Given two partial functions  $\beta_1$  and  $\beta_2$ , we use  $\beta_1 \circ \beta_2$  to represent their composition. Specifically, we have the following definition.

$$\beta_1 \circ \beta_2(a) = \begin{cases} \beta_1(a) & \text{if } \beta_2(a) \text{ is undefined.} \\ \beta_2(a) & \text{otherwise} \end{cases}$$

For our purpose, a partial strategy vector is a mapping from  $\{1, 2\}$  to  $\Sigma$  and can be undefined for some players in  $\{1, 2\}$ . It is for a party  $A \subseteq \{1, 2\}$  if it is defined only for players in  $A$  and represents a collaborative strategy of the players with a defined strategy in  $A$ . It is total if it is defined for all players.

For convenience, we also define partial move vectors as mappings from  $\{1, 2\}$  to  $E$ . A partial move vector is for a party  $A \subseteq \{1, 2\}$  if it is defined only for players in  $A$ . It is total if it is defined for all players in  $\{1, 2\}$ . Given two partial move vectors  $\gamma_1$  and  $\gamma_2$ , we define  $\gamma_1 \circ \gamma_2$  to represent the composition of the two vectors.

Given an  $S$ , we propose to construct  $\text{sfrch}_k(S)$  in an induction on  $k$ . We need the following preliminary concepts for the presentation.

**Definition 7. (Traps)** For  $A \subseteq \{1, 2\}$ , a *trap* for  $A$  is a subset  $Q' \subseteq Q$  that party  $\{1, 2\} \setminus A$  has a strategy vector  $\beta$  to keep all plays from leaving  $Q'$ . Formally, we require that, for every  $q \in Q'$  and partial move vector  $\gamma$  for  $A$ , there exists a partial move vector  $\gamma'$  for  $\{1, 2\} \setminus A$  such that  $\delta(q, \gamma \circ \gamma'(1), \dots, \gamma \circ \gamma'(m)) \in Q'$ .  $\square$

### Base case, $\text{sfrch}_0(S)$

In the base case,  $\text{sfrch}_0(S)$  characterizes those states, from which the protagonist can direct the plays to  $S$  and stay there via a protagonist strategy when there is no error injected by the antagonist. Thus  $\text{sfrch}_0(S)$  is the greatest trap for the antagonist to  $S$  when no error happens and the greatest solution to the following equation.

$$X = \left\{ q \left| \begin{array}{l} q \in X \cap S, e \in E_1, \\ \forall e' \in E_2 (e' \neq \text{noerr} \Rightarrow \delta(q, e, e') \in X) \end{array} \right. \right\}.$$

In AMCE, we can alternatively define  $\text{sfrch}_0(S)$  as follows.

$$\text{sfrch}_0(S) \stackrel{\text{def}}{=} \nu X. (S \wedge \langle 1 \rangle \bigcirc^{\neg 2:\text{error}} X).$$

This is the usual safety kernel of  $S$ , which consists of those states, from which any controlled transition is safe. It can be computed by the usual greatest fixed point construction.



Table 6.1: Algorithm for  $\text{sfrch}_0(S)$  by iterative elimination

---

```

 $\text{sfrch}_0(S)$ 
1: for  $q \in S$  do  $c_q = |E_1|$  end for
2: for  $q \in S, e \in E_1$  do  $b_{[q,e]} = \text{true}$  end for
3: Let  $i = 0$  and  $K_0 = Q - S$ .
4: while  $K_i \neq \emptyset$  do
5:   Let  $K_{i+1} = \emptyset$ .
6:   for  $q \in K_i$  and  $[q', e, f] \in L_q$  do
7:     if  $b_{[q',e]}$  is true then
8:       Let  $c_{q'} = c_{q'} - 1$ .
9:       if  $c_{q'}$  is 0 then add  $q'$  to  $K_{i+1}$ . end if
10:    end if
11:    Set  $b_{[q',e]}$  to false.
12:   end for
13:   Increment  $i$  by 1.
14: end while
15: return  $S - (K_0 \cup \dots \cup K_i)$ .

```

---

**Lemma 34.**  $\text{sfrch}_0(S)$  can be constructed, together with a suitable memoriless control strategy, in time linear to the size of  $\mathcal{G}$ .

**Proof :** A state  $q \in S$  can stay in  $\text{sfrch}_0(S)$  if there is a choice  $e \in E_1$  such that for all  $f \in E_2$ ,  $\delta(q, e, f) \in \text{sfrch}_0(S)$ . Basically, we can use the typical approach of iterative elimination to calculate  $\text{sfrch}_0(S)$ . That is, we first let  $K_0 = Q - S$ . Then we a sequence of mutually disjoint sets  $K_1, K_2, \dots, K_i, \dots$  such that for all  $i \geq 1$ , states in  $K_{i+1}$  can be shown to be not in  $\text{sfrch}_0(S)$  by evidences of states in  $K_i \cup \dots \cup K_0$ . Linear time can be achieved with careful book-keeping of the choices of moves at all states in  $S$ . We need a counter  $c_q$  for each  $q \in S$  initialized to  $|E_1|$  for the initial number of candidate choices of moves. Then for each  $[q, e] \in S \times E_1$ , we need a Boolean flag  $b_{[q,e]}$  initialized to *true* to represent that  $\{[e, f] \mid f \in E_2\}$  is still a valid choice of moves at  $q$  to satisfy  $\text{sfrch}_0(S)$ . For each state  $q$ , we also need to maintain a list of transition source states. That is, for each  $\delta(q', e, f) = q$ , we need record  $[q', e, f]$  in list  $L_q$ . Then the iterative elimination proceeds as the algorithm in table 6.1. The algorithm is linear time since each transition  $\delta(q, e, f)$  is checked exactly once. □

## Inductive cases, $\text{sfrchk}(S)$

Now we explain how to define the inductive cases of  $\text{sfrch}_k(S)$ . The condition is for those states from which plays can be directed to  $S$  via a recovery segment in  $Q \setminus (S \cup F)$  with  $k$  or less errors injected by the antagonist. An intermediate step for the construction of  $k$ - $\text{sfrch}$  states is the construction of an attractor that controls, through controlled moves, the play prefixes to stay in a subset  $L \subseteq Q \setminus F$  of non-failure states. As only controlled (non-error) moves are allowed, this is merely a backward reachability cone.

The *controlled limited attractor set* of a set  $X$  for a limited region  $L \subseteq Q$ , denoted  $\text{cone}_L(X)$  is the set from which there is a protagonist strategy to move to  $X$  without leaving  $L$  and errors injected by the antagonist. Technically,  $\text{cone}_L(X)$  is the least solution to equation:

$$Y = X \cup \left\{ q \left| \begin{array}{l} q \in L, e \in E_1, \\ \forall e' \in E_2 \setminus \{\text{error}\} (\delta(q, e, e') \in Y) \end{array} \right. \right\}.$$

The controlled limited attractor set  $\text{cone}_L(X)$  can be constructed using simple backward reachability for  $X$  of controlled transitions through states of  $L$ . In AMCE, this can be constructed as follows.

$$\text{cone}_L(X) \stackrel{\text{def}}{=} \mu Y. (X \vee (L \wedge \langle 1 \rangle \bigcirc^{\neg 2:\text{error}} Y))$$

Note that the protagonist must use the same move irrespective of the move of the antagonist to both stay in  $L$  and approach  $X$ , provided that the antagonist does not inject an error.

The controlled limited attractor set  $\text{cone}_L(X)$  is used in the construction of  $\text{sfrch}_k(S)$ . We further construct a descending chain  $V_0 \supseteq V_1 \supseteq \dots \supseteq V_{k-1}$  of limited attractors  $V_i$ . From  $V_i$  we have an attractor strategy towards  $S$  for the protagonist, which can tolerate up to  $i$  further errors. The respective  $V_i$  are attractors that avoid failure states. Moreover, from a state in  $V_i$  with  $i > 1$ , any error transition leads to  $V_{i-1}$ .

A state  $q \in Q$  is *fragile* for a set  $B \subseteq Q$  if, for all moves of the protagonist, at least one of its successors is outside of  $B$ . (The intuition is that this is an error move, and for simple safety resilience game structures, we can restrict the definition to failure states.) The set of fragile states for  $B$  is

$$\text{frag}(B) \stackrel{\text{def}}{=} \{q \mid \forall e \in E_1 \exists e' \in E_2 (\delta(q, e, e') \notin B)\}.$$

In AMCE, we have the following formulation of  $\text{frag}(B)$ .

$$\text{frag}(B) \stackrel{\text{def}}{=} [1] \bigcirc \neg B.$$

Technically, it is, however, easier to construct its dual

$$Q \setminus \text{frag}(B) = \langle 1 \rangle \bigcirc B.$$

This dual can be constructed using a controlled backward reachability to  $B$  with any strategy of the protagonist.

The limited regions  $L_i$  of states allowed when approaching  $S$  also form a descending chain  $L_0 \supseteq L_1 \supseteq \dots \supseteq L_k$ . Using these building blocks, we can compute the  $k$ -sfrch states as follows. The states in  $L_{i+1}$  are the non-failure states from which all error transitions lead to a state in  $V_i$ . The sets  $V_i$  contain the states from which there is a controlled path to  $S$  that progresses through  $L_i$ ; all error transitions originating from any state of this path lead to  $V_{i-1}$ .  $V_0$  is therefore just the set of states from which there is a controlled path to  $S$ .

From all states in  $V_{k-1}$ , the protagonist therefore has an optimal strategy in the recovery segment of the game described earlier: if the antagonist can play at most  $k-1$  errors, then the protagonist can make sure that  $S$  is reached.

Starting with  $L_0 \stackrel{\text{def}}{=} Q \setminus F$  that characterizes cones on the way to  $S$  without any errors, we define the  $V_k$ 's and  $L_k$ 's inductively by

$$L_k \stackrel{\text{def}}{=} L_0 \setminus \text{frag}(Q \setminus \text{cone}_{L_{k-1}}(S)),$$

In AMCE, this can be defined inductively as follows.

$$\begin{aligned} L_0 &\stackrel{\text{def}}{=} \neg F \\ L_k &\stackrel{\text{def}}{=} L_0 \wedge \langle 1 \rangle \bigcirc \text{cone}_{L_{k-1}}(S). \end{aligned}$$

Finally, we choose  $\text{sfrch}_k(S) \stackrel{\text{def}}{=} \text{sfrch}_0(S \cap L_k)$ . In AMCE, this can be expressed as follows.

$$\text{sfrch}_k(S) \stackrel{\text{def}}{=} \text{sfrch}_0(S \wedge L_k).$$

$$\begin{aligned}
L_0 &\stackrel{\text{def}}{=} \neg F \\
L_k &\stackrel{\text{def}}{=} \neg F \wedge \langle 1 \rangle \bigcirc \mu y. S \vee (L_{k-1} \wedge \langle 1 \rangle \bigcirc^{\text{error}} y \wedge \bigcirc L_{k-1}) \\
\text{sfrch}_0(S) &\stackrel{\text{def}}{=} \nu x. (S \wedge \langle 1 \rangle \bigcirc^{\text{error}} x) \\
\text{sfrch}_k(S) &\stackrel{\text{def}}{=} \text{sfrch}_0(S \wedge L_k) \\
\text{res}_k(\mathcal{G}) &\stackrel{\text{def}}{=} \nu S. ((Q \setminus F) \wedge \text{sfrch}_k(S)) : \text{the set of } k\text{-resilient states}
\end{aligned}$$

Table 6.2: Algorithm for  $k$ -resilient states

### Algorithm for the set of $k$ -resilient states

Finding a control strategy for  $k$ -sfrch control within  $\text{sfrch}_k(S)$  is simple: as long as we remain in  $\text{sfrch}_k(S) = \text{sfrch}_0(S \cap L_k)$ , we can choose any control move that does not leave  $\text{sfrch}_k(S)$ . Once  $\text{sfrch}_k(S)$  is left through an error transition to  $V_{k-1}, V_{k-2}, \dots$ , we determine the maximal  $i$  for which it holds that we are in  $V_i$  and follow the attractor strategy of  $\text{cone}_{L_i}(S)$  towards  $S$ .

In summary, we present our algorithms for the set of  $k$ -resilient states in Table 6.2. In fact, we have presented two algorithms. The first constructs  $\text{sfrch}_k(S)$ , which can be used for checking whether the safety region  $S$  provided by the users is indeed a good one. The way to do it is to simply check whether  $S$  is a solution to  $\text{sfrch}_k(x) = x$ .

Then our second algorithm calculates  $\text{res}_k(\mathcal{G})$  as the greatest fixed point  $S$  of  $\text{sfrch}_k(\cdot)$  as the recommendation for the safety region:

$$\text{res}_k(\mathcal{G}) = \bigcup \{S \subseteq Q \mid S = \text{sfrch}_k(S) \text{ and } S \cup F = \emptyset\}.$$

In this way, the users do not have to calculate and provide the safety region, which would be error prone. According to the argument and lemmas from above, we get the following theorem.

**Theorem 35.**  $\mathcal{G}$  is  $k$ -resilient if, and only if,  $r \in \text{res}_k(\mathcal{G})$ . □

### 6.4.3 Complexity

A rough complexity of our resilience level checking algorithm straightforwardly follows the complexity of AMC model-checking. Specifically, the following lemma explains

the maximal resilience level that we need consider. For convenience, let  $k_{\max}$  be the maximal resilience level of  $\mathcal{G}$ .

**Lemma 36.**  $k_{\max}$  is either infinite or no greater than  $|Q \setminus F|$ .

**Proof :** We assume that  $k_{\max}$  is greater than  $|Q \setminus F|$  but not infinite. This means that there exists a failed recovery segment  $\rho$  with  $k + 1$  errors injected by the antagonist. Since the protagonist can only use memoryless strategies, there must be two position indices  $i < j < |\rho| - 1$  with  $\rho(i) = \rho(j)$  in the recovery segment such that at  $\rho(i)$  and  $\rho(j)$ , the protagonist makes the same move while the antagonist makes different moves. This implies the existence of a shorter failed recovery segment  $\rho[0, i]\rho[j + 1, |\rho| - 1]$ . By repeating the above argument, we can eventually identify a failed recovery segment of length  $\leq |Q \setminus F|$  that contradicts the assumption and establishes the lemma.  $\square$

With Lemma 36, we can use the complexity of AMC model-checking problem [3] to straightforwardly establish the  $O(k_{\max}|E|)^2 = O(|Q \setminus F| \cdot |E|)^2$  complexity of  $\text{res}_k(\mathcal{G})$  when  $k$  is  $k_{\max}$ . In the following, we present a more detailed analysis of the complexity of our resilience level checking algorithm. All individual steps in the construction (intersection, difference, predecessor, and attractor) are linear in the size of the safety resilience game, and there are  $O(k)$  of these operations in the construction. This provides a bi-linear (linear in  $k$  and  $|\mathcal{G}|$ ) algorithm for the construction of  $\text{sfrch}_k$  and a strategy for the protagonist.

**Lemma 37.** A memoryless control strategy for the states in  $\text{sfrch}_k(S)$  can be constructed in time linear in both  $k$  and the size  $|\mathcal{G}|$  of the safety resilience game  $\mathcal{G}$ .  $\square$

The construction of  $\text{res}_k(\mathcal{G})$  uses the repeated execution of  $(Q \setminus F) \wedge \text{sfrch}_k(\cdot)$ . The execution of  $\text{sfrch}_k(\cdot)$  needs to be repeated at most  $|Q \setminus F|$  times until a fixed point is reached, and each execution requires at most  $O(k \cdot |\mathcal{G}|)$  steps by Lemma 37.

For the control strategy of the protagonist, we can simply use the control strategy from  $\text{sfrch}_k(S_\infty)$  from the fixed point  $S_\infty$ . This control strategy is memoryless (cf. Lemma 37).

**Lemma 38.**  $\text{res}_k(\mathcal{G})$  and a memoryless  $k$ -resilient control strategy for  $\text{res}_k(\mathcal{G})$  can be constructed in  $O(k \cdot |Q \setminus F| \cdot |\mathcal{G}|)$  time.  $\square$

Finding the resilience level  $k_{\max}$  for the initial state  $r$  requires at most  $O(\log k_{\max})$  many constructions of  $\text{res}_i(\mathcal{G})$ . We start with  $i = 1$ , double the parameter until  $k_{\max}$  is exceeded, and then use logarithmic search to find  $k_{\max}$ .

**Corollary 39.** For the initial state  $r$ , we can determine the resilience level  $k_{\max} = \max\{i \in \mathbb{N} \mid r \in \text{res}_i(Q \setminus F)\}$  of  $r$ ,  $\text{res}_{k_{\max}}(Q \setminus F)$ , and a memoryless  $k_{\max}$ -resilient control strategy for  $\text{res}_{k_{\max}}(Q \setminus F)$  in  $O(|Q \setminus F| \cdot |\mathcal{G}| \cdot k_{\max} \log k_{\max})$  time.  $\square$

**Simple safety resilience game structures.** For simple safety resilience game structures, checking if a state is in  $\text{sfrch}_0(S)$  is NL-complete.

**Lemma 40.** Testing if a state is in  $\text{sfrch}_0(S)$  is NL-complete.

*Proof.* NL completeness can be shown by reduction to and from the repeated ST-reachability [22] (the question whether there is a path from a state  $S$  to a state  $T$  and from  $T$  to itself in a directed graph).  $\square$

Likewise, the controlled limited attractor set  $\text{cone}_L(S)$  can be constructed using simple backwards reachability for  $G$  of controlled transition through states of  $L$ . For  $A = \text{cone}_L(S)$ , determining whether a state is in  $A$  is NL-complete (see [22]).

The complexity of determining whether or not a state  $q$  is in  $\text{sfrch}_k(S)$  thus depends on whether or not we consider  $k$  to be a fixed parameter. Considering  $k$  to be bounded (or fixed) is natural in our context, because  $k$  is bounded by the redundancy.

**Lemma 41.** For a fixed parameter  $k$ , testing if a state  $s$  of a simple safety resilience game structures is in  $\text{sfrch}_k(S)$  is NL-complete.

*Proof.* Testing if a state is in  $L_0$  is in NL. By an inductive argument, we can show that

- provided that testing if a state is in  $L_i$  is in NL, we can test if a state is in  $A_i = \text{cone}_{L_i}(S)$  by using the nondeterministic power to guess a path towards  $S$ , while verifying that we are in  $L_i$  in every state we pass before  $S$  is reached; and
- if we can check if a state is in  $A_i$  in NL, then we can check if it is in  $Q \setminus A_i$  [25], in  $\text{frag}(Q \setminus A_i)$  (with one nondeterministic transition), and in  $L_{i+1} = L_0 \setminus \text{frag}(S \setminus A_i)$  [25] in NL.

Testing that a state is in  $S \cap L_k$  is therefore in NL and testing if it is in  $\text{sfrch}_0(S \cap L_k)$  reduces to guessing a state  $t$  in  $\text{sfrch}_k(G)$  and an ST path (a path from  $s$  to  $t$  followed by a loop from  $t$  to  $t$ ), verifying for all states on the path that they are in  $S \cap L_k$ .

For hardness, note that the last step of the construction alone is NL-complete (Lemma 34).

□

If  $k$  is considered an input, then reachability in AND-OR graphs can easily be encoded in LOGSPACE: It suffices to use the nodes of an AND-OR graph as the states, the outgoing edges of OR nodes as the result of the choice of the protagonist only (while the move of the antagonist has no influence on the outcome, no matter whether or not she induces an error), and to model the AND nodes as a state, where the no-error move of the antagonist will lead in cycling in the state, while the antagonist can choose the successor from the graph when inducing an error. Choosing  $k$  to be the number of nodes of the AND-OR graph and  $F$  to be the target nodes of the AND-OR graph, the target nodes of the AND-OR graph are not reachable from a state  $s$  iff  $s \in \text{sfrch}_k(Q \setminus F)$ .

Given that reachability in AND-OR graphs is PTIME-complete [21], this provides:

**Lemma 42.** If  $k$  is considered an input parameter, then testing if a state  $s$  of a simple safety resilience game structures is in  $\text{sfrch}_k(S)$  is PTIME-complete. □

The complexity of  $\text{res}_k(S)$  is (almost) independent of the parameter  $k$ :

**Theorem 43.** The problem of checking whether or not a state  $s$  is  $k$ -resilient for a set  $S$  is PTIME-complete for all  $k > 0$  and NL-complete for  $k = 0$ .

*Proof.* We have shown inclusion in PTIME in Lemma 38. For hardness in the  $k > 0$  case, we can use the same reduction from the reachability problem in AND-OR graphs as for  $\text{sfrch}_k(S)$ .

For  $k = 0$ ,  $\text{sfrch}_0(G) = \text{sfrch}_0(\text{sfrch}_0(G))$  implies  $\text{res}_0(G) = \text{sfrch}_0(G)$ . The problem of checking if a state is in  $\text{res}_0(G)$  is therefore NL-complete by Lemma 40. □

**Hardness for general safety resilience game structures.** For general resilience game structures, we can again use a LOGSPACE reduction from the reachability in AND-OR

graphs: We again use the nodes of an AND-OR graph as the states, and the outgoing edges of OR nodes are selected based on the choice of protagonist only. For the AND nodes, we leave the choice to the antagonist only, without the need to invoke an error. (That is, errors play no role in this reduction. The antagonist may be allowed to insert one, but she can always obtain the same transition without doing so.)

Marking  $F$  as the target nodes, we get  $\text{res}_k(Q \setminus F) = \text{sfrch}_l(Q \setminus F)$  for all non-negative integers  $k, l$ , and  $s \in \text{sfrch}_0(Q \setminus F)$  iff the target nodes of the AND-OR graph are not reachable from  $s$ . With Lemmas 37 and 38, we get the following theorem.

**Theorem 44.** For all  $k \geq 0$ , the problems of checking whether or not a state  $s$  is in  $\text{res}_k(Q \setminus F)$  and  $\text{res}_k(Q \setminus F)$ , respectively, are PTIME-complete for general safety resilience game structures.



# Chapter 7

## Implementation Experiment

### 7.1 BSIL and TCL

#### 7.1.1 Implementation

We implemented semi-symbolic model-checkers of BSIL and TCL with **REDLIB** [23, 72–75] which is a free library for symbolic model-checking based on decision diagrams. **REDLIB** supports symbolic pre-condition and post-condition calculation of discrete transitions. Our model-checker starts from a symbolic representation of the initial condition. Then it repeatedly applies the post-condition procedure of **REDLIB** to explore the symbolic state representations in the computation tree. Our model-checker uses the result in section ?? to bound the exploration depth of tree.

#### 7.1.2 Benchmarks and their experiment report

Then we experimented with two parameterized benchmarks. The first is the prisoners' dilemma described in Example 1 with the number of prisoners as a parameter. Then we applied the model-checker to check whether the model satisfies formula (A) and (B) in the introduction. The time and memory usage of the model-checker is reported in Table 7.1.

The second benchmark is for campaign strategies of 2 political parties, Party 1 and Party 2, for seats in the congress. There are several parameters in our model, the number of seats in the congress, the amount of budget of each party in a round, and the number of

Table 7.1: Experiment data for the prisoners' dilemma model

#prisoners	Formula (A)		Formula (B)	
	Time	Mem	Time	Mem
2	0.50s	72M	0.51s	71M
3	0.92s	75M	0.88s	75M
4	1.23s	83M	1.14s	80M
5	3.19s	146M	2.90s	140M
6	6.71s	388M	6.52s	361M
7	25.10s	1043M	23.11s	979M

s: seconds (computation time);

M: megabytes (memory usage);

The experiment is conducted on a PC with Intel i7-2600k 3.4GHZ CPU and 8G RAM running Ubuntu 12.04.

candidates of each party in a round. The game is played by the chiefs of the parties, Chief 1 (a gentleman) and Chief 2 (a lady), who decide the amount of support that a candidate receive in a round. Candidates with more budget in a round will be elected in the round. If two or more candidates get the same amount of support, the winner will be decided randomly.

In our model, in each round, chief 1 first decides his budget allotment, then chief 2 decides her budget allotment, and then the election result is determined. For example, with 3 dollars in his budget for two candidates, the strategy of a Chief can be written as  $(x, y) \in \{(3, 0), (2, 1), (1, 2), (0, 3)\}$  where  $x$  and  $y$  respectively denote the support (in dollars) allotted to the first and the second candidates in his party. If Chief 1 uses strategy  $(2, 1)$  and Chief 2 uses strategy  $(0, 2)$  (Assume she has less budget.), the result of the round is that the two parties both win one seat in the round. On the other hand, if Chief 2 chooses strategy  $(1, 1)$  instead, then there is no strategy for Chief 1 to win more seats for his party.

We use the following propositions and predicates to write specifications. For each  $i \in [1, 2]$  and  $j \in [1, c]$  where  $c$  is the number of candidates,  $elect_{i,j}$  means the  $j$ 'th candidate of party  $i$  is elected now. Then we can use propositions  $elect_{i,j}$  to construct a predicate  $win_i$  that means party  $i$  wins more seats than party  $2-i$  in the round. Assume that the first candidate in each party is the chief. Then the following BSIL formula specifies that Chief 1 has a strategy to make sure candidate 1 to be elected and in the same time allow the opponent party to win more seats in every round.

$$\langle 1 \rangle (((\langle + \rangle \Diamond elect_{1,1}) \wedge (\langle +2 \rangle \Box win_2)) \quad (N)$$

Similarly, the following formula specifies that the chief of party 1 has a strategy to assure that his party can win more seats than his opponent party while allowing party two to assure the winning of a candidate.

$$\langle 1 \rangle (((\langle + \rangle \Diamond win_1) \wedge (\langle +2 \rangle \Diamond elected_{2,1})) \quad (O)$$

With different values of budget and candidate count, we may use our model-checker to analyze the strategy for achieving various goals. The time and memory usage data for checking formula (N) and (O) is in Table 7.2.

Table 7.2: Experiment Data Election

Parameters					Formula (N)			Formula (O)		
$c_1$	$b_1$	$c_2$	$b_2$	$s$	Result	Time	Mem	Result	Time	Mem
2	2	2	4	1	UNSAT	0.52s	61M	UNSAT	0.53s	61M
2	4	2	4	2	UNSAT	1.20s	75M	UNSAT	1.15s	75M
2	4	2	4	3	SAT	0.73s	75M	UNSAT	1.14s	75M
3	6	3	6	3	SAT	1.51s	211M	UNSAT	0.93s	211M
3	9	3	6	3	UNSAT	9.14s	837M	SAT	5.22s	681M
3	6	3	9	3	SAT	8.74s	502M	UNSAT	8.49s	502M
3	9	3	9	4	SAT	158s	6631M	UNSAT	93.17s	5082M

$c_1$ : # candidates of party 1

$b_1$ : total budget in a round of party 1

$c_2$ : # candidates of party 2

$b_2$ : total budget in a round of party 2

$s$ : # of seats in the congress

s: seconds (computation time)

M: megabytes (memory usage)

The experiment is conducted on a PC with  
Intel i7-2600k 3.4GHZ CPU and 8G RAM  
running Ubuntu 12.04.

The experiment shows that some possibilities of using our tool to flexibly support the analysis and synthesis of collaborating strategies among several agents. In the experiment with the prisoners' dilemma, the prisoners are modelled as individual agents. In the experiment with the election game, the chairpersons of the parties are modelled as individual agents and their candidates and respective allocated budgets are modelled as numbers. It would be interesting to see how we can explore the techniques in modelling and specifying concurrent games with our tool.

Table 7.3: Performance data of model-checking the TCL fragment

properties \ $m$	2	3	4	5	6	7	8	9	10
(I)	0.71s	0.94s	5.41s	66.3s	945s	>1000s			
	163M	165M	185M	350M	1307M				
(J)	0.50s	0.52s	0.61s	0.71s	1.11s	1.62s	5.77s	20.9s	68.1s
	163M	163M	164M	165M	168M	176M	214M	270M	376M
(K)	0.51s	0.51s	0.6s	0.82s	1.01s	1.81s	5.54s	18.2s	48.3s
	163M	163M	164M	165M	168M	176M	200M	241M	318M
(L)	0.5s	0.51s	0.57s	0.74s	1.01s	1.79s	7.41s	33.8s	141s
	163M	163M	164M	165M	168M	175M	232M	312M	430M
(M)	0.51s	0.66s	19.1s	>1000s					
	163M	164M	194M						
(P)	0.51s	0.53s	0.61s	0.71s	1.01s	1.70s	5.38s	15.2s	53.7s
	163M	163M	163M	165M	168M	175M	202M	243M	295M
(Q)	0.52s	0.52s	0.65s	0.72s	1.03s	1.85s	4.86s	16.1s	93.5s
	163M	163M	164M	165M	169M	177M	189M	208M	235M

s: seconds; M: megabytes.

The models are with 1 policeman and  $m$  prisoners. The experiment was carried out on an Intel i5 2.4G notebook with 2 cores and 4G memory, running ubuntu Linux version 11.10.

As in the case of BSIL, we use the parametrised models of the iterated prisoners' dilemma as our benchmarks to check the performance of our implementation on TCL model checker. We use seven benchmark formulas on these models in our experiments. The first five benchmarks are taken from the examples (I) through (M) from the introduction. Benchmarks (P) and (Q) are the following two properties, taken from [76].

- Property (P) specifies that all prisoners except Prisoner 1 can collaborate to release Prisoner 1 and let Prisoner 1 decide their fate.

$$\langle 2, \dots, m \rangle \left( (\langle + \rangle \Diamond \neg \text{jail}_1) \wedge \bigwedge_{i \in \{2, \dots, m\}} (\langle +1 \rangle \Diamond \neg \text{jail}_i) \wedge (\langle +1 \rangle \Box \text{jail}_i) \right) \quad (\text{P})$$

- Property (Q) specifies that Prisoner 1 has a strategy to put all other prisoners in jail while leaving her fate to them.

$$\langle 1 \rangle \left( (\bigwedge_{i \in \{2, \dots, m\}} \langle + \rangle \Box \text{jail}_i) \wedge (\langle 2, \dots, m \rangle \Diamond \neg \text{jail}_1) \wedge \langle 2, \dots, m \rangle \Box \text{jail}_1 \right) \quad (\text{Q})$$

For these benchmarks, we have collected the performance data for various parameter values in Table 7.4. For small models, the memory usage is dominated by the normal overhead, such as the representation of variable tables, state-transition tables, formula struc-

tures, etc. The data shows that our prototype can handle the various benchmarks, and scales well on five of the seven benchmarks. Ignoring the overhead, it also shows the exponential growth. The models, however, are growing exponentially, too. We assume that this growth is the main cause of the exponential growth of the response time.

## 7.2 Resilience

## 7.3 Tool implementation

In the following, we report our implementation and experiment about checking a system’s resilience to dense errors.

We adopt *CEFSM* (*communicating extended finite-state machine*) [26] as a convenient language for the description of abstract models of our concurrent game structures. A CEFSM consists of several finite-state machines extended with shared variables for the modeling of shared memory and with synchronizations for the modeling of message-passing in distributed systems. This is justifiable since the fault-tolerant algorithms may themselves be subject to restrictions in concurrent or distributed computation. Indeed, we found CEFSM very expressive in modeling the benchmarks from the literature [27, 28].

The translation from our CEFSMs to state transition systems, such as finite Kripke structures, is standard in the literature. All state spaces, conditions, preconditions, postconditions, fixed points, etc., are represented as logic formulas. The logic formulas are then implemented with multi-value decision diagrams (MDD) [29].

We then took advantage of the support of REDLIB for writing down template automatas for constructing complex models. We specified a template automata with REDLIB to describe the moves of the players. Conceptually, the player automatas are constructed as an instance of the template automata. Then the whole game structure is constructed as the product of all player automatas. Finally, we use the API of REDLIB to do on-the-fly construction of the game structure which can be advantageous since unreachable states will never be generated.

### 7.3.1 Benchmarks and their experiment report

We use the following five parameterized benchmarks to check the performance of our techniques. Each benchmark has parameters for the number of participating modules in the model. Such parameterized models come in handy for the evaluation of the scalability of our techniques with respect to concurrency and model sizes.

1. We use the example of a fault-tolerant computer architecture (Example 2) as our first benchmark. An important feature of this benchmark is that there is an assumed mechanism for detecting errors of the modules. Once an error is detected, a processor can be assigned to recover the module, albeit to the cost of a reduced redundancy in the executions.
2. Voting is a common technique for fault tolerance through replication when there is no mechanism to detect errors of the modules [17]. In its simplest form, a system can guarantee correctness, provided less than half of its modules are faulty. This benchmark implements this simple voting mechanism. Every time a voting is requested, the modules submit their ballots individually. Then we check how many module failures the system can endure and recover.
3. This is a simplified version of the previous voting benchmark, where we assume that there is a blackboard for the client to check the voting result.
4. *Practical Byzantine fault-tolerance (PBFT)* algorithm: We use an abstract model of the famous algorithm by Castro and Liskov [27]. It does not assume the availability of an error-detection mechanism but uses voting techniques to guarantee the correctness of computations when less than one third of the voters are faulty. This algorithm has impact on the design of many protocols [30–34] and is used in Bitcoin [35], a peer-to-peer digital currency system.
5. *Fault-tolerant clock synchronization algorithm*: Clock synchronization is a central issue in distributed computing. In [28], Ramanathan, Shin, and Butler presented several fault-tolerance clock synchronization algorithms in the presence of Byzan-

tine faults with high probability. We use a nondeterministic abstract model of the convergence averaging algorithm from their paper. The algorithm is proven correct when no more than one third of the local clocks can drift to eight time units from the median of all clock readings.

### 7.3.2 Modelling of the fault-tolerant systems

Appropriate modelling of the benchmarks is always important for the efficient verification of real-world target systems. Many unnecessary details can burden the verification algorithm and blow up the computation, while sketchy models can then give too many false alarms and miss correct benchmarks. We have found that there is an interesting aspect in the modelling of the above benchmarks. Replication and voting are commonly adopted techniques for achieving fault-tolerance and resilience. Such fault-tolerant algorithms usually consist of several identical modules that use the same behavior templates. This observation implies that the identity of individual modules can be unimportant for some benchmarks. For such benchmarks, we can use counter abstraction [36, 37] in their models. Specifically, with counter abstraction, we can model all system players with one player that keeps a counter  $c(l)$  for each control location  $l$  in the template automatas. Then at a state of the whole game graph,  $c(l)$  records the number of system players at location  $l$ . With this technique, a system with  $m - 1$  system player and one error model player is then reduced to two players: one counter-abstraction player for all the system players and one remaining error model player. If a system player enters a location  $l$  in a global transition, then in the model,  $c(l)$  is incremented by one in the abstract global transition. If a system player leaves  $l$  in the global transition, then  $c(l)$  is decremented by one in the abstract global transition. But the succession of location movements of a particular player is omitted from the abstraction.

We found that we can use counter abstraction to prove the correctness of benchmarks 1, 2, and 3. In contrast, the PBFT and the clock synchronization algorithms use counters for each module to model the responses received from its peer modules. As a result, we decided not to use counter abstraction to model these two algorithms in this work.

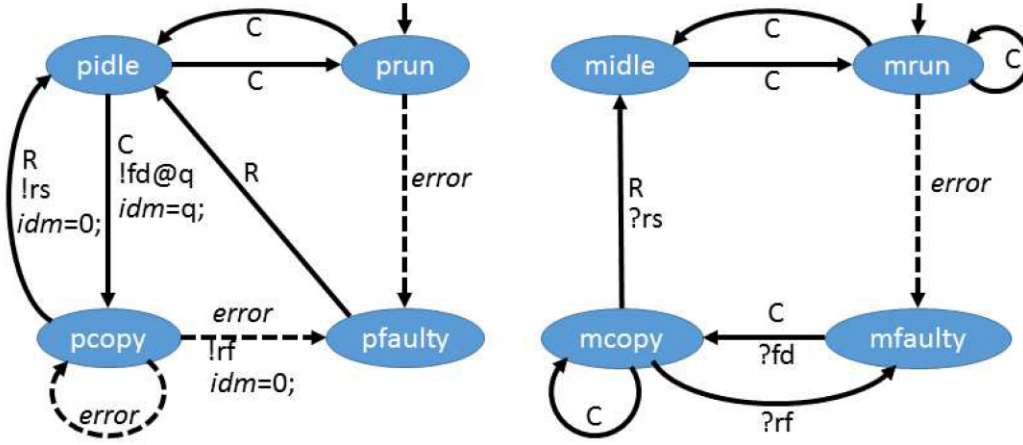


Figure 7.1: CEFSM templates of  $n$  processors and  $m$  memory copies

In the following, we explain how to apply our techniques to analyze the resilience levels of the avionic systems in Example 2. The application is achieved in three steps. We first model the system under analysis either as a plain CEFSM or with counter abstraction (if our analysis tool cannot handle the complexity of the plain CEFSM). We then build the product automaton of the CEFSM as the resilience game structure except for the move vectors. Finally, we convert the labels on the transitions of the product automaton to move vectors of the two players. Note that the moves may not correspond to the transition labels of the CEFSM.

### Step 1: the construction of the CEFSM

We first present the CEFSM model template of Example 2 in Figure 7.1. The CEFSM model has  $n$  processors and  $m$  memory modules. Figures 7.1(a) and (b) are for the abstraction of processors and memory copies, respectively. The ovals represent local states of a processor or a memory module, while the arrows represent transitions. The transitions of a CEFSM are labeled with ‘error’, ‘C’ (for Control), or ‘R’ (for recovery).

We also use synchronizers to bind process transitions. For example, when a memory module moves into a faulty state, an idle processor may issue an  $\text{fd}$  (error-detected) event and try to repair the module by copying memory contents from normal memory modules. Such error-detection is usually achieved with standard hardware. Note that the



benchmarks are models that reflect the recovery mechanism, abstracting away the details of the original systems. A central issue in the design of this recovery mechanism is then the resilience level of the controlled systems. We need three synchronizers: *fd* for error detection by a processor, *rs* for recovery success, and *rf* for recovery failure. The three synchronizers are used to bind a transition from a processor and another from a memory module into a synchronized transition. For example, a processor at state *pidle* and a memory module at state *mfaulty* may simultaneously enter their *pcopy* and *mcopy* states respectively through synchronizers *!fd* (for sending the synchronizer) and *?fd* (for receiving). We also conveniently use a variable *q* in this synchronized transition to capture the identifier of the memory module receiving the synchronizer. A transition without synchronization labels is considered a trivial synchronized transition. The transition system of the CEFSM operates with interleaving semantics at the abstraction level of the synchronized transitions.

For counter abstraction, we need four global variables *crp*, *cfp*, *crm*, and *cfm* respectively to keep track of the numbers of running processors, faulty processors, running memory modules, and faulty memory modules. We also need a local variable *idm* for each processor to record the faulty memory module identifier that the processor is responsible for recovery. We label the controllable, error, and recovery transitions respectively with ‘C’, ‘error’, and ‘R’. We also label each transition with synchronizers and actions. At any moment, the processors and the memory modules may enter their running states, execute a task, and generate the outcome. A processor starts its execution from state *prun* while a memory module starts from state *mrn*.

## Step 2: building the product automata

The product automata is a Kripke structure whose states are of is a vector  $[p_1, \dots, p_n, i_1, \dots, i_n, s_1, \dots, s_n]$  of  $2n + m$  elements. For all  $k$ ,  $p_k$  and  $i_k$  respectively represent the current location and the current *idm* value of processor  $k$  while  $s_k$  represents the current location of memory module  $k$ . Then interleaving semantics that each time only a global transition (a single local process transition without synchronizers or two local process transition bound by a




Figure 7.2: Counter abstraction of the CEFM templates of  $n$  processors and  $m$  memory copies

synchronizer) is executed is adopted to determine the transition relation from one state to another. Such techniques are standard in model construction. REDLIB can help in this regard by constructing the Kripke structure in an on-the-fly style to avoid the construction of those states not reachable from the initial state.

### Step 3: the labeling of the move vectors

After the second step, we have the game structure ready except for the move vectors on the transitions. We use  $E_1 = \{C, R, nop\}$ , where *nop* represents “no operation,” and  $E_2 = \{noerr, error\}$ . Then we use the following three rules to label move vectors.

- Every global transition with one component local process transition labeled with *error* is labeled with move vector  $[nop, error]$ .
- Every global transition with a component local process transition labeled with *R* is labeled with move vector  $[R, noerr]$ .
- All other global transitions are labeled with move vector  $[C, noerr]$ .

### Counter abstraction of the example

We also use the CEFM in figure 7.1 to explain counter abstraction. We need eight counter variables: *pr*, *pi*, *pc*, *pf*, *mr*, *mi*, *mc*, and *mf* to respectively record the number of processes in location *prun*, *pidle*, *pcopy*, *pfaulty*, *mr*un, *mi*dle, *mc*opy, and *mf*aulty in a state. Then the counter abstraction of the CEFM is in Figure 7.2. The initial state are

specified with constraint:  $pr = n \wedge pi = 0 \wedge pc = 0 \wedge pf = 0 \wedge mr = m \wedge mi = 0 \wedge mc = 0 \wedge mf = 0$  on the counters. The state in the product automata must satisfy the following constraints:  $pr + pi + pc + pf = n \wedge mr + mi + mc + mf = m$ . As can be seen, we do not care which processor is in the idle mode, in the running mode, etc., in this abstraction. Similarly, we do not care which memory module is in the idle mode, in the running mode, and etc. The local state transition only keeps tracks of the number of processors in each mode and the number of memory modules in each mode. We also do not care which processor is in charge of the recovery of which memory module. Such an abstraction can be done automatically.

The labeling of the move vectors on the transitions in the Kripke structure (product automaton) follows the same rules for the product automaton from the CEFSM in Figure 7.1.

## Analysis of the game structure

The majority outcome of the processors and memory copies is used as the outcome of the system. A processor may enter the faulty state. A memory module may also enter the faulty state. Processors may control to recover themselves or a faulty memory module by copying the contents of a functioning memory module to the faulty one. At any moment, we want to make sure that we can always recover to a global condition with the following two restrictions.

- There are at least two more processors in the running mode than the processors in the faulty mode.
- There are at least two more memory copies in the running mode than memory copies in the faulty mode.

Together, the failure condition is  $crp - cfp < 2 \vee crm - cfm < 2$ . That is, all states in the transition system satisfying  $crp - cfp < 2 \vee crm - cfm < 2$  are in set  $F$ .

We report the performance data in Table 7.4 for the resilience algorithms described in Section ?? against the parameterized benchmarks in the above with various parameters.

Table 7.4: Performance data for resilience calculation  
megabytes

s: seconds; M:

benchmarks	concurrency	$k$	game sizes		sfrch <sub>k</sub>		res <sub>k</sub>	
			#nodes	#edges	time	memory	time	memory
avionics	2 processors & 2 memory modules	2	118	750	0.62s	114M	0.85s	116M
	2 processors & 3 memory modules	2	414	3252	0.94s	139M	1.10s	153M
	3 processors & 3 memory modules	3	1540	15090	4.67s	225M	8.38s	267M
	3 processors & 4 memory modules	3	5601	63889	42.86s	815M	155s	846M
avionics (counter abstraction)	6 processors & 6 memory modules	2	1372	6594	2.89s	129M	3.54s	516M
	7 processors & 7 memory modules	3	2304	11396	10.7s	216M	23.4s	808M
	8 processors & 8 memory modules	3	3645	18432	43.8s	1009M	135s	2430M
voting	1 client & 20 replicas	9	9922	23551	7.01s	260M	36.7s	297M
	1 client & 26 replicas	12	20776	49882	19.9s	474M	79.6s	611M
simple voting	1 client & 150 replicas	74	458	1056	0.71s	159M	31.7s	219M
	1 client & 200 replicas	99	608	1406	1.06s	161M	162s	337M
	1 client & 250 replicas	124	758	1756	1.36s	163M	307s	499M
PBFT	1 client & 6 replicas	2	577	897	0.34s	72M	1.05s	193M
	1 client & 9 replicas	4	2817	4609	13.3s	564M	58.5s	1657M
clock	1 client & 15 servers	7	16384	229376	45.1s	3075M	62.4s	3264M
sync	1 client & 17 servers	8	65536	1070421	870s	14725M	915s	15433M

The second column shows the concurrency sizes. The third column shows the values of  $k$  for the rows. The fourth and fifth columns show the sizes of the concurrent game structures. The sixth and seventh columns show the time and spaces used to calculate  $\text{sfrch}_k()$ . Similarly, the eighth and ninth columns show the time and spaces for calculating the  $\text{res}_k()$ .

The benchmark in Figure 7.1 does not have nodes in  $\text{sfrch}_2(G)$  and  $\text{res}_2(G)$ . So we changed the benchmark to see how we check our implementation with  $k > 1$ . The change is that the recovery transition from state *pcopy* to *pidle* of processors are relabeled as controllable. This change significantly limits the ability of the system errors to derail the system. For the avionics system, the resilience level  $k$  is set to one less than half the number of processors. For the voting and simple voting benchmarks, the value of  $k$  is set to one less than half the number of replicas (voters). For the PBFT and clock synchronization algorithm, we choose  $k$  to be one less than one third of the number of replicas.

The performance data has been collected with a Virtual Machine (VM) running open-suse 11.4 x86 on Intel i7 2600k 3.8GHz CPU with 4 cores and 8G memory. The VM only uses one core and 4G memory.

The time and space used to calculate resilience is a little bit more than that to check for  $\text{sfrch}$ . The reason is that  $\text{sfrch}_k$  is a pre-requisite for calculating  $\text{res}_k$ . In our experiment,

$\text{sfrch}_k$  is usually very close to  $\text{res}_k$  and does not require much extra time in calculating  $\text{res}_k$  out of  $\text{sfrch}_k$ .

The experiments show that our techniques scale to realistic levels of redundancy. For fault-tolerant hardware, usually the numbers of replicas are small, for example, less than 10 replicas. Thus our techniques seem very promising for the verification and synthesis of hardware fault-tolerance.

On the other hand, nowadays, software fault-tolerance through networked computers can create huge numbers of replicas. Our experiment shows that counter abstraction can be a useful techniques for the modeling and verification of software resilience. Specifically, for the avionics benchmark, we can verify models of much higher concurrency and complexity with counter abstraction than without.

# Chapter 8

## Conclusions

### 8.1 Conclusion

BSIL and TCL can be useful in describing combinations of strategy profiles in a multi-agent system. They are carefully designed for low model-checking cost while maintaining sufficient expressiveness to specify some useful properties. We have thoroughly investigated the theoretical aspects of BSIL and TCL, including the expressive power, the model checking complexity. Our experiment report also points out new research directions, including further extension to BSIL&TCL and performance-enhancing techniques to our model-checking algorithms.

We have introduced an approach for the development of a control of safety critical systems that maximizes the number of *dense* errors the system can tolerate. Our techniques are inspired by the problem of controlling systems with redundancy: in order to deflect the effect of individual errors, safety critical systems are often equipped with multiple copies of various components. If one or more components fail, such systems can still work properly as long as the correct behavior can be identified.

This has inspired the two-phase formulation of the safety resilience problems in this article. In the first phase, we identify a  $k$ -resilient region, while we develop a control strategy for recovery in the second phase. After an error, the controller can recover to the  $k$ -resilient region without encountering a system failure, unless the error is part of a group of more than  $k$  errors that happen in close succession. Such a recovering strategy

is memoryless. Being memoryless on a small abstraction in particular implies that the recovery is fast.

The system can, once recovered, tolerate and recover from  $k$  further dense errors, and so forth. Consequently, our control strategy allows for recovery from an arbitrary number of errors, provided that the number of dense errors is restricted. This is the best guarantee we can hope for: our technique guarantees to find the optimal parameter  $k$ .

This parameter is bound to be small (smaller than the number of redundant components). Optimizing it is computationally inexpensive, but provides strong guarantees: the likelihood of having more than  $k$  errors appear in short succession after an error occurred are, for independent errors, exponential in  $k$ . As errors are few and far between, each level of resilience gained reduces the likelihood of system-level failures significantly.

# Bibliography

- [1] Ian Sommerville. *Software Engineering: (Update) (8th Edition) (International Computer Science)*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2006.
- [2] John M. Rushby. Formal specification and verification of a fault-masking and transient-recovery model for digital flight-control systems. In Jan Vytopil, editor, *FTRTFT*, volume 571 of *Lecture Notes in Computer Science*, pages 237–257. Springer, 1992.
- [3] Rajeev Alur, Thomas A. Henzinger, and Orna Kupferman. Alternating-time temporal logic. *Journal of the ACM (JACM)*, 49(5):672–713, September 2002.
- [4] Arnaud Da Costa, François Laroussinie, and Nicolas Markey. Atl with strategy contexts: Expressiveness and model checking. In *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS 2010)*, volume 8 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 120–132. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2010.
- [5] Krishnendu Chatterjee, Thomas A. Henzinger, and Nir Piterman. Strategy logic. *Information and Computation*, 208:677–693, 2010.
- [6] Fabio Mogavero, Aniello Murano, and Moshe Y. Vardi. Reasoning about strategies. In *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS 2010)*, volume 8 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 133–144. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2010.



- [7] Lawrence J. Stockmeyer. *The complexity of decision problems in automata theory and logic*. MIT, 1974.
- [8] A.P. Sistla and E.M. Clarke. The complexity of propositional linear temporal logics. *JACM*, 32, 1985.
- [9] Edmund m. clarke and ernest allen emerson and a. prasad sistla.
- [10] E. Allen Emerson and Joe Y. Halpern. Decision procedures and expressiveness in the temporal logic of branching time. *Journal of Computer and System Sciences*, 30(1):1–24, Feb. 1985.
- [11] E. Allen Emerson and Joe Y. Halpern. ‘sometimes’ and ‘not never’ revisited: On branching versus linear time temporal logic. *Journal of ACM*, 33(1):151–178, Jan. 1986.
- [12] E. Allen Emerson and Chin-Laung Lei. Modalities for model checking: Branching time logic strikes back. *Science of Computer Programming*, 8:275–306, 1987.
- [13] M. Ben-Ari, Amir Pnueli, and Zohar Manna. The temporal logic of branching time. *Acta Informatica*, 20:207–226, 1983.
- [14] E. Allen Emerson and Edmund M. Clarke. Characterizing correctness properties of parallel programs as fixpoints. In *7th Colloquium on Automata, Language, and Programming*, volume LNCS 85. Springer-Verlag, 1980.
- [15] Leslie Lamport. Sometimes is sometimes “not never”—on the temporal logic of programs. In *7th Annual ACM Symp. on Principles of Programming Languages*, pages 174–185, 1980.
- [16] Christel Baier, Tomas Brázdil, Marcus Gröser, and Antonín Kucera. Stochastic game logic. In *QEST*, pages 227–236. IEEE Computer Society, 2007.
- [17] D.K. Pradhan. *Fault-tolerant computer system design*. Prentice-Hall, Inc., 1996.

- [18] Rüdiger Ehlers and Ufuk Topcu. Resilience to intermittent assumption violations in reactive synthesis. In Martin Fränzle and John Lygeros, editors, *Proceedings of the 17th International Conference on Hybrid Systems: Computation and Control (part of CPS Week), HSCC'14, Berlin, Germany, April 15-17, 2014*, pages 203–212. ACM, 2014.
- [19] Roderick Bloem, Rüdiger Ehlers, Swen Jacobs, and Robert Könighofer. How to handle assumptions in synthesis. In Krishnendu Chatterjee, Rüdiger Ehlers, and Susmit Jha, editors, *Proceedings of the 3rd Workshop on Synthesis, SYNT 2014, Vienna, Austria, July 23-24, 2014*, volume 157 of *EPTCS*, pages 34–50, 2014.
- [20] Chung-Hao Huang, Doron Peled, Sven Schewe, and Farn Wang. Rapid recovery for systems with scarce faults. In *Proceedings of the 3rd International Symposium on Games, Automata, Logics and Formal Verification (GandALF 2012), 6–8 September, Naples, Italy*, volume 96 of *Electronic Proceedings in Theoretical Computer Science*, pages 15–28, 2012.
- [21] N. Immerman. Number of quantifiers is better than number of tape cells. *Journal of Computer and System Sciences*, 22(3):65–72, 1981.
- [22] C.H. Papadimitriou. *Computational Complexity*. Addison-Wesley, 1994.
- [23] Farn Wang. Model-checking fair dense-time systems with propositions and events. *International Journal on Software Tools for Technology Transfer*, April 2014.
- [24] Farn Wang. Model-checking distributed real-time systems with states, events, and multiple fairness assumptions. In *International Conference on Algebraic Methodology and Software Technology (AMAST)*, volume LNCS 3116. Springer-Verlag, 2004.
- [25] N. Immerman. Nondeterministic space is closed under complementation. *SIAM Journal on Computing*, 17:935–938, 1988.
- [26] Ferenc Belina and Dieter Hogrefe. The ccitt-specification and description language sdl. *Computer Networks and ISDN Systems*, 16(4):311–341, March 1989.

- [27] M. Castro and B. Liskov. Practical byzantine fault tolerance. In *Proceedings of the 3rd USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, February 1999.
- [28] P. Ramanathan, K. G. Shin, and R. W. Butler. Fault-tolerant clock synchronization in distributed systems. *IEEE Computer*, 23(10):33–42, 1990.
- [29] D.M. Miller and R. Drechsler. Implementing a multiple-valued decision diagram package. In *The 28th International Symposium on multiple-valued logic (ISMVL)*, pages 52–57, 1998.
- [30] M. Abd-El-Malek, G. Ganger, G. Goodson, M. Reiter, and J. Wylie. Fault-scalable byzantine fault-tolerant services. In *ACM Symposium on Operating Systems Principles*, 2005.
- [31] J. Cowling, D. Myers, B. Liskov, R. Rodrigues, and L. Shriram. Hq replication: A hybrid quorum protocol for byzantine fault tolerance. In *USENIX OSDI Symposium*, 2006.
- [32] R. Kotla, L. Alvisi, M. Dahlin, A. Clement, and E. Wong. Zyzzyva: Speculative byzantine fault tolerance. *ACM Transactions on Computer Systems*, 27(4), 2009.
- [33] R. Guerraoui, N. Knežević, M. Vukolić, and V. Quéma. The next 700 bft protocols. In *Proceedings of the 5th European conference on Computer systems (EuroSys)*, 2010.
- [34] A. Clement, E. Wong, L. Alvisi, M. Dahlin, and M. Marchetti. Making byzantine fault tolerant systems tolerate byzantine faults. In *USENIX Symposium on Networked Systems Design and Implementation*, April 22-24, 2009.
- [35] Bitcoin website. "<http://bitcoin.org/>".
- [36] E. Allen Emerson and R. Trefler. From asymmetry to full symmetry: New techniques for symmetry reduction in model-checking. In *Correct Hardware Design and Verification Methods (CHARME)*, pages 142–156, 1999.

- [37] B. Lubachevsky. An approach to automating the verification of compact parallel coordination programs. *Acta Informatica*, 1984.
- [38] A. Church. Logic, arithmetic and automata. In *International Congress of Mathematicians*, pages 23–35, August 1962 (Stockholm 1963).
- [39] Amir Pnueli and Roni Rosner. On the synthesis of a reactive module. In *16th ACM SIGPLAN-SIGACT symposium on principles of programming languages (POPL)*, pages 179–190, 1989.
- [40] M.O. Rabin. Decidability of second order theories and automata on infinite trees. *Transaction of the American Mathematical Society*, 141:1–35, 1969.
- [41] J.R. Büchi. On a decision method in restricted second order arithmetic. In *International Congress on Logic, Methodology, and Philosophy of Science*, pages 1–11, 1962.
- [42] J.R. Büchi and L.H. Landweber. Solving sequential conditions by finite-state strategies. *Trans. AMS*, 138(4):295–311, 1969.
- [43] P.C. Attie, A. Arora, and E. Allen Emerson. Synthesis of fault-tolerant concurrent programs. *ACM Transactions on Programming Languages and Systems*, 26(1):125–185, 2004.
- [44] A. Ebzenasir, S.S. Kulkarni, and A. Arora. Ftsyn: a framework for automatic synthesis of fault-tolerance. *STTT*, 10(5):455–471, 2008.
- [45] A. Girault and E. Rutten. Automating the addition of fault tolerance with discrete controller synthesis. *Formal Methods in System Design*, 35(2):190–225, 2009.
- [46] Orna Kupferman and Moshe Y. Vardi. Synthesis with incomplete informatio. In *Advances in Temporal Logic*, pages 109–127. Kluwer Academic Publishers, 2000.
- [47] Amir Pnueli and Roni Rosner. On the synthesis of an asynchronous reactive module. In *Proceedings of the 16th International Colloquium on Automata, Languages*

*and Programming (ICALP)*, volume LNCS 372, pages 652–671, London, UK, 1989. Springer-Verlag.

- [48] J.M. Rushby. Formal specification and verification of a fault-masking and transient-recovery model for digital flight-control systems. In *FTRTFT*, pages 237–257, 1992.
- [49] S. Schewe and B. Finkbeiner. Synthesis of asynchronous systems. In *LOPSTR*, pages 127–142, 2006.
- [50] H. Jin, K. Ravi, and Fabio Somenzi. Fate and free will in error traces. *International Journal on Software Tools for Technology Transfer (STTT)*, 6(2):102–116, 2004.
- [51] Eugene Asarin, Oded Maler, and Amir Pnueli. Symbolic controller synthesis for discrete and timed systems. In *Hybrid Systems II*, volume LNCS 999. Springer-Verlag, 1995.
- [52] Wolfgang Thomas. Finite-state strategies in regular infinite games. In *FSTTCS*, pages 149–158, 1994.
- [53] A. Arora and M.G. Gouda. Closure and convergence: A foundation of fault-tolerant computing. *IEEE Transactions on Software Engineering (TSE)*, 19(11):1015–1027, 1993.
- [54] Edgar W. Dijkstra. A belated proof of self-stabilization. *Distributed Computing*, 1(1):5–6, 1986.
- [55] L. Doyen, T.A. Henzinger, A. Legay, and D. Nickovic. Robustness of sequential circuits. In *ACSD*, pages 77–84, 2010.
- [56] R. Bloem, K. Greimel, and Barbara Jobstmann T.A. Henzinger. Synthesizing robust systems. In *FMCAD*, pages 85–92, 2009.
- [57] J. Fearnley, M.N. Rabe, Sven Schewe, and L. Zhang. Efficient approximation of optimal control for continuous-time markov games. In *Proc. 31st IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS)*, pages 399–410, 2011.

- [58] Orna Kupferman, Moshe Y. Vardi, and Pierre Wolper. Module checking. *Information and Computation*, 164(2):322–344, 2001.
- [59] T. Brihaye, A.D.C. Lopes, F. Laroussinie, and Nicolas Markey. Atl with strategy contexts and bounded memory. In *LFCS*, volume LNCS 5407, pages 92–106. Springer-Verlag, 2009.
- [60] S. Pinchinat. A generic constructive solution for concurrent games with expressive constraints on strategies. In *Automated Technology for Verification and Analysis (ATVA)*, volume LNCS 4762, pages 253–267. Springer-Verlag, 2007.
- [61] François Laroussinie and Nicolas Markey. Satisfiability of atl with strategy contexts. In *Workshop on Games, Automata, Logics and Formal Verification (GANDALF)*, volume EPTCS 119, pages 208–223, 2013.
- [62] Fabio Mogavero, Aniello Murano, Giuseppe Perelli, and Moshe Y. Vardi. What makes atl\* decidable ? a decidable fragment of strategy logic. In *Concurrency theory (CONCUR 2012)*, volume LNCS 7454, pages 193–208. Springer-Verlag, 2012.
- [63] Fabio Mogavero, Aniello Murano, and Luigi Sauro. On the boundary of behavioral strategies. In *ACM/IEEE LICS*, 2013.
- [64] Michael R. Garey and David S. Johnson. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. W. H. Freeman & Company, 1979.
- [65] Daniel Kirsten. Alternating tree automata and parity games. In Erich Gradel, Wolfgang Thomas, and Thomas Wilke, editors, *Automata, Logics, and Infinite Games*, volume LNCS 2500, pages 3?V–21. Springer, 2002.
- [66] Júlia Zappe. Modal  $\mu$ -calculus and alternating tree automata. In Erich Gradel, Wolfgang Thomas, and Thomas Wilke, editors, *Automata, Logics, and Infinite Games*, volume LNCS 2500, pages 3?V–21. Springer, 2002.

- [67] David E. Muller and Paul E. Schupp. Simulating alternating tree automata by non-deterministic automata: new results and new proofs of the theorems of Rabin, McNaughton and Safra. *Theoretical Computer Science*, 141(1-2):69–107, 1995.
- [68] Shmuel Safra. On the complexity of  $\omega$ -automata. In *Proceedings of the 29th Annual Symposium on Foundations of Computer Science (FOCS 1988), 24–26 October, White Plains, New York, USA*, pages 319–327. IEEE Computer Society Press, 1988.
- [69] Nir Piterman. From nondeterministic Büchi and Streett automata to deterministic parity automata. *Journal of Logical Methods in Computer Science*, 3(3:5), 2007.
- [70] Sven Schewe. Tighter bounds for the determinisation of Büchi automata. In *Proceedings of the Twelfth International Conference on Foundations of Software Science and Computation Structures (FoSSaCS 2009), 22–29 March, York, England, UK*, volume 5504 of *Lecture Notes in Computer Science*, pages 167–181. Springer-Verlag, 2009.
- [71] Thomas Wilke.  $\text{CTL}^+$  is exponentially more succinct than CTL. pages 110–121, 1999.
- [72] Farn Wang. Efficient verification of timed automata with BDD-like data-structures. *International Journal of Software Tools for Technology Transfer (STTT)*, 6(1), 2004. special issue for the 4th International Conference on Verification, Model Checking, and Abstract Interpretation (VMCAI), Jan. 2003, LNCS 2575, Springer-Verlag.
- [73] Farn Wang. Time-progress evaluation for dense-time automata with concave path conditions. In *Automated Technology for Verification and Analysis (ATVA)*, volume LNCS 5311. Springer-Verlag, 2008.
- [74] Farn Wang. Efficient model-checking of dense-time systems with time-convexity analysis. In *IEEE Real-Time System Symposium (RTSS)*. IEEE Computer Society, 2008.
- [75] Farn Wang. Efficient model-checking of dense-time systems with time-convexity analysis. *Theoretical Computer Science (TCS)*, 467:89–V108, 7 January 2013.

- [76] Farn Wang, Chung-Hao Huang, and Fang Yu. A temporal logic for the interaction of strategies. In *22nd Concurrency Theory (CONCUR)*, volume LNCS 6901. Springer-Verlag, Sept. 2011.
- [77] M Y Vardi and L Stockmeyer. Improved upper and lower bounds for modal logics of programs. In *Proceedings of the Seventeenth Annual ACM Symposium on Theory of Computing*, STOC '85, pages 240–251, New York, NY, USA, 1985. ACM.
- [78] Larry J. Stockmeyer and Ashok K. Chandra. Provably difficult combinatorial games. *SIAM Journal on Computing (SICOMP)*, 8(2):151–174, 1979.
- [79] David E. Muller and Paul E. Schupp. Simulating alternating tree automata by non-deterministic automata: New results and new proofs of the theorems of rabin, mc-naughton and safra. *Theor. Comput. Sci.*, 141(1&2):69–107, 1995.
- [80] Krishnendu Chatterjee and Monika Henzinger. An  $o(n^2)$  time algorithm for alternating büchi games. In *Proceedings of the Twenty-Third Annual ACM-SIAM Symposium on Discrete Algorithms (SODA 2012)*, Kyoto, Japan, January 17-19, 2012, pages 1386–1399. SIAM, 2012.
- [81] Sven Schewe. Solving parity games in big steps. In *Proceedings of the 27th Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS 2007)*, 12–14 December, New Delhi, India, volume 4805 of *Lecture Notes in Computer Science*, pages 449–460. Springer-Verlag, 2007.